

CS 5112

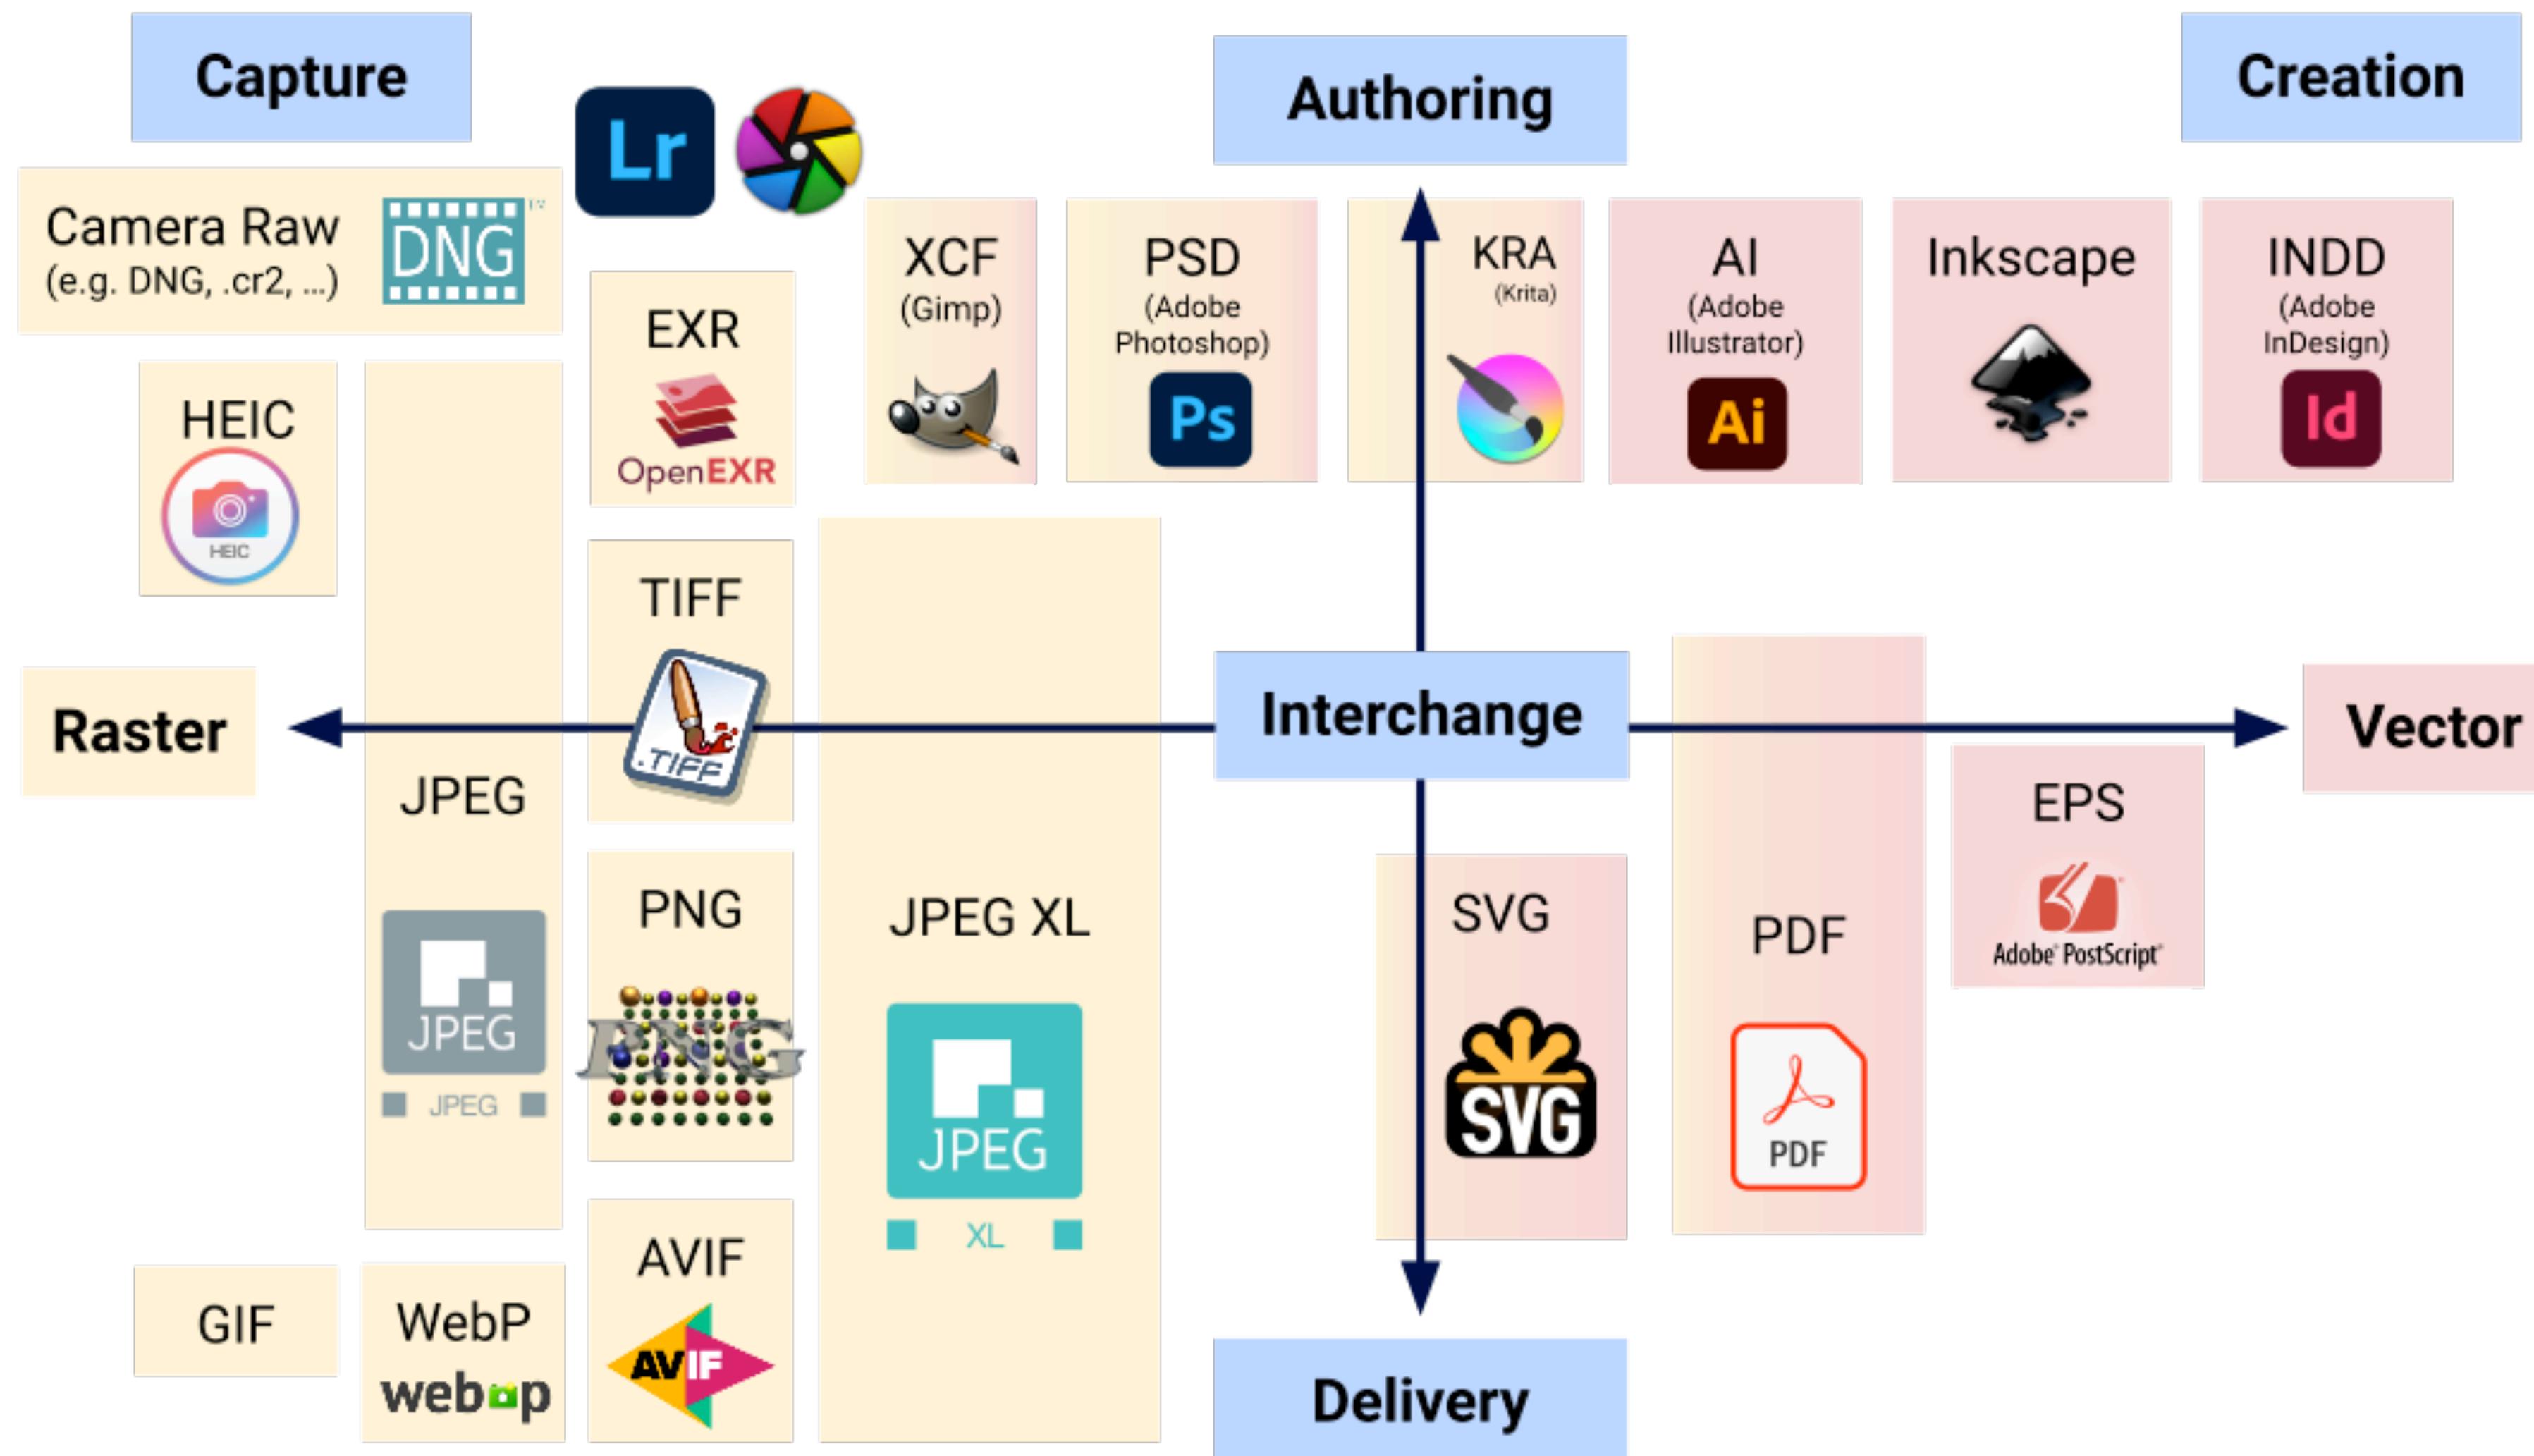
Algorithms for Applications

Huffman Coding Application:

JPEG Compression

How do you store an image in a file?

How do you store an image in a file?



How do you store an image in a file?

What's the simplest way?

How do you store an image in a file?

What's the simplest way?

```
3 3 3 3 3 3 3 3  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
0 5 5 5 5 5 5 0  
0 5 5 5 5 5 5 0  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
2 2 2 2 2 2 2 2
```



0	000000	
1	FF0000	
2	00FF00	
3	0000FF	
4	FFFFFF	
5	FFFF00	
6	FF00FF	
7	00FFFF	
8	FF0080	
9	FF8040	
A	804000	
B	008080	
C	800000	
D	800080	
E	8080FF	

Uncompressed bitmap

How do you store an image in a file?

What's the simplest way?

```
3 3 3 3 3 3 3 3  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
0 5 5 5 5 5 5 0  
0 5 5 5 5 5 5 0  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
2 2 2 2 2 2 2 2
```



0	000000	
1	FF0000	
2	00FF00	
3	0000FF	
4	FFFFFF	
5	FFFF00	
6	FF00FF	
7	00FFFF	
8	FF0080	
9	FF8040	
A	804000	
B	008080	
C	800000	
D	800080	
E	8080FF	

Store the values of each pixel in an array

Uncompressed bitmap

How do you store an image in a file?

What's the simplest way?

```
3 3 3 3 3 3 3 3  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
0 5 5 5 5 5 5 0  
0 5 5 5 5 5 5 0  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
2 2 2 2 2 2 2 2
```



0	000000	
1	FF0000	
2	00FF00	
3	0000FF	
4	FFFFFF	
5	FFFF00	
6	FF00FF	
7	00FFFF	
8	FF0080	
9	FF8040	
A	804000	
B	008080	
C	800000	
D	800080	
E	8080FF	

Store the values of
each pixel in an array

Problem:
these are large

Uncompressed bitmap

How do you store an image in a file?

What's the simplest way?

```
3 3 3 3 3 3 3 3  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
0 5 5 5 5 5 5 0  
0 5 5 5 5 5 5 0  
0 1 4 1 4 1 4 0  
0 4 1 4 1 4 1 0  
2 2 2 2 2 2 2 2
```



0	000000	
1	FF0000	
2	00FF00	
3	0000FF	
4	FFFFFF	
5	FFFF00	
6	FF00FF	
7	00FFFF	
8	FF0080	
9	FF8040	
A	804000	
B	008080	
C	800000	
D	800080	
E	8080FF	

Uncompressed bitmap

Store the values of each pixel in an array

Problem:
these are large

Solution:
use compression

Compressed Bitmaps

Take a bitmap and apply lossless compression



Lempel-Ziv-Welch (LZW)



DEFLATE
(LZ77 and Huffman coding)

Compressed Bitmaps

Take a bitmap and apply lossless compression



Lempel-Ziv-Welch (LZW)



DEFLATE
(LZ77 and Huffman coding)

JPEG Files

Joint Photographic Experts Group

Introduced in 1992

Most widely used image compression standard

“JPEG was largely responsible for the proliferation of digital images and digital photos across the Internet and later social media”



JPEG

JPEG Files

Joint Photographic Experts Group

Introduced in 1992

Most widely used image compression standard

“JPEG was largely responsible for the proliferation of digital images and digital photos across the Internet and later social media”

Why?

What's so great about JPEG?

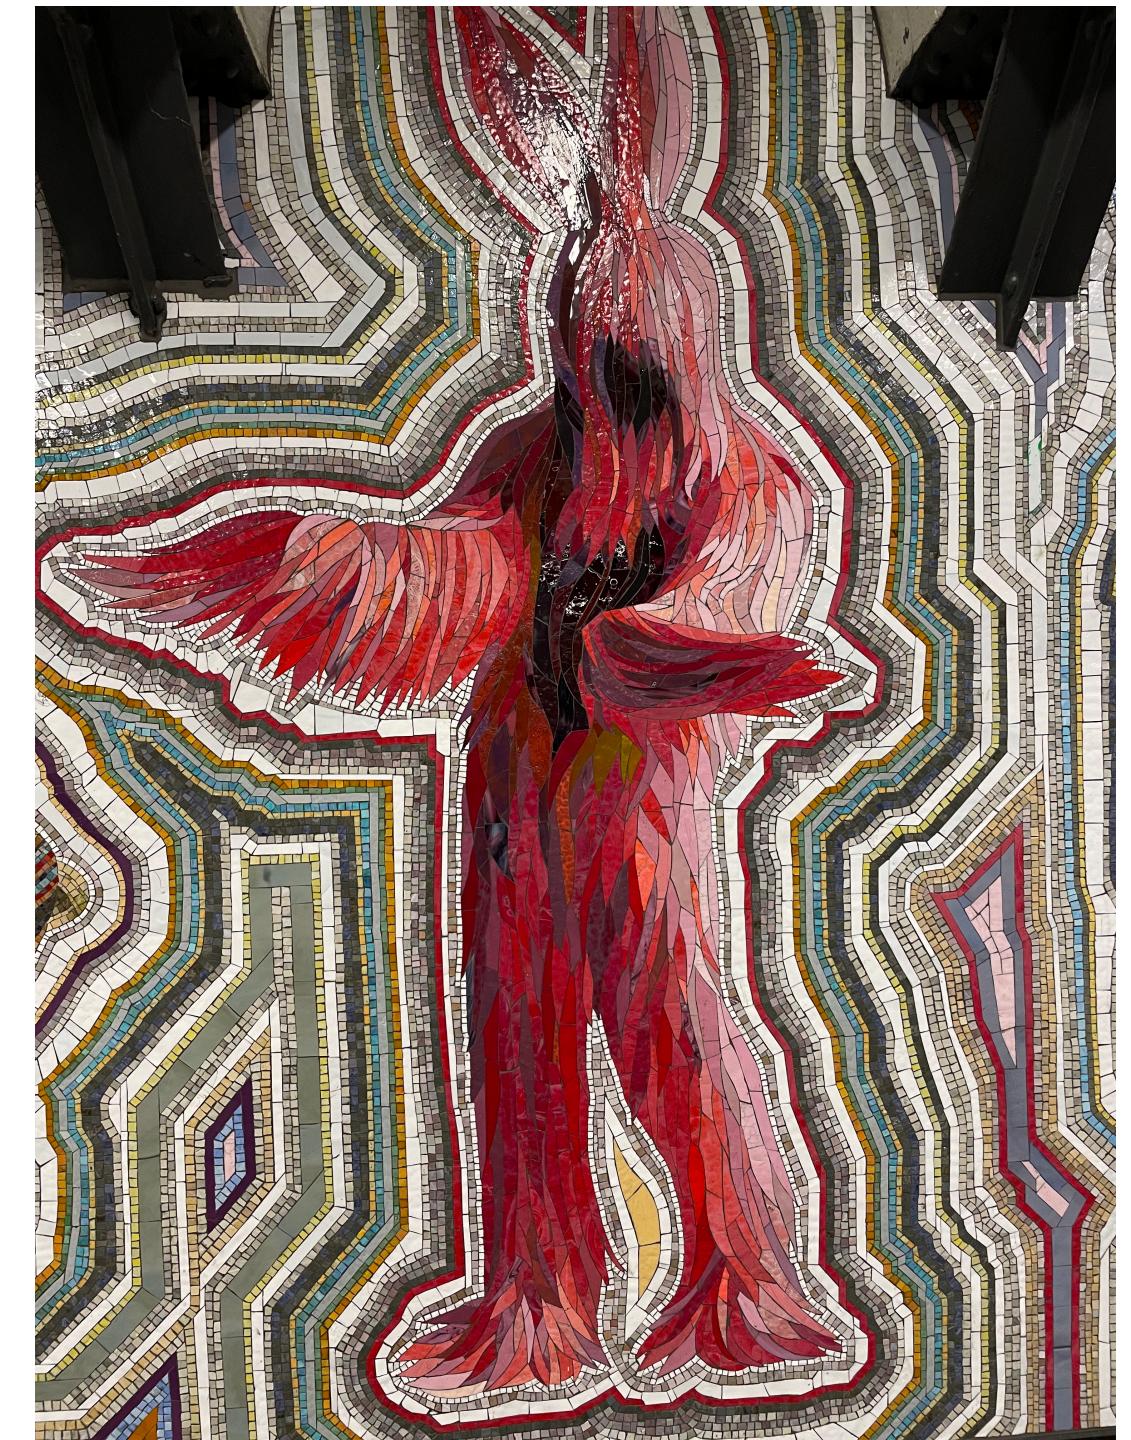
Image File Size



Bitmap (BMP)
37MiB

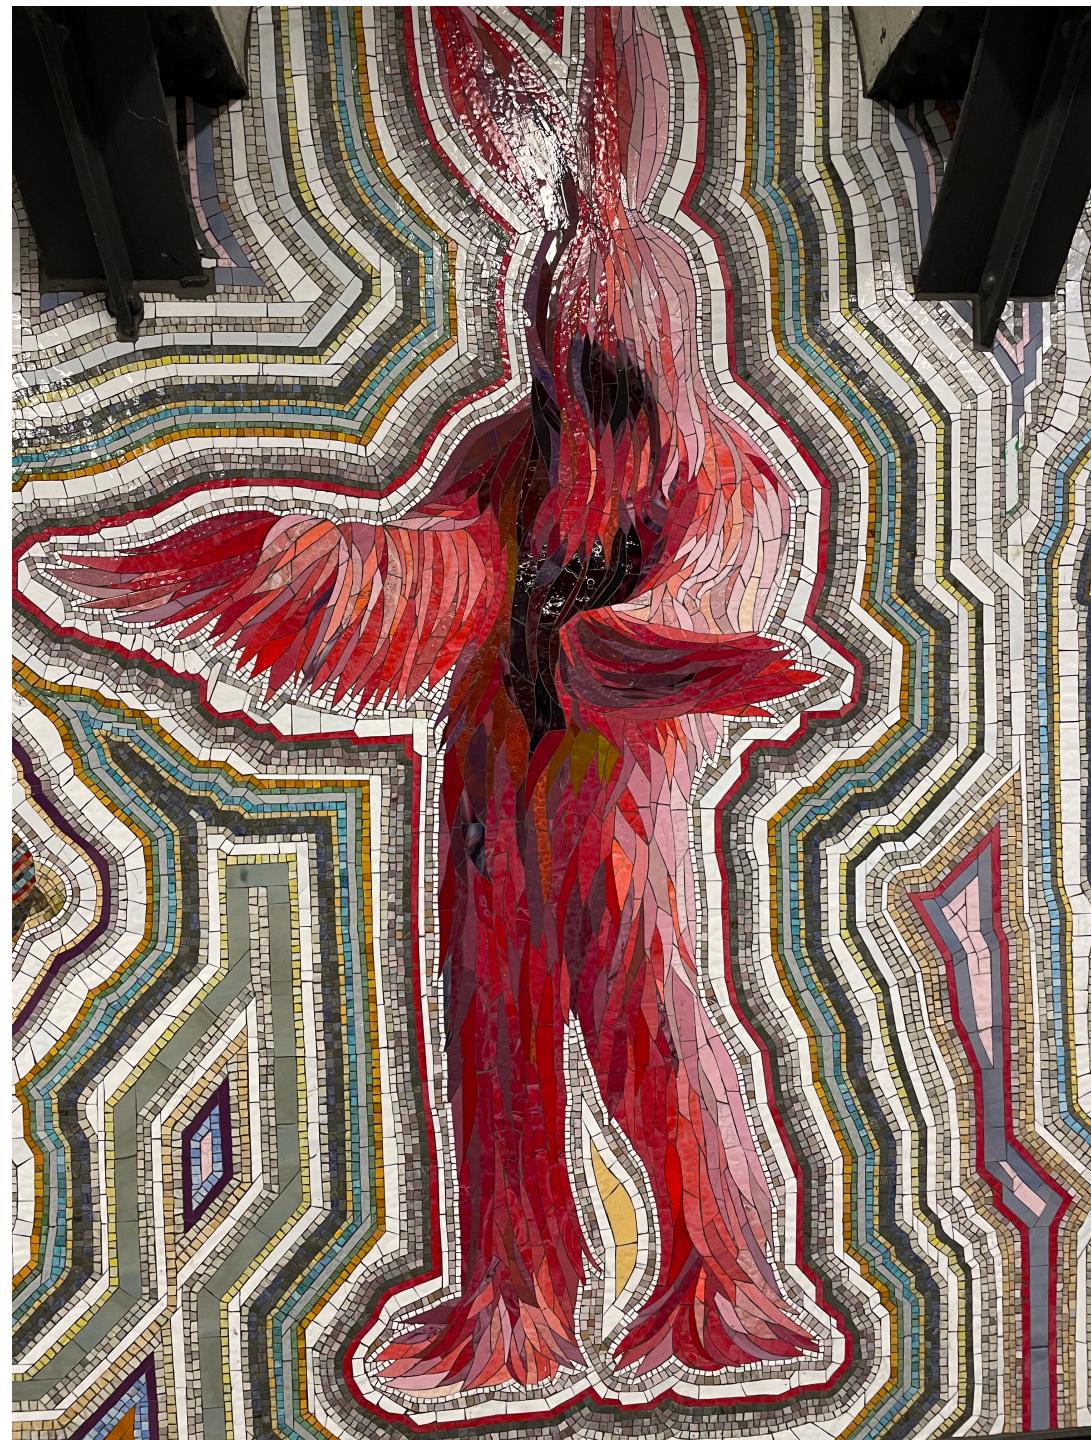


Compressed Bitmap (PNG)
21MiB

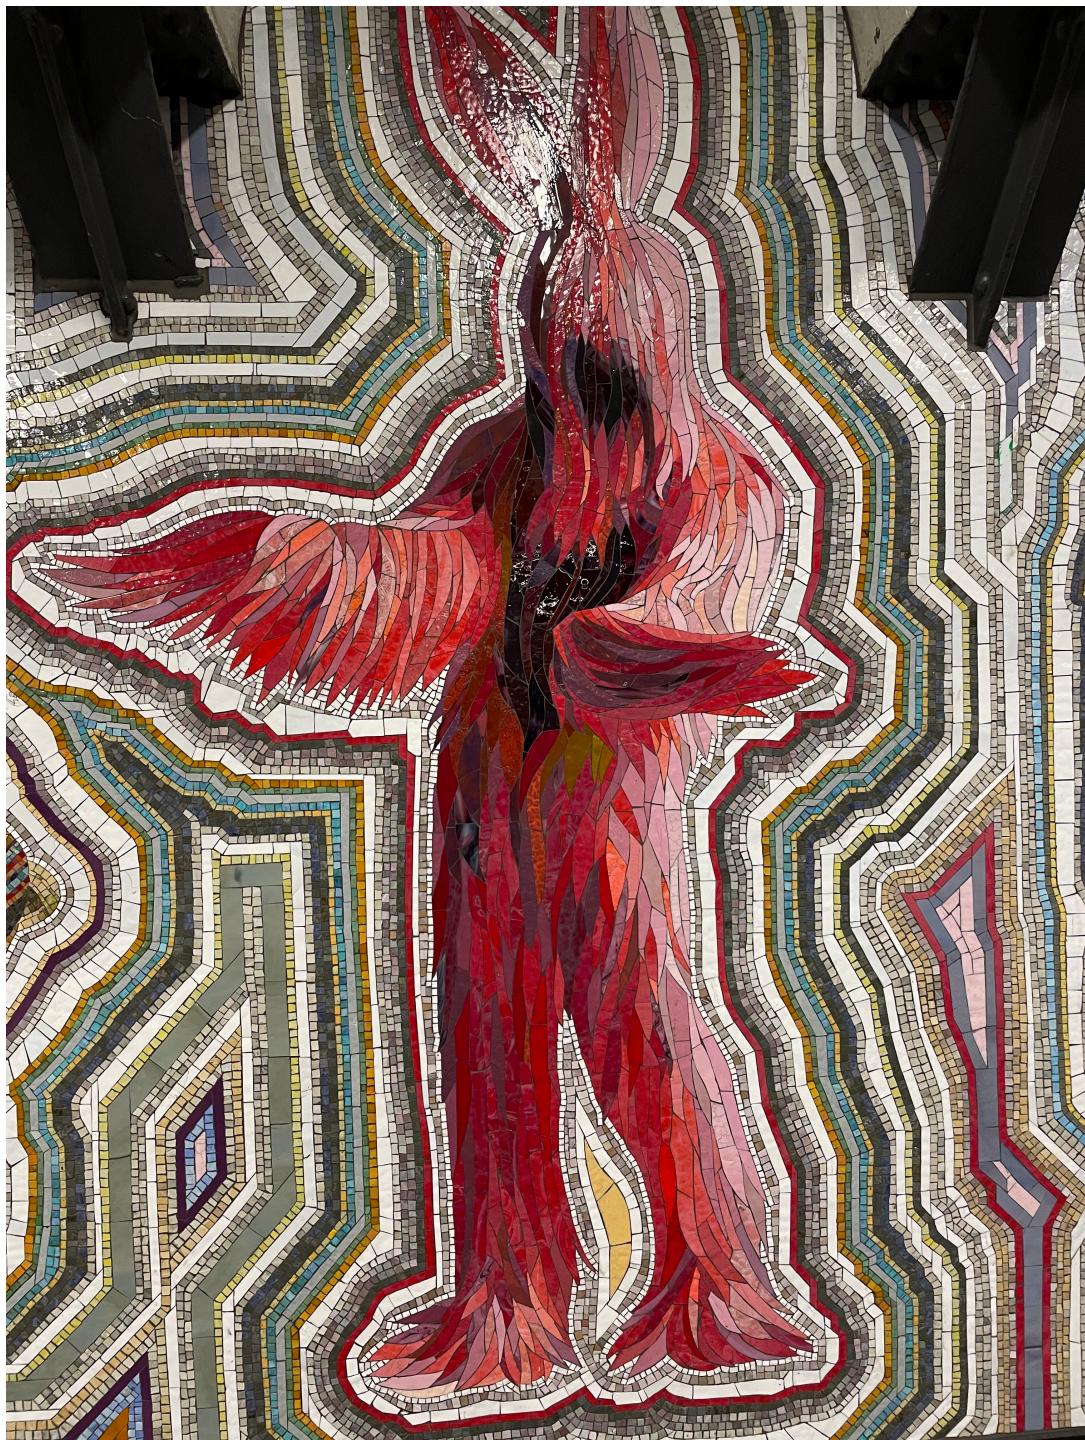


JPEG
12MiB

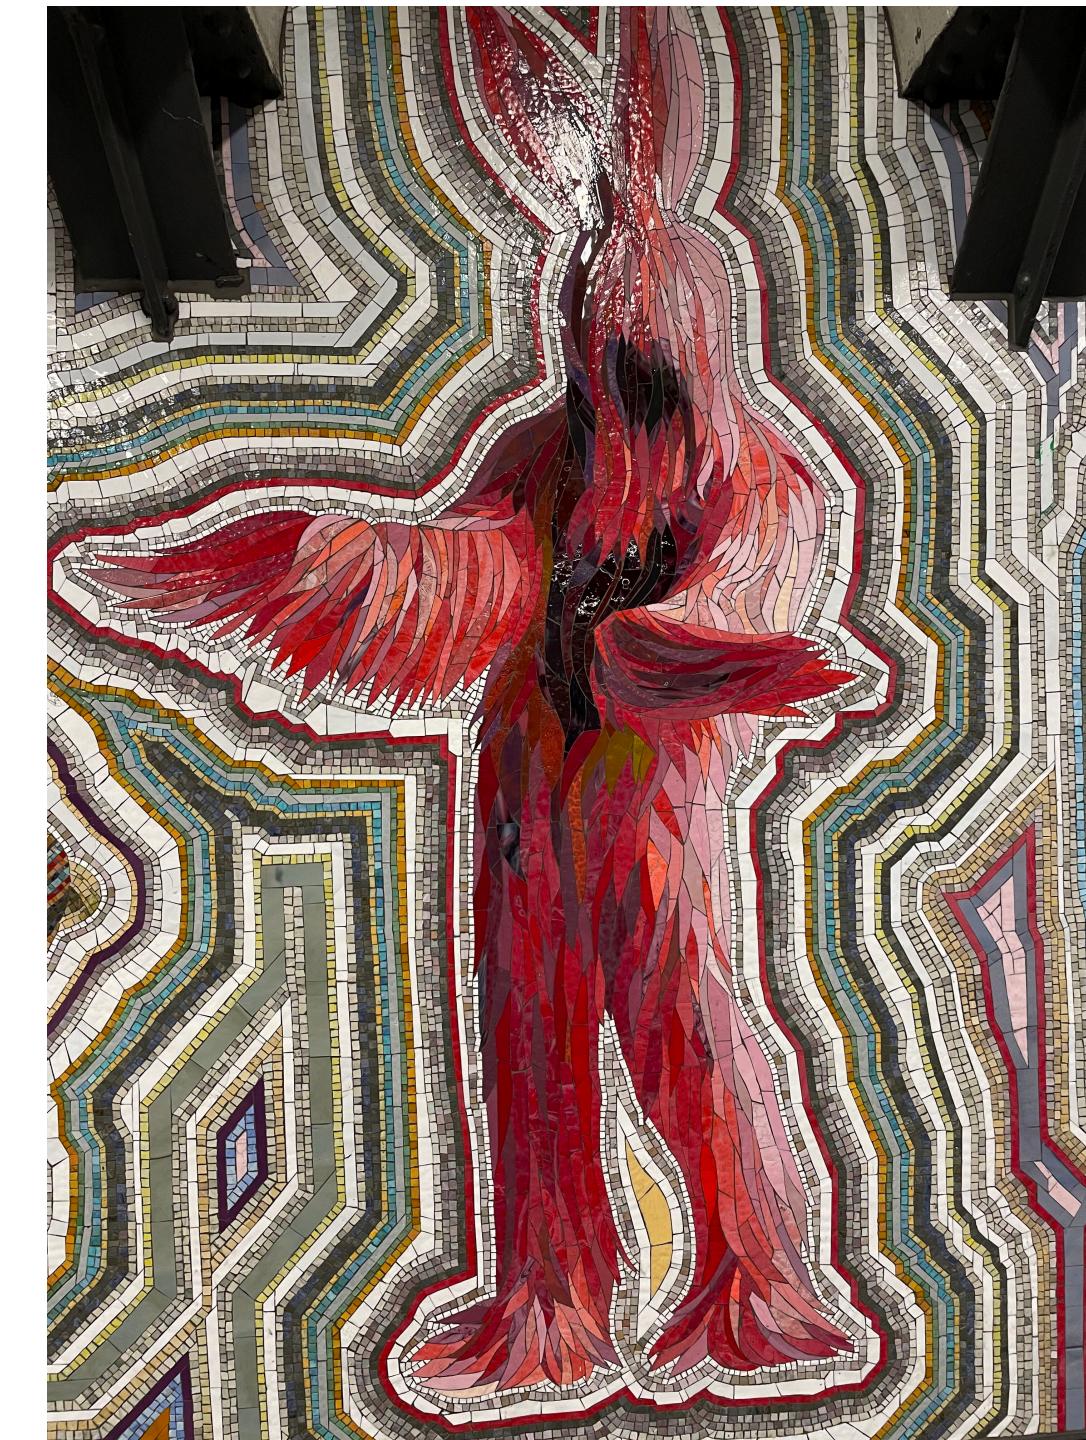
JPEG Variable Compression



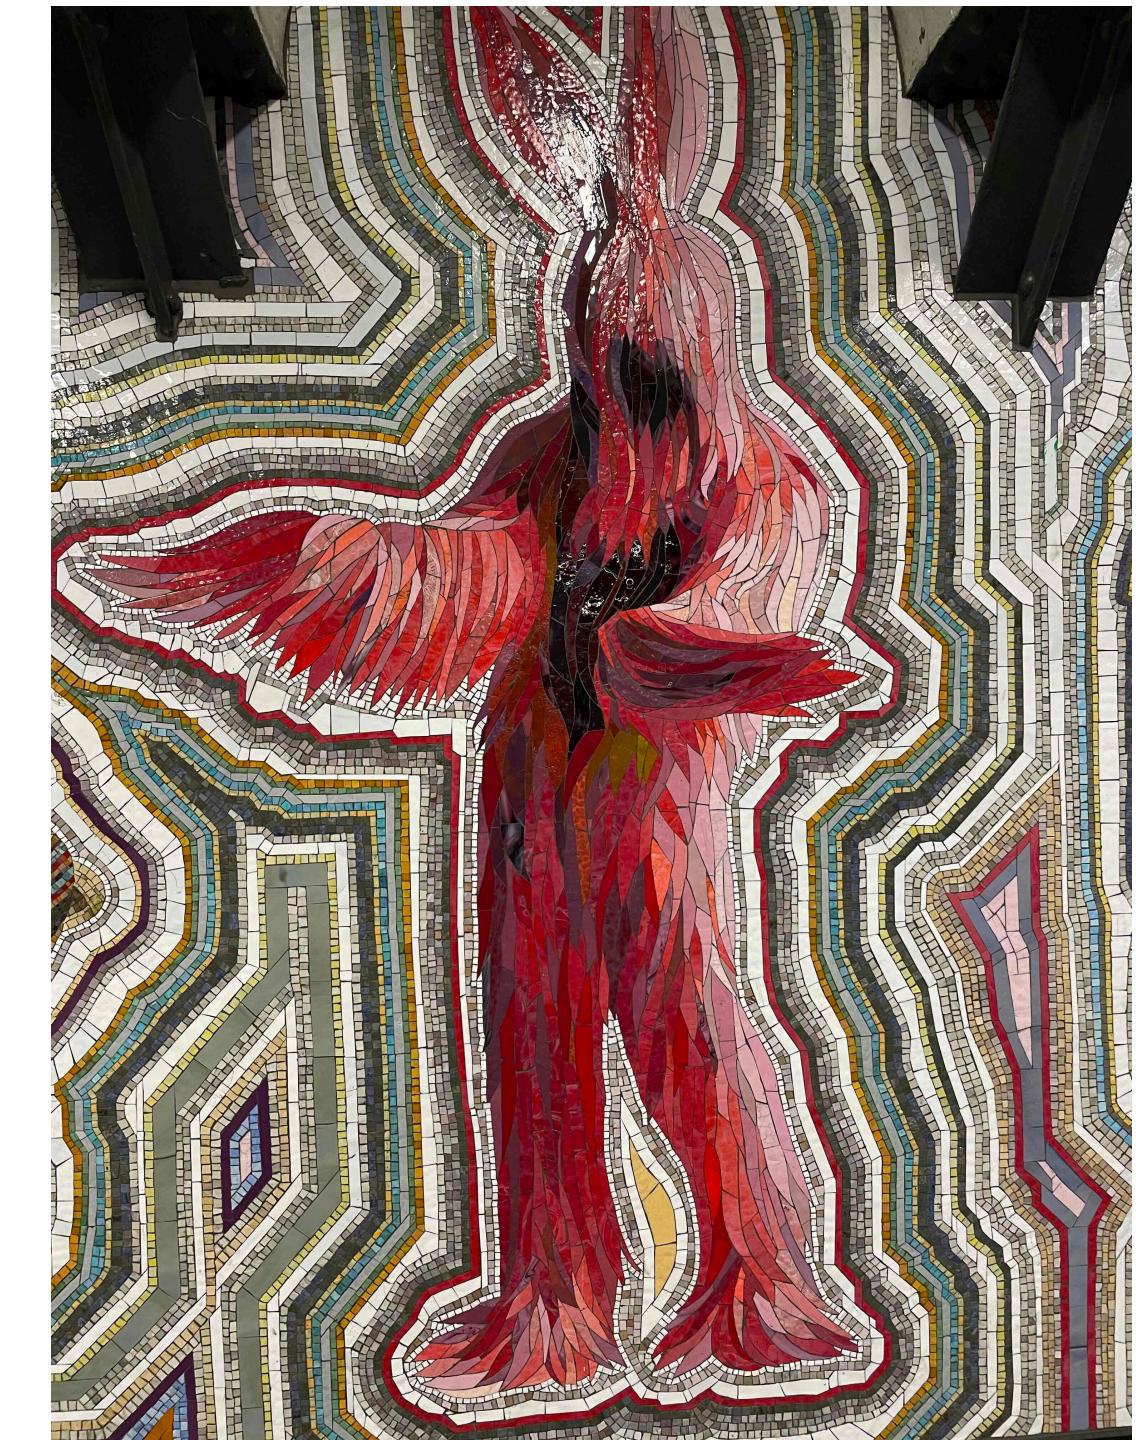
Maximum
12MiB



High
5.4MiB

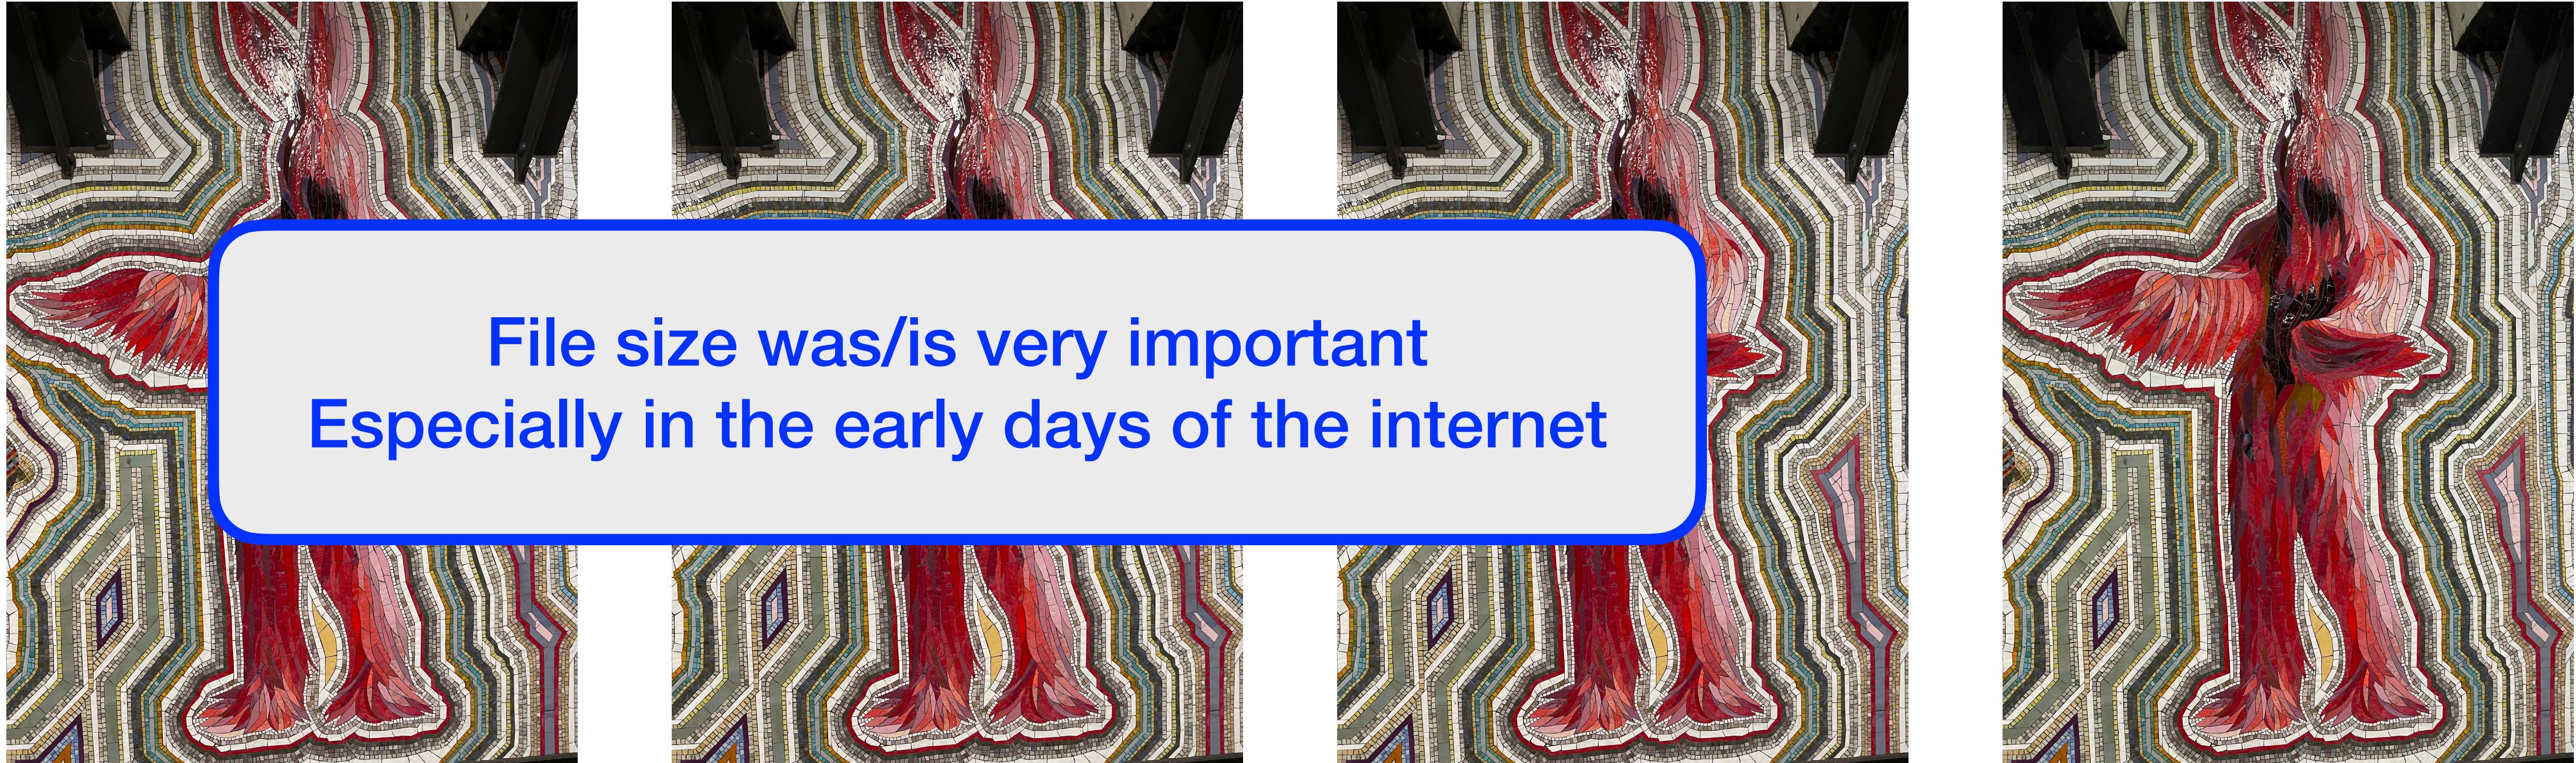


Medium
3.2MiB



Low
1.9MiB

JPEG Variable Compression



File size was/is very important
Especially in the early days of the internet

Maximum
12MiB

High
5.4MiB

Medium
3.2MiB

Low
1.9MiB

Times Square—42nd St



Each One, Every One, Equal All – Nick Cave

JPEG Compression

How does it work?

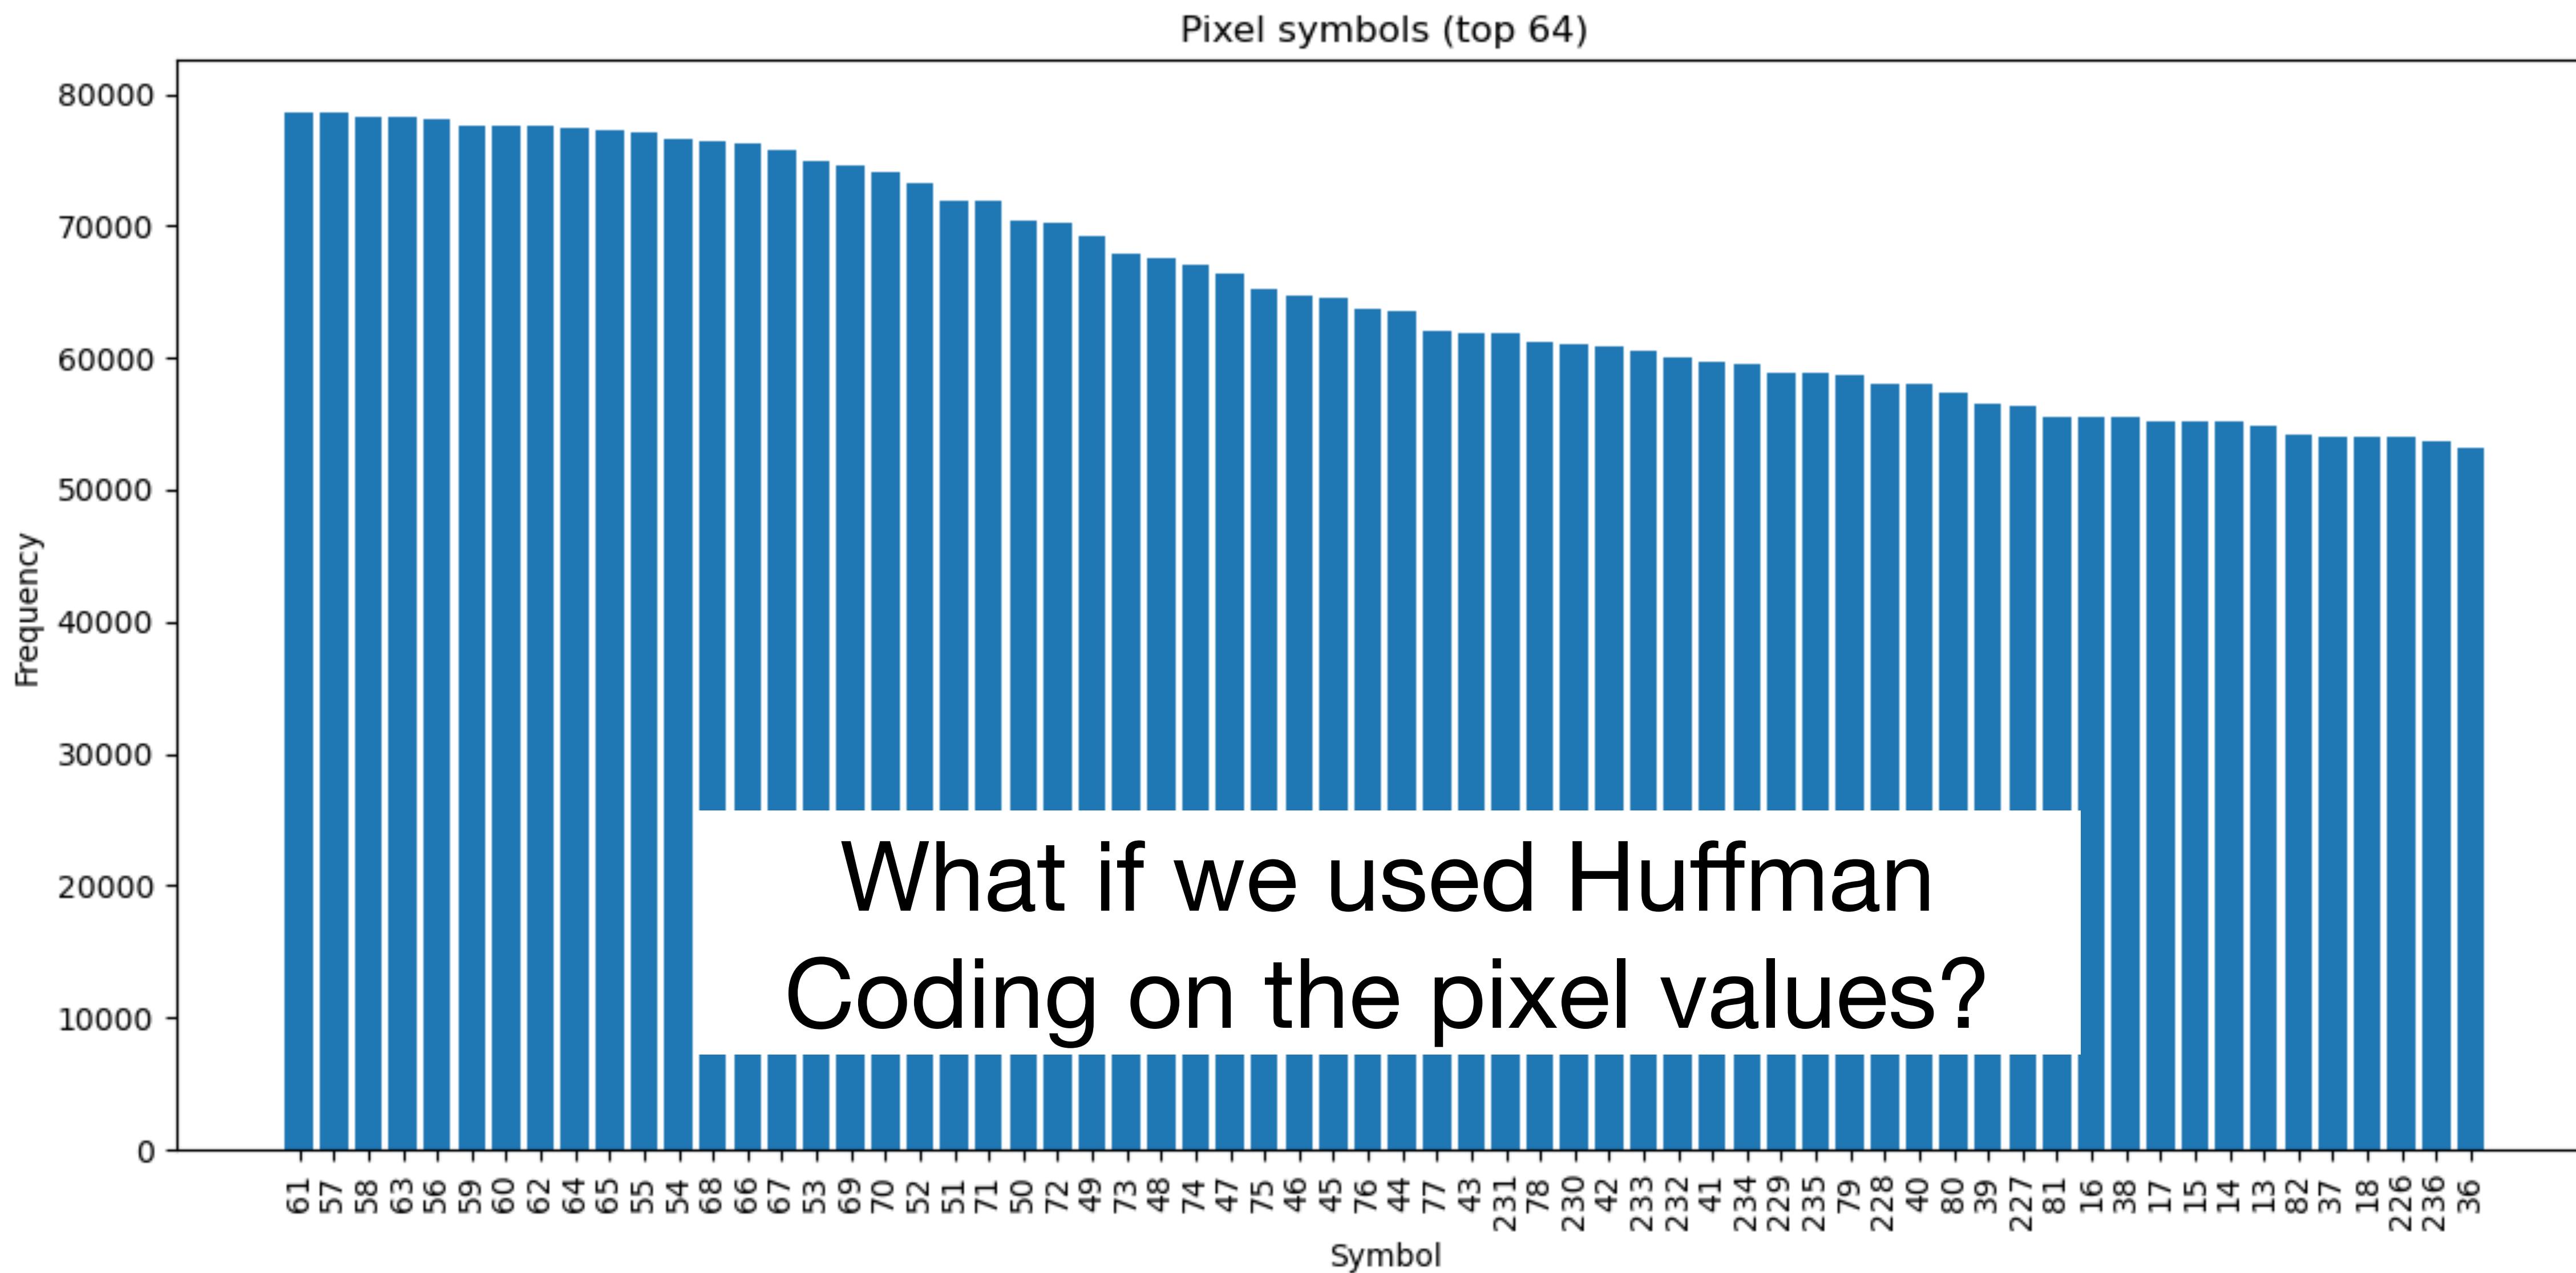
Why does it work so well?

We just learned about Huffman Coding...

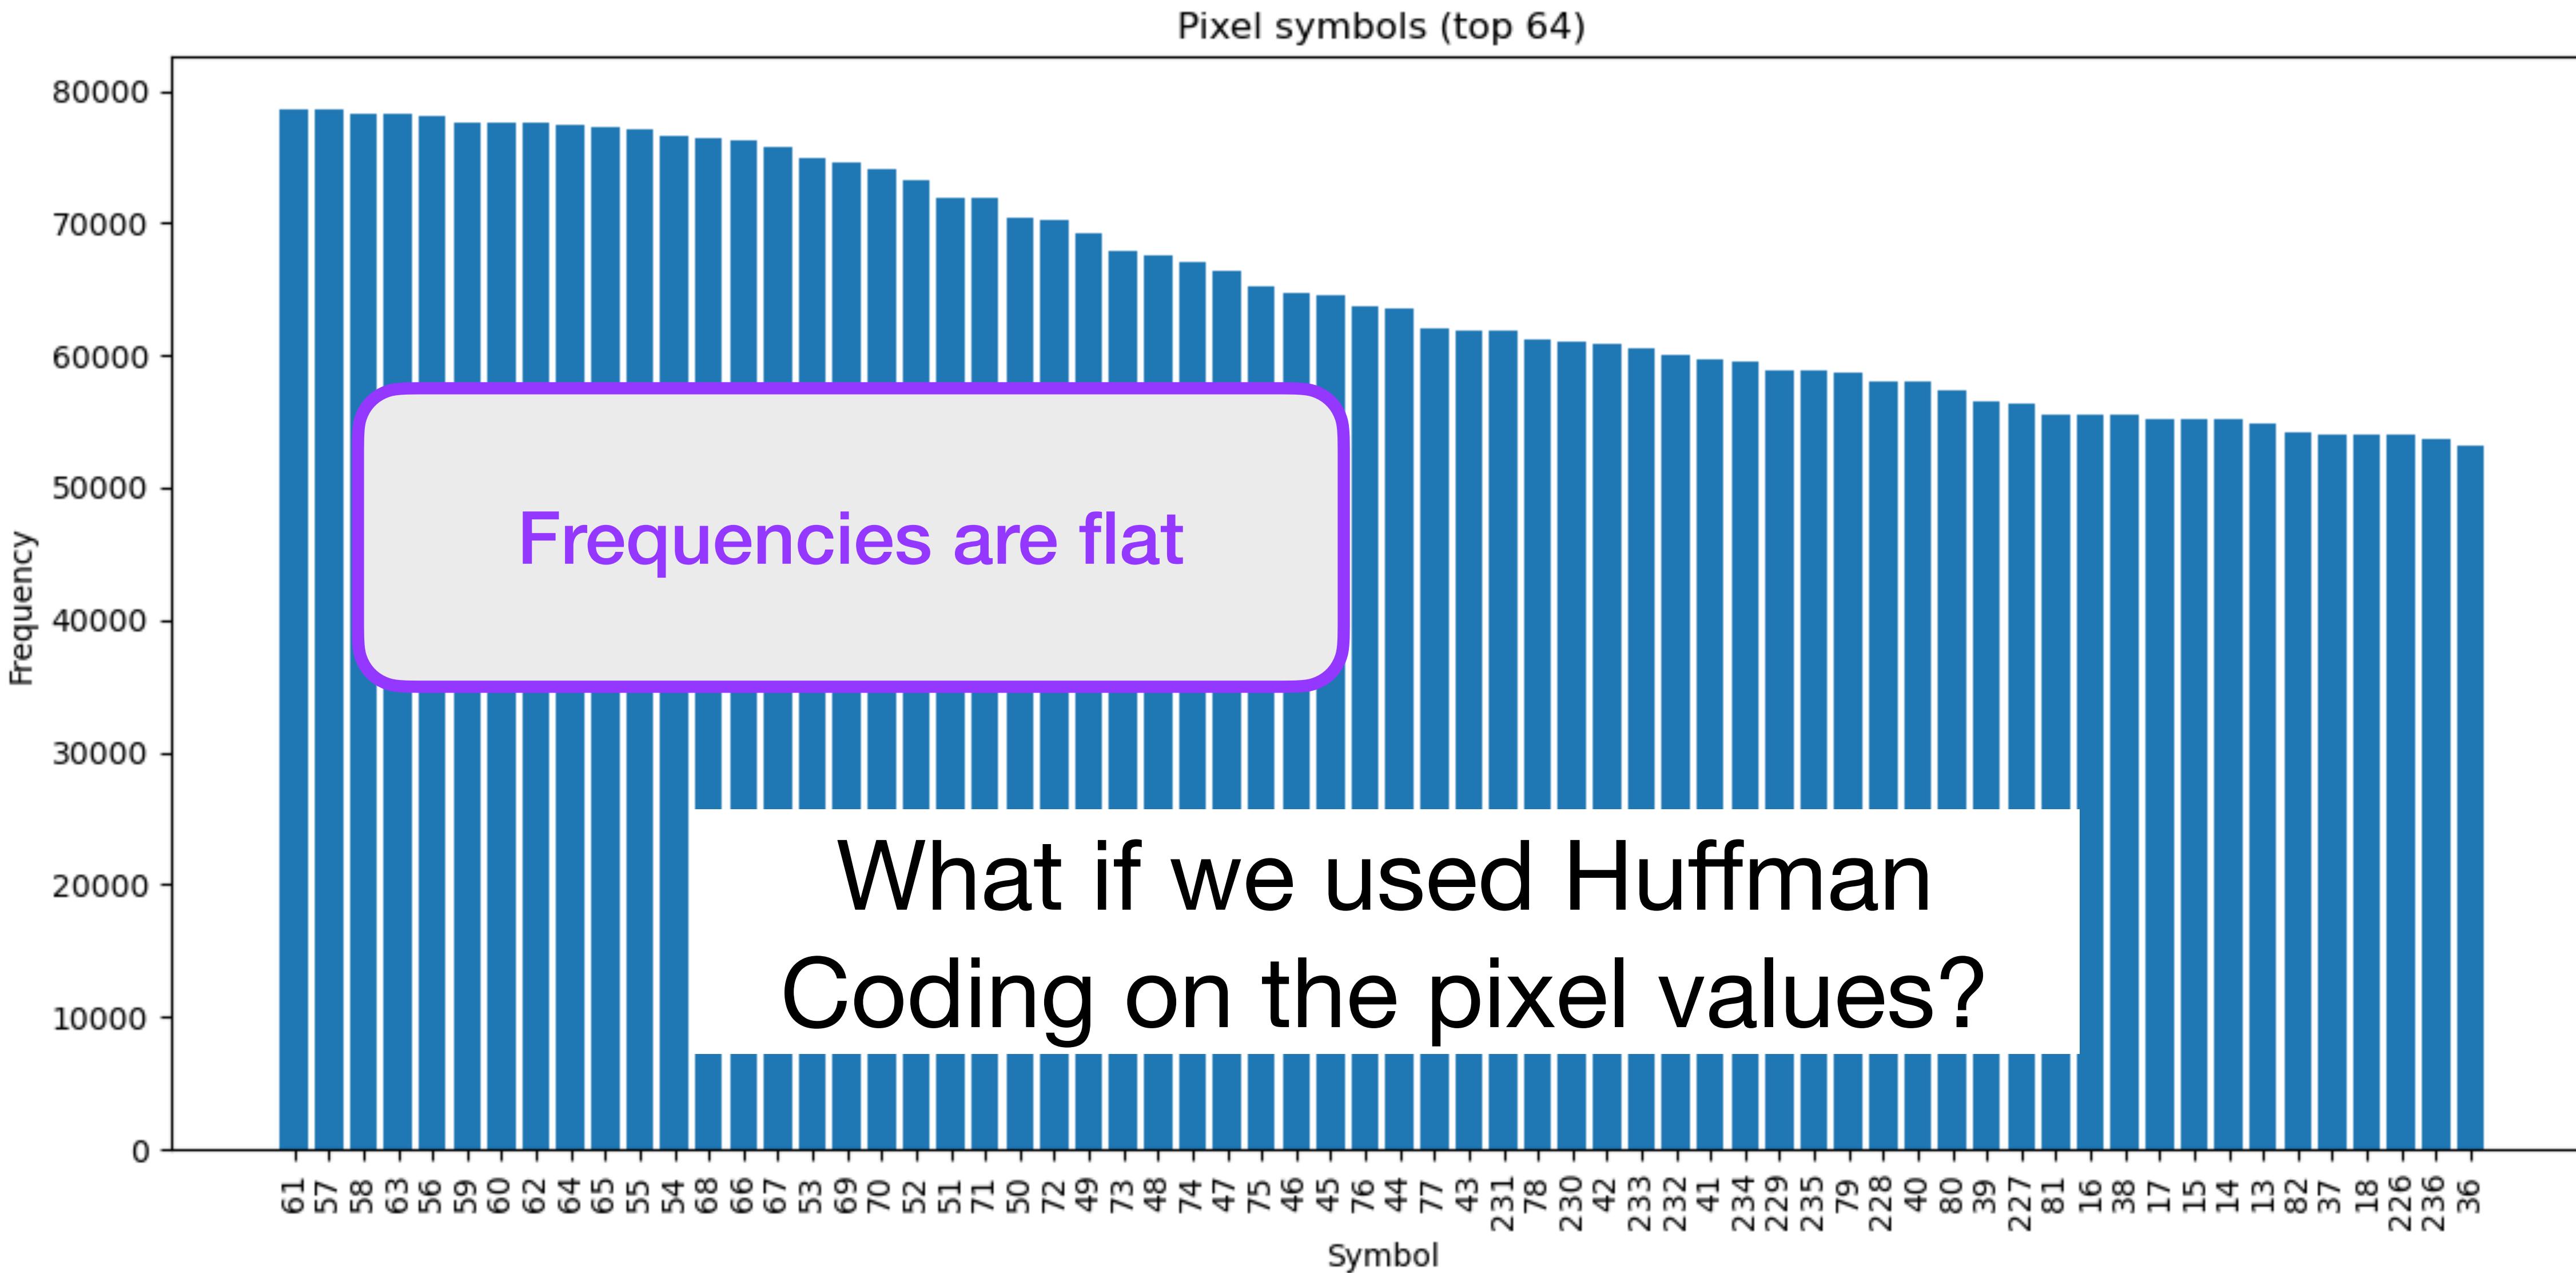


What if we used Huffman
Coding on the pixel values?

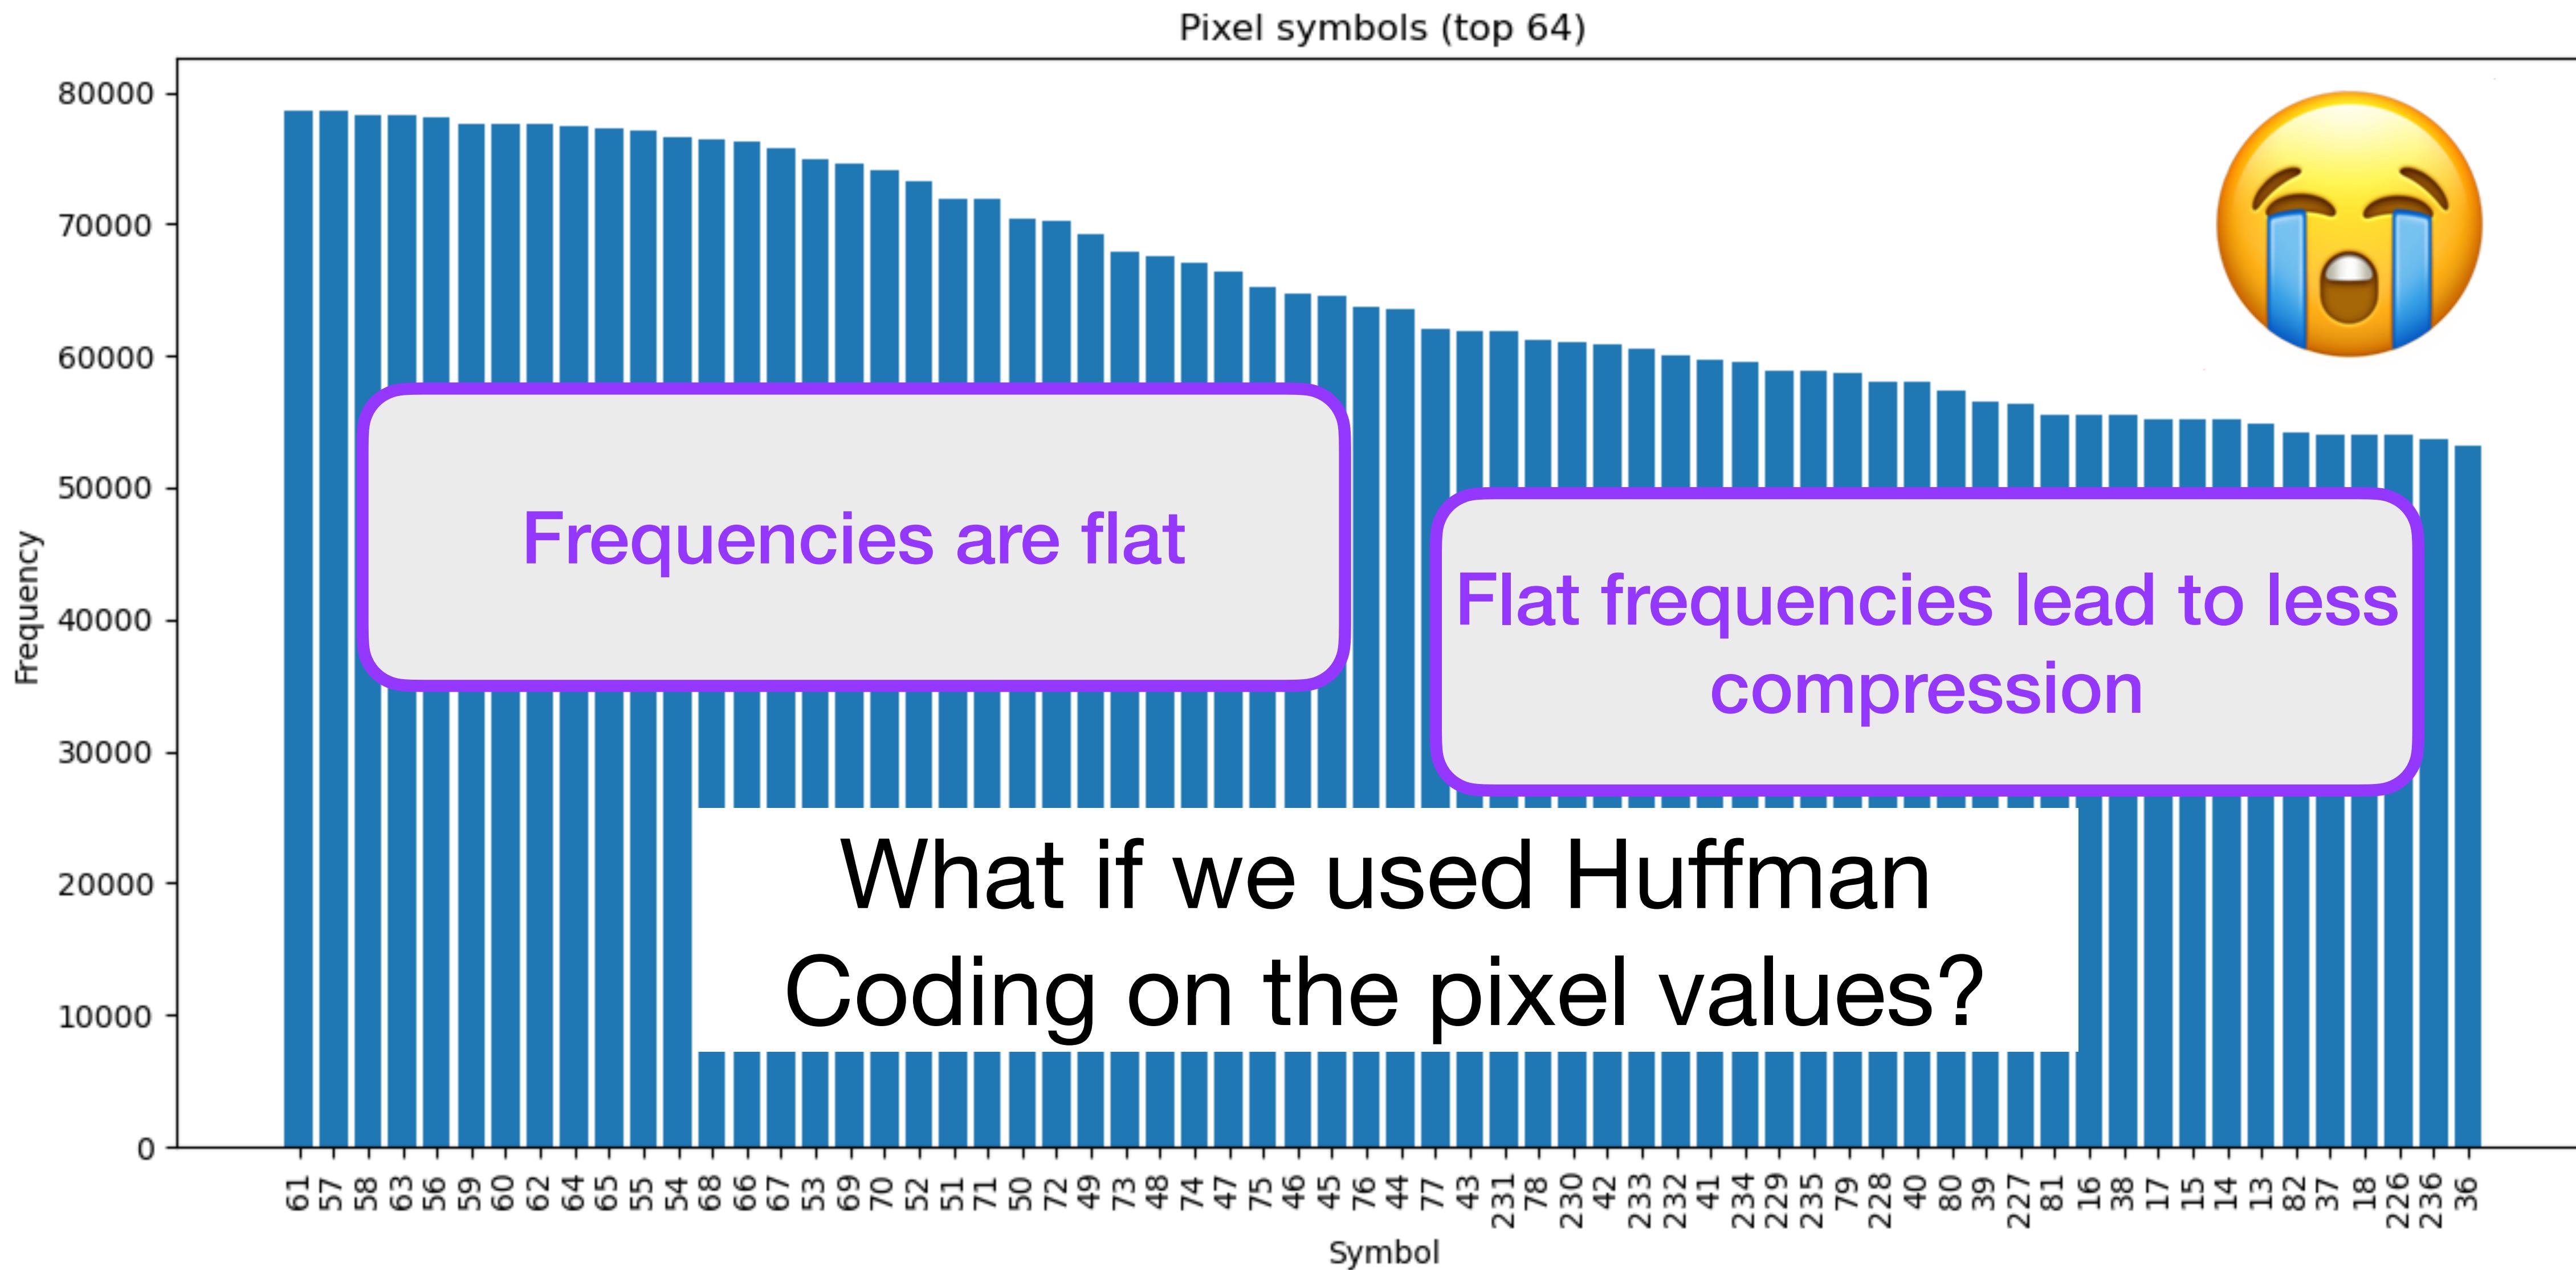
We just learned about Huffman Coding...



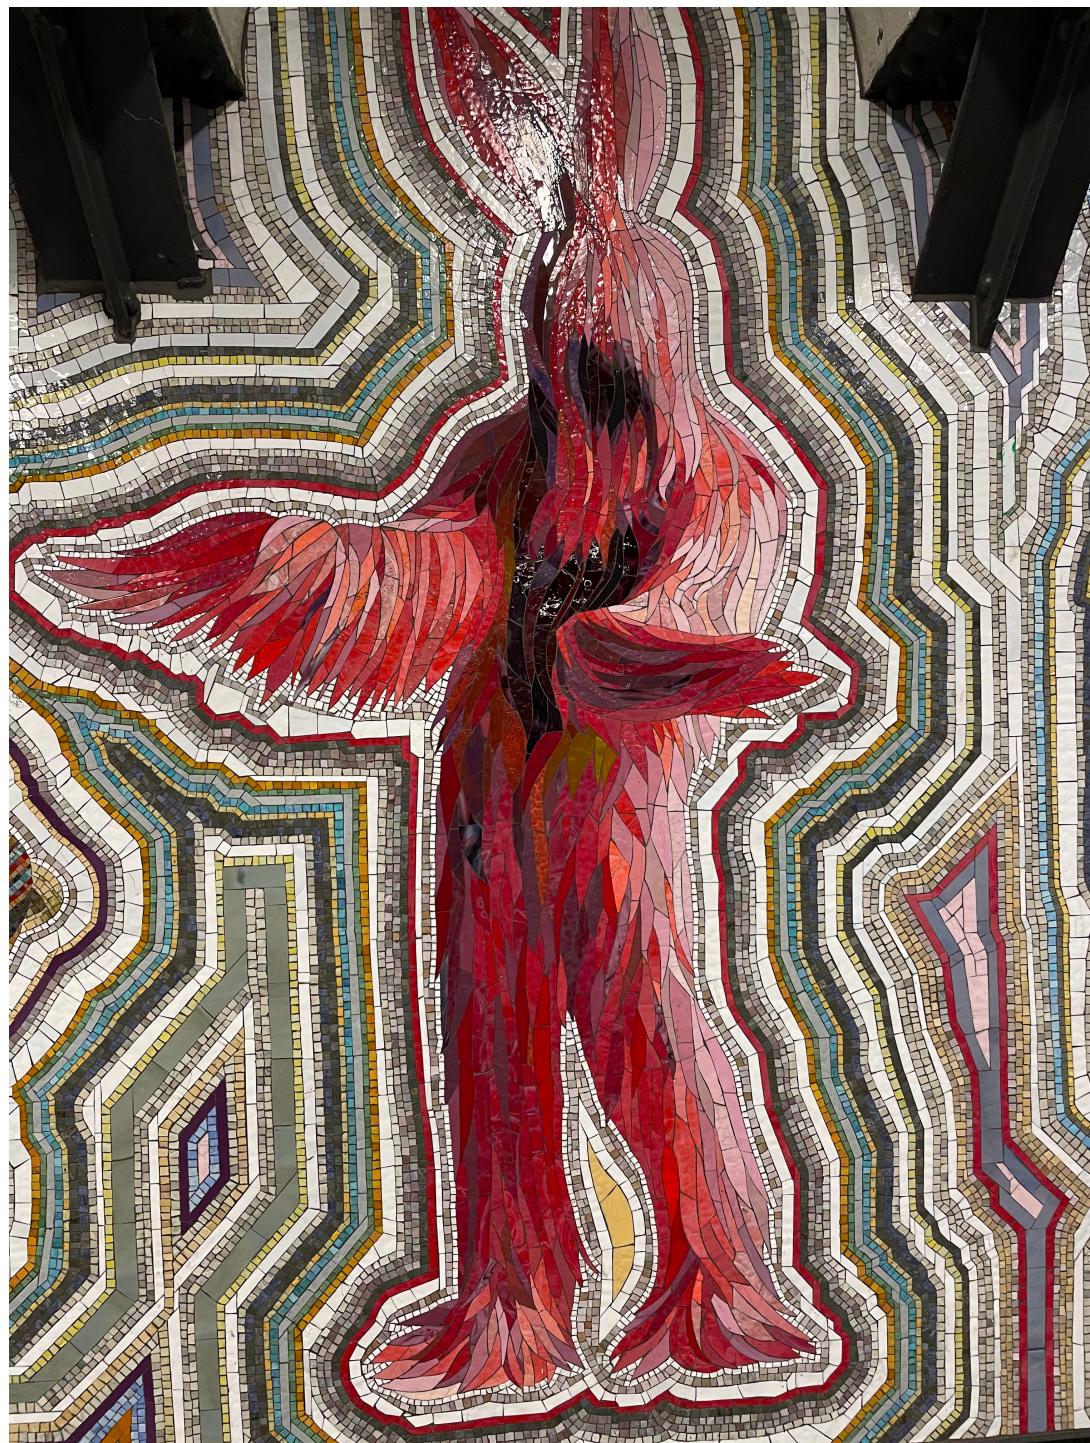
We just learned about Huffman Coding...



We just learned about Huffman Coding...



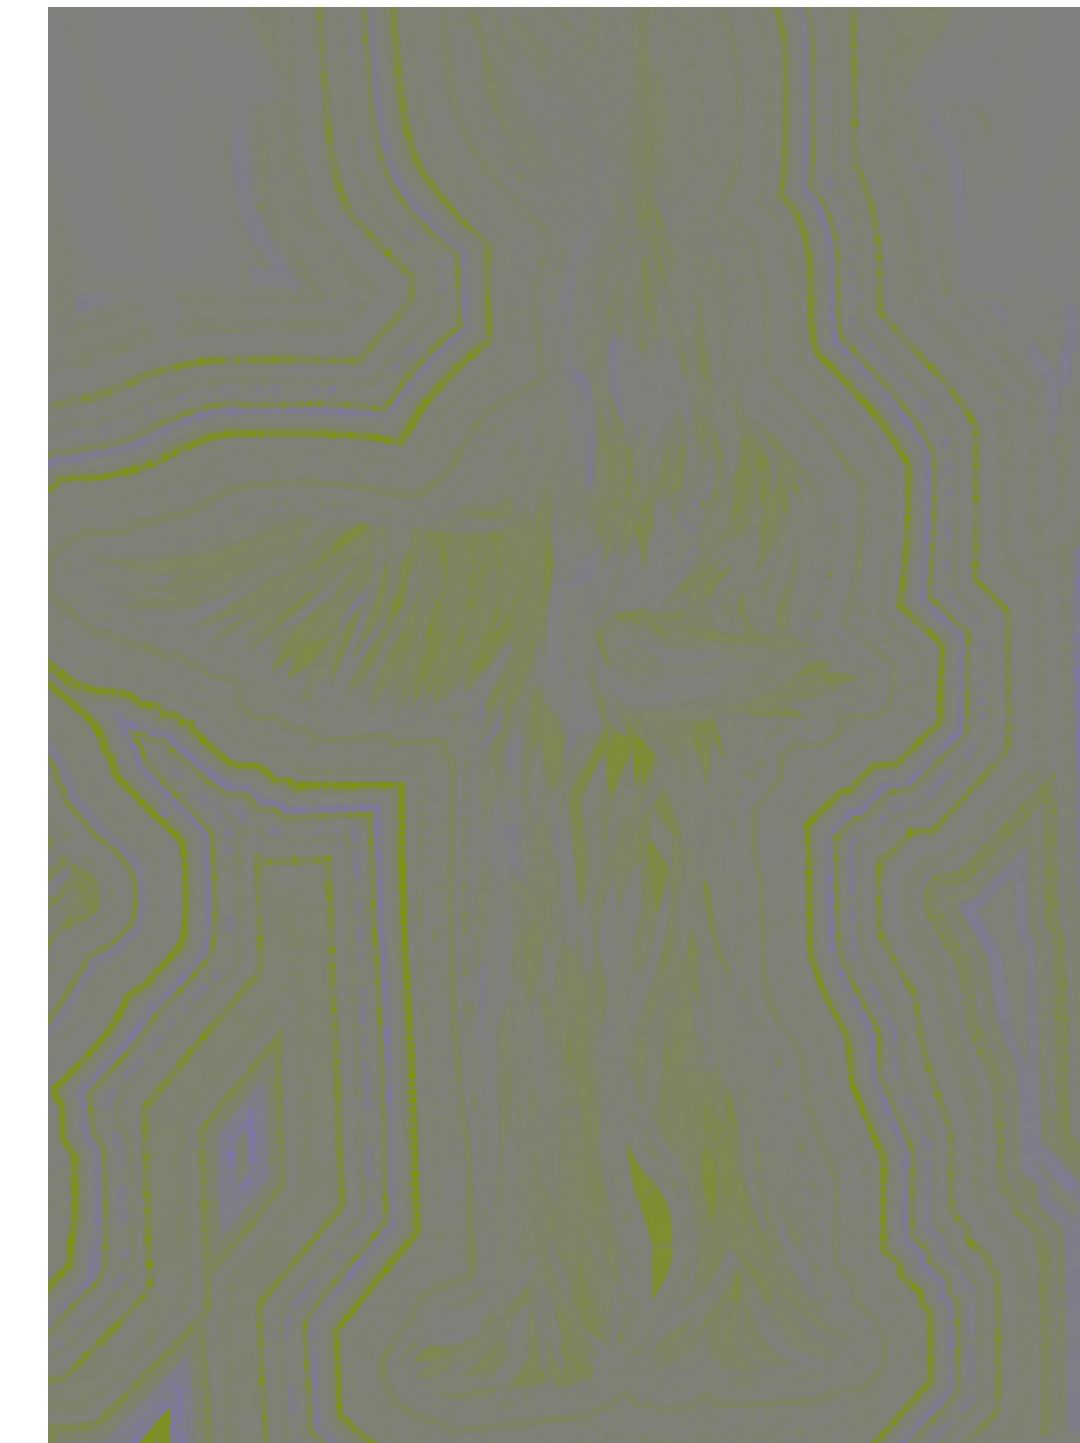
Start by separating channels



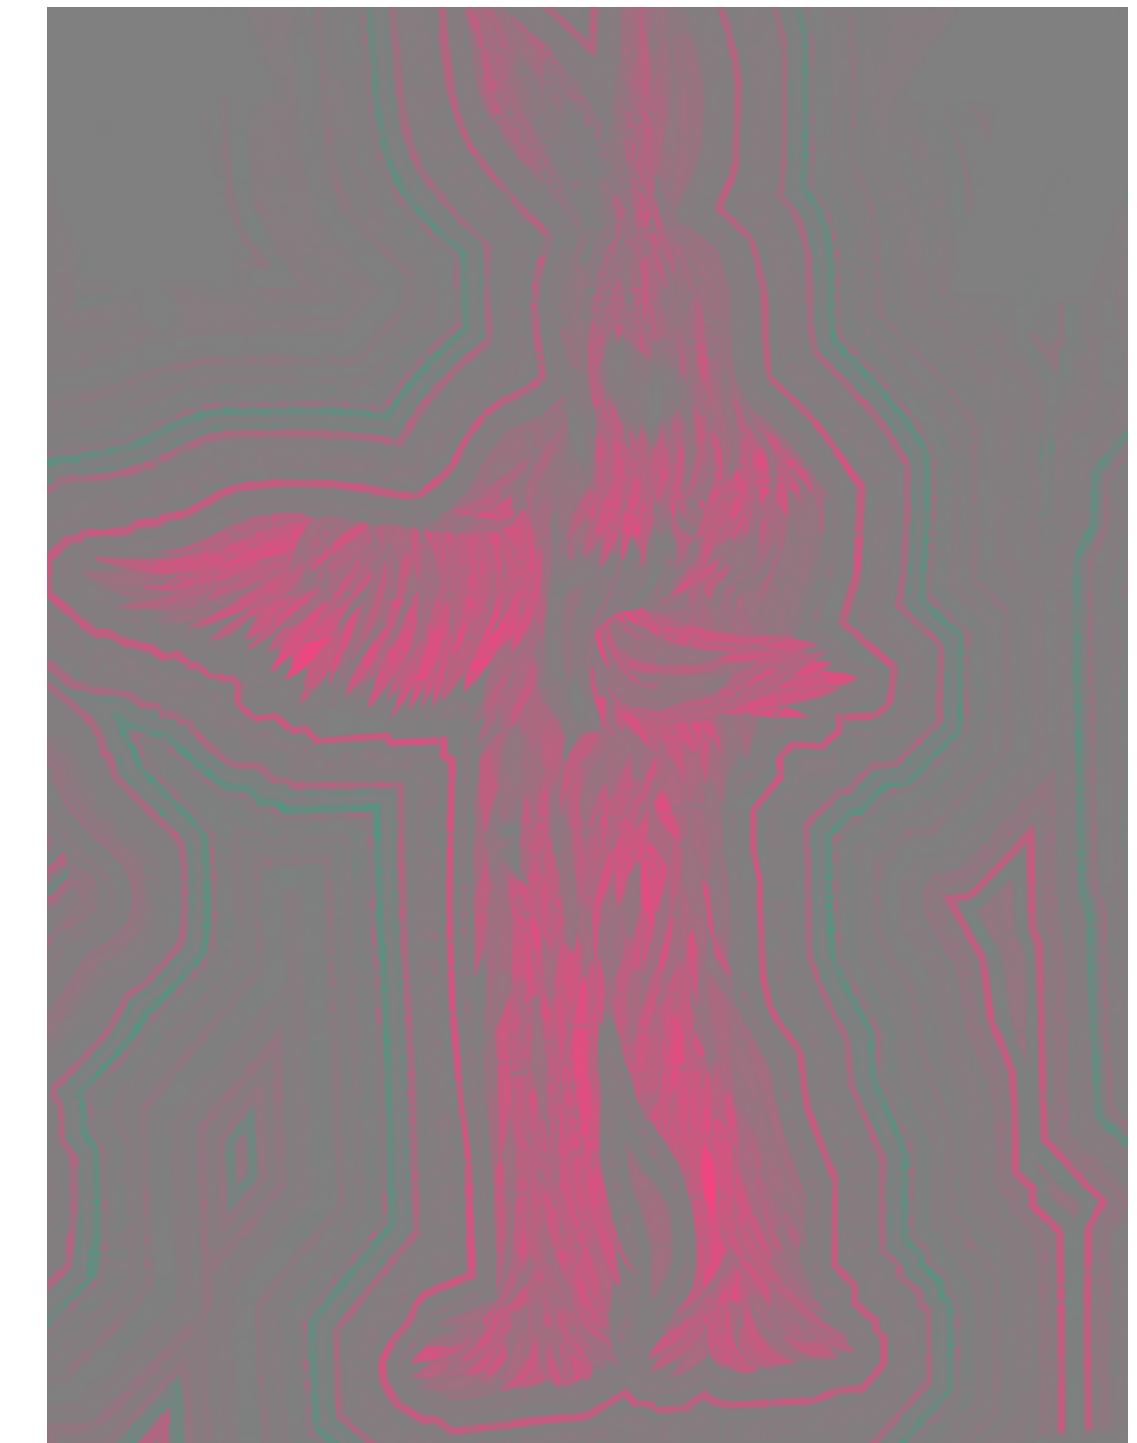
Combined



Y
Luminance



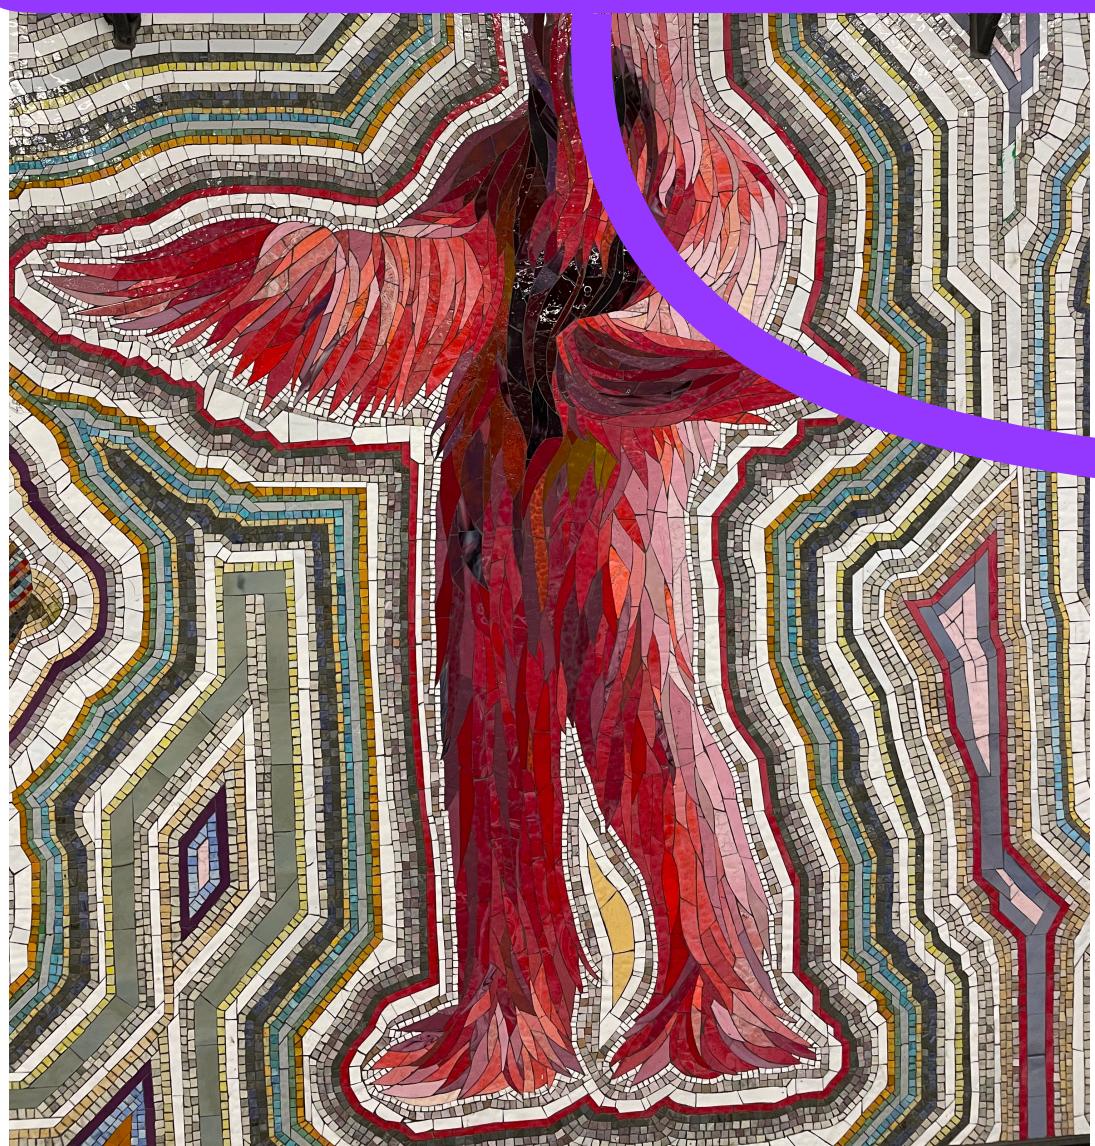
Cb
Blue difference



Cr
Red difference

Start by separating channels

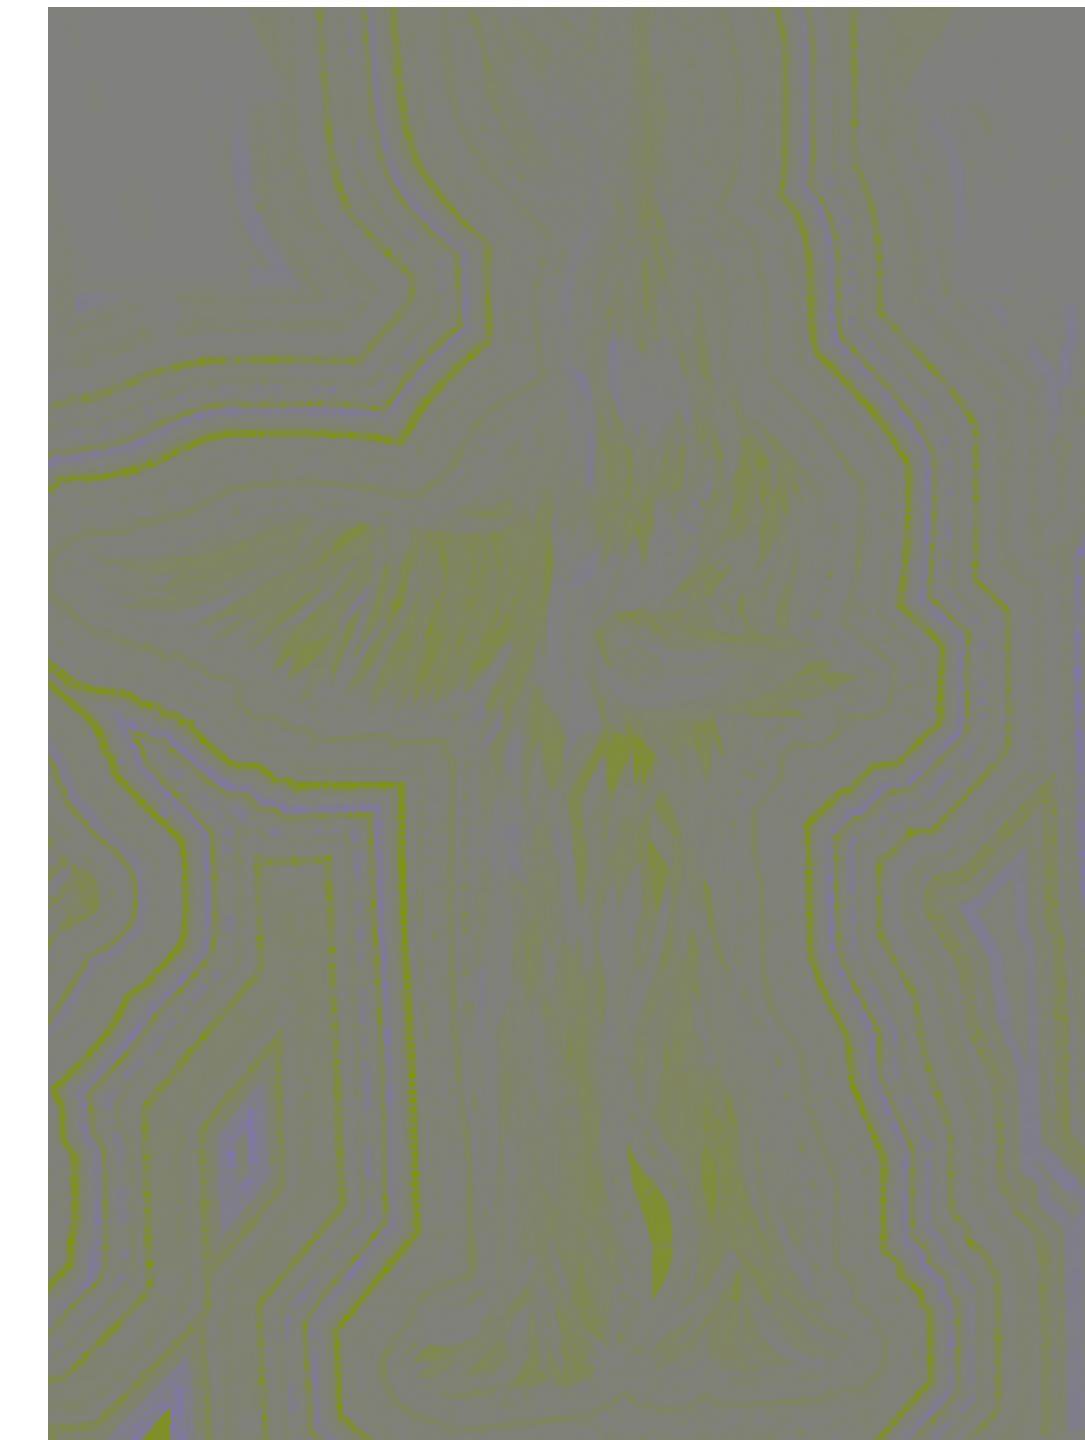
Most detail in Y



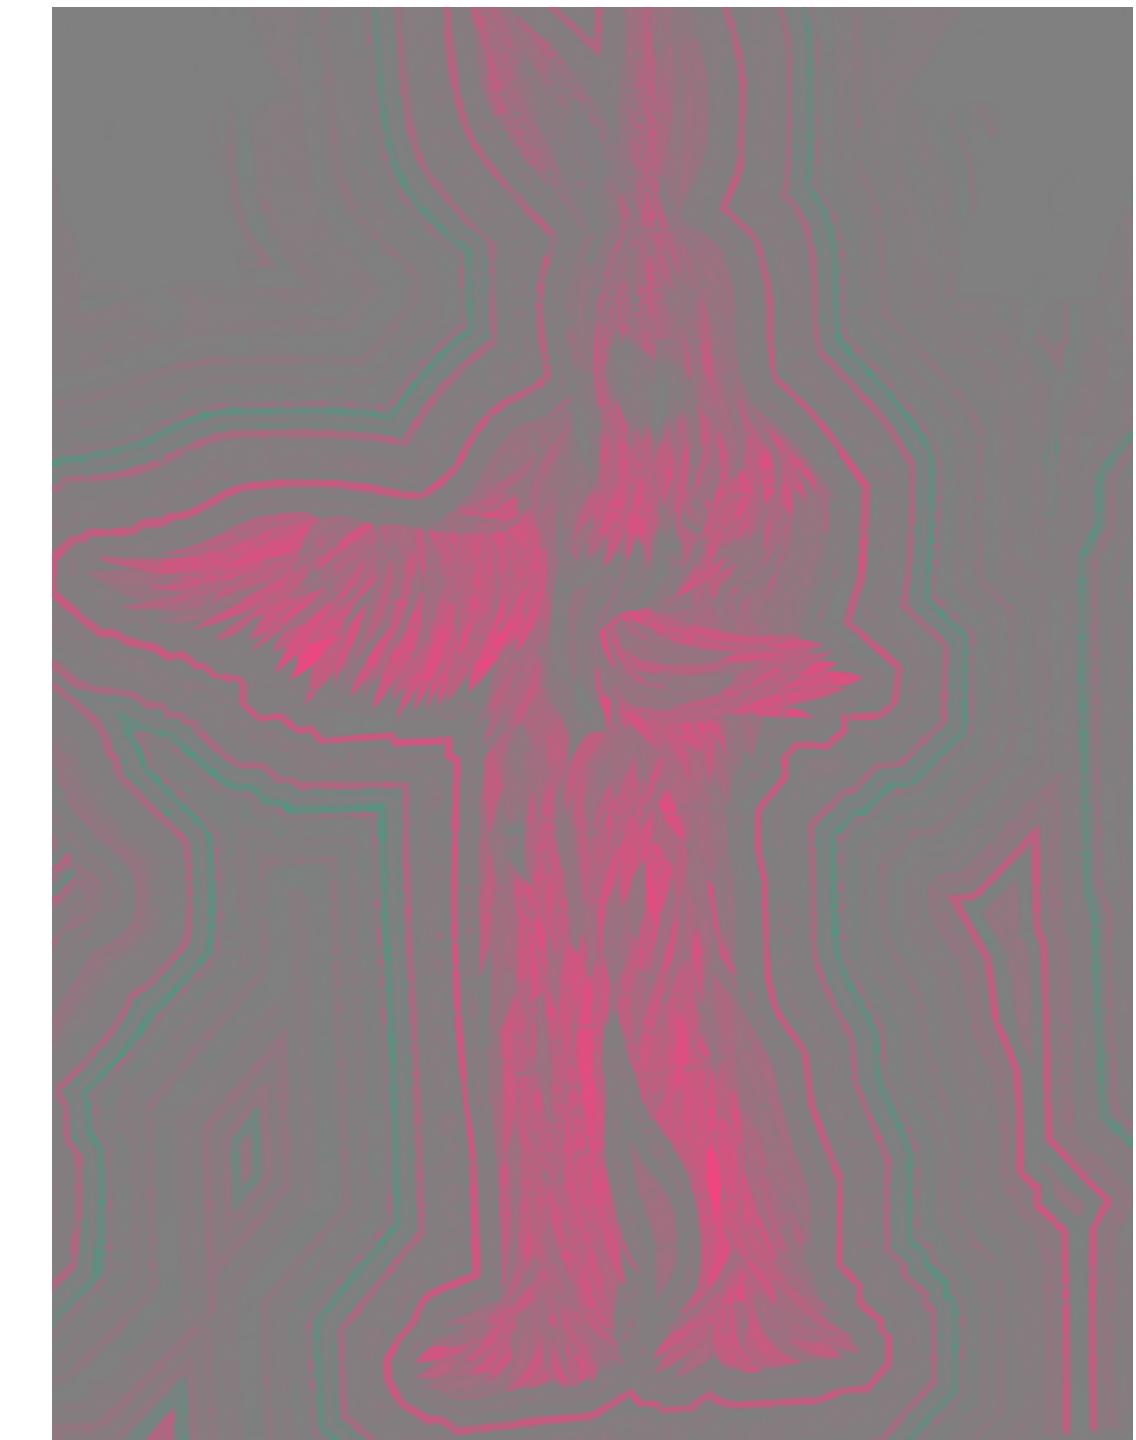
Combined



Y
Luminance



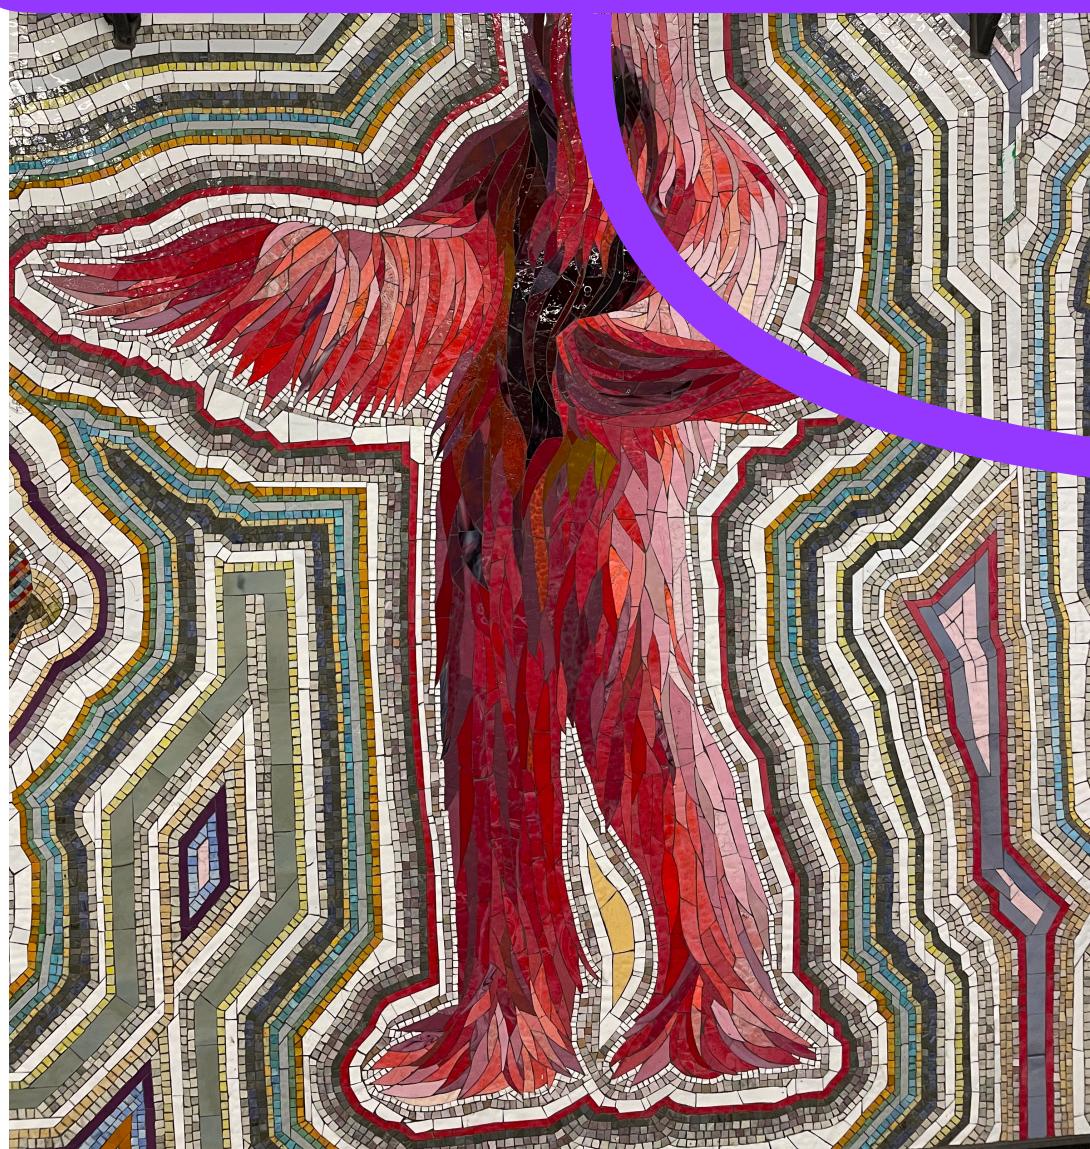
Cb
Blue difference



Cr
Red difference

Start by separating channels

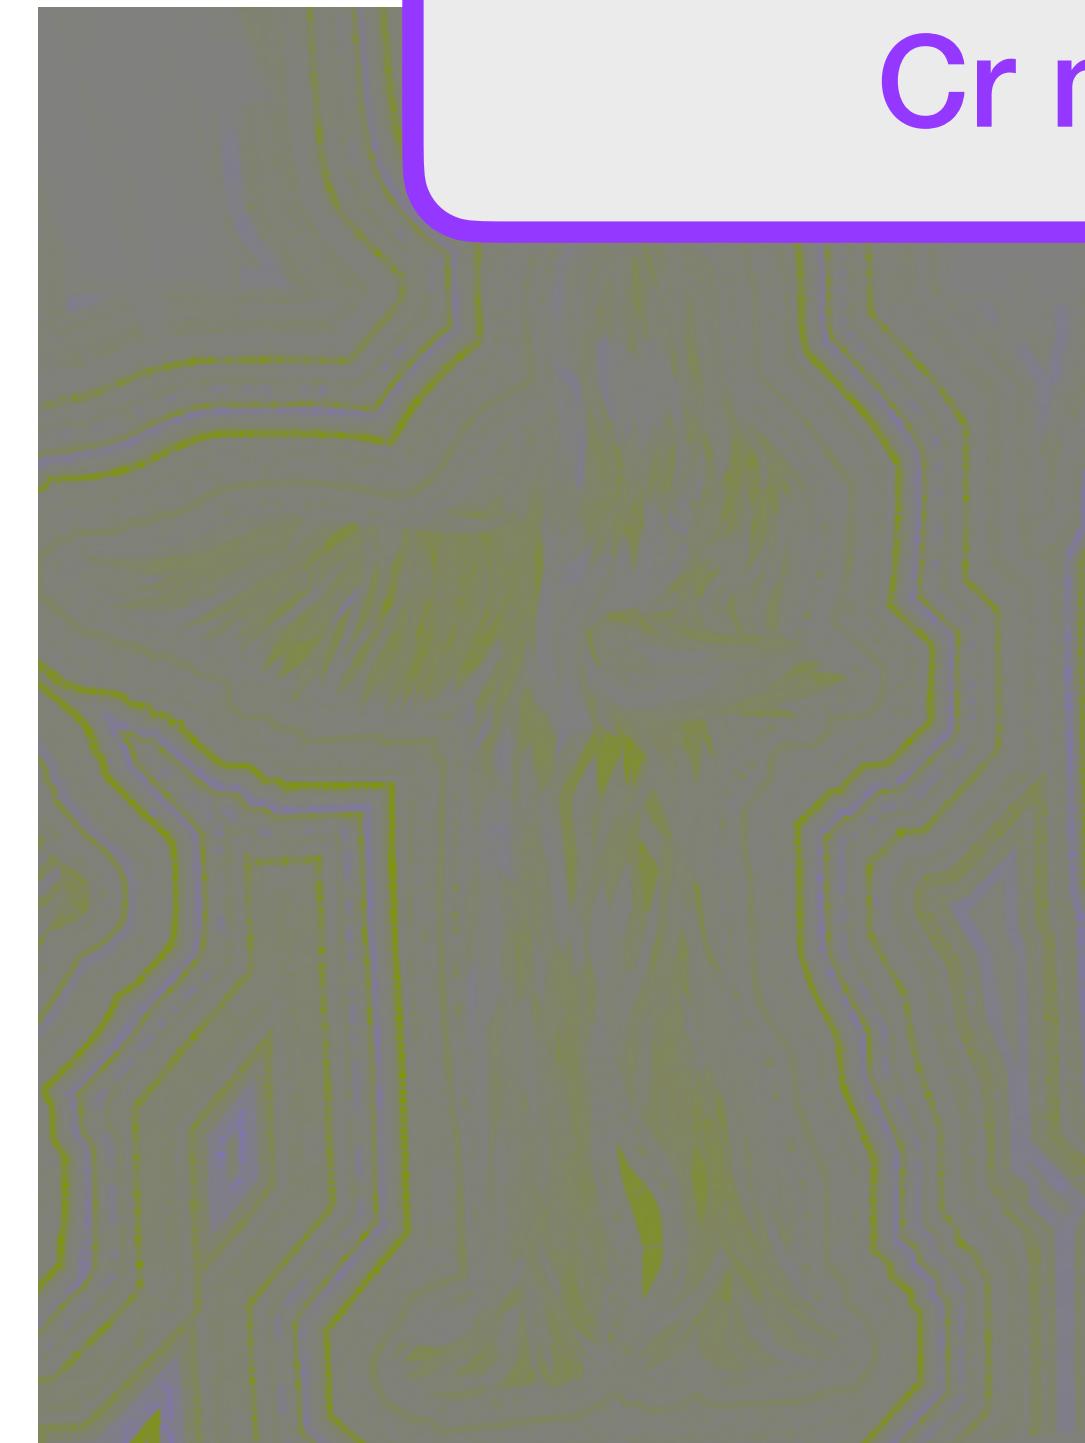
Most detail in Y



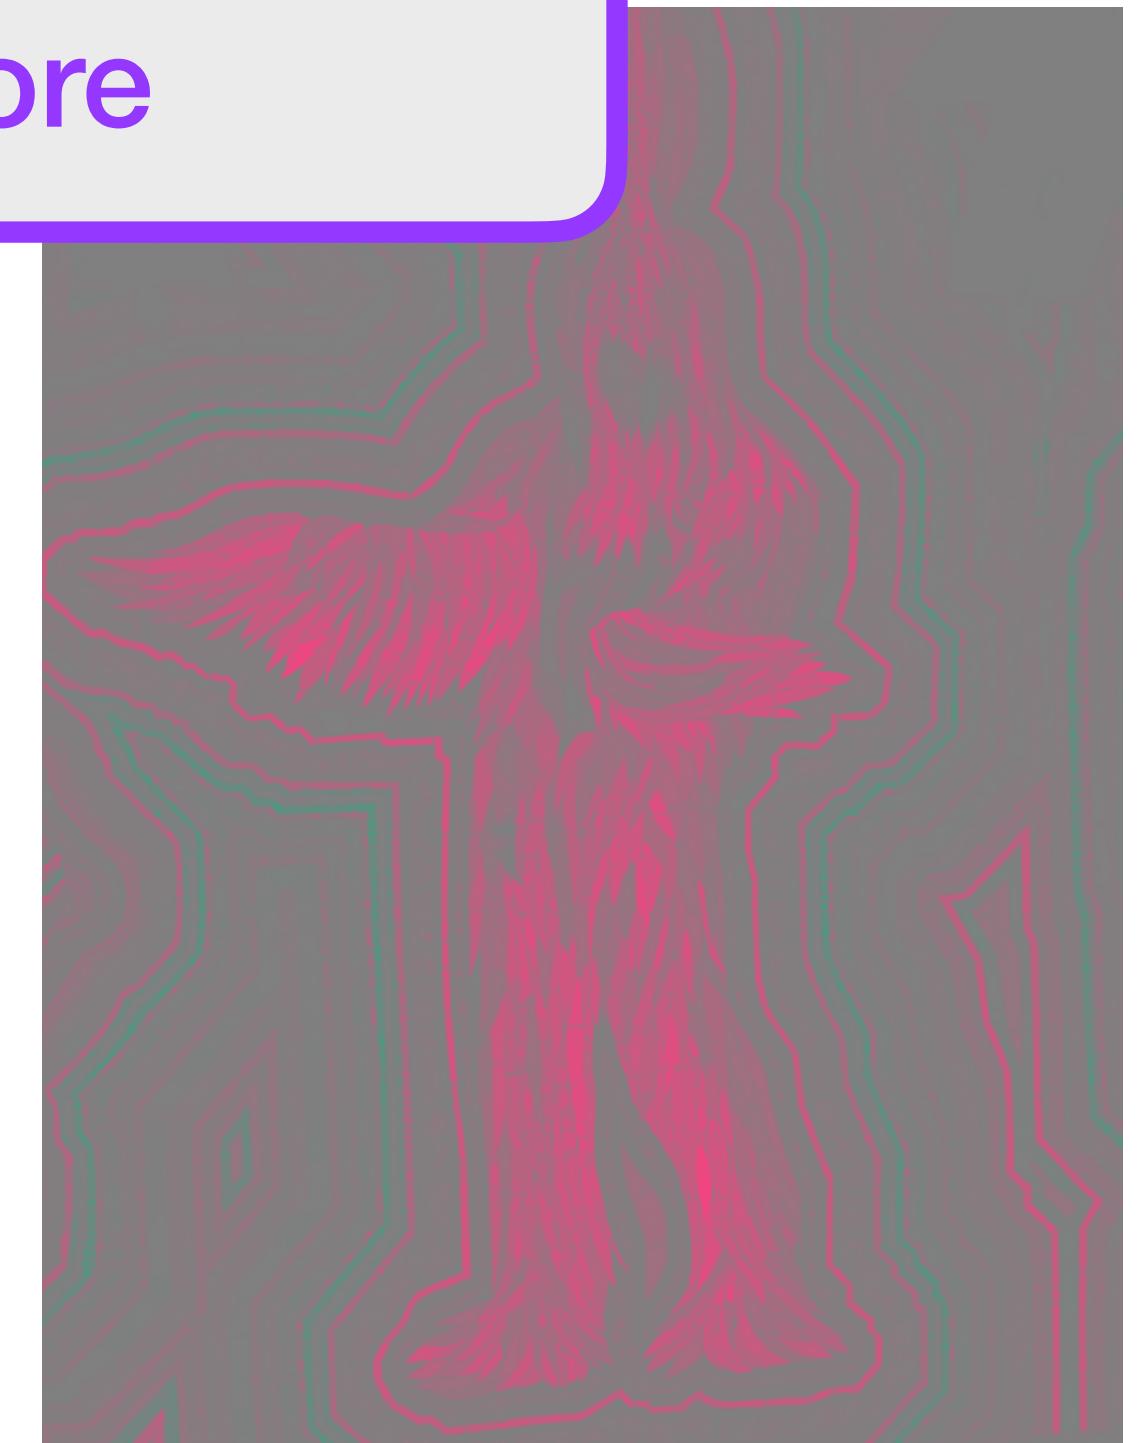
Combined

Y
Luminance

Can compress Cb and
Cr more



Cb
Blue difference



Cr
Red difference

Human Perception of Color



We are much more
sensitive to luminance
boundaries than color
boundaries

Focus on Y – Luminance

Even though this image
has lots of detail



Focus on Y – Luminance

Large areas look
“uniform”

Even though this image
has lots of detail



Focus on Y – Luminance

Large areas look
“uniform”



Even though this image
has lots of detail

When we zoom in, over
small blocks there is a
lot of uniformity



240 x 240

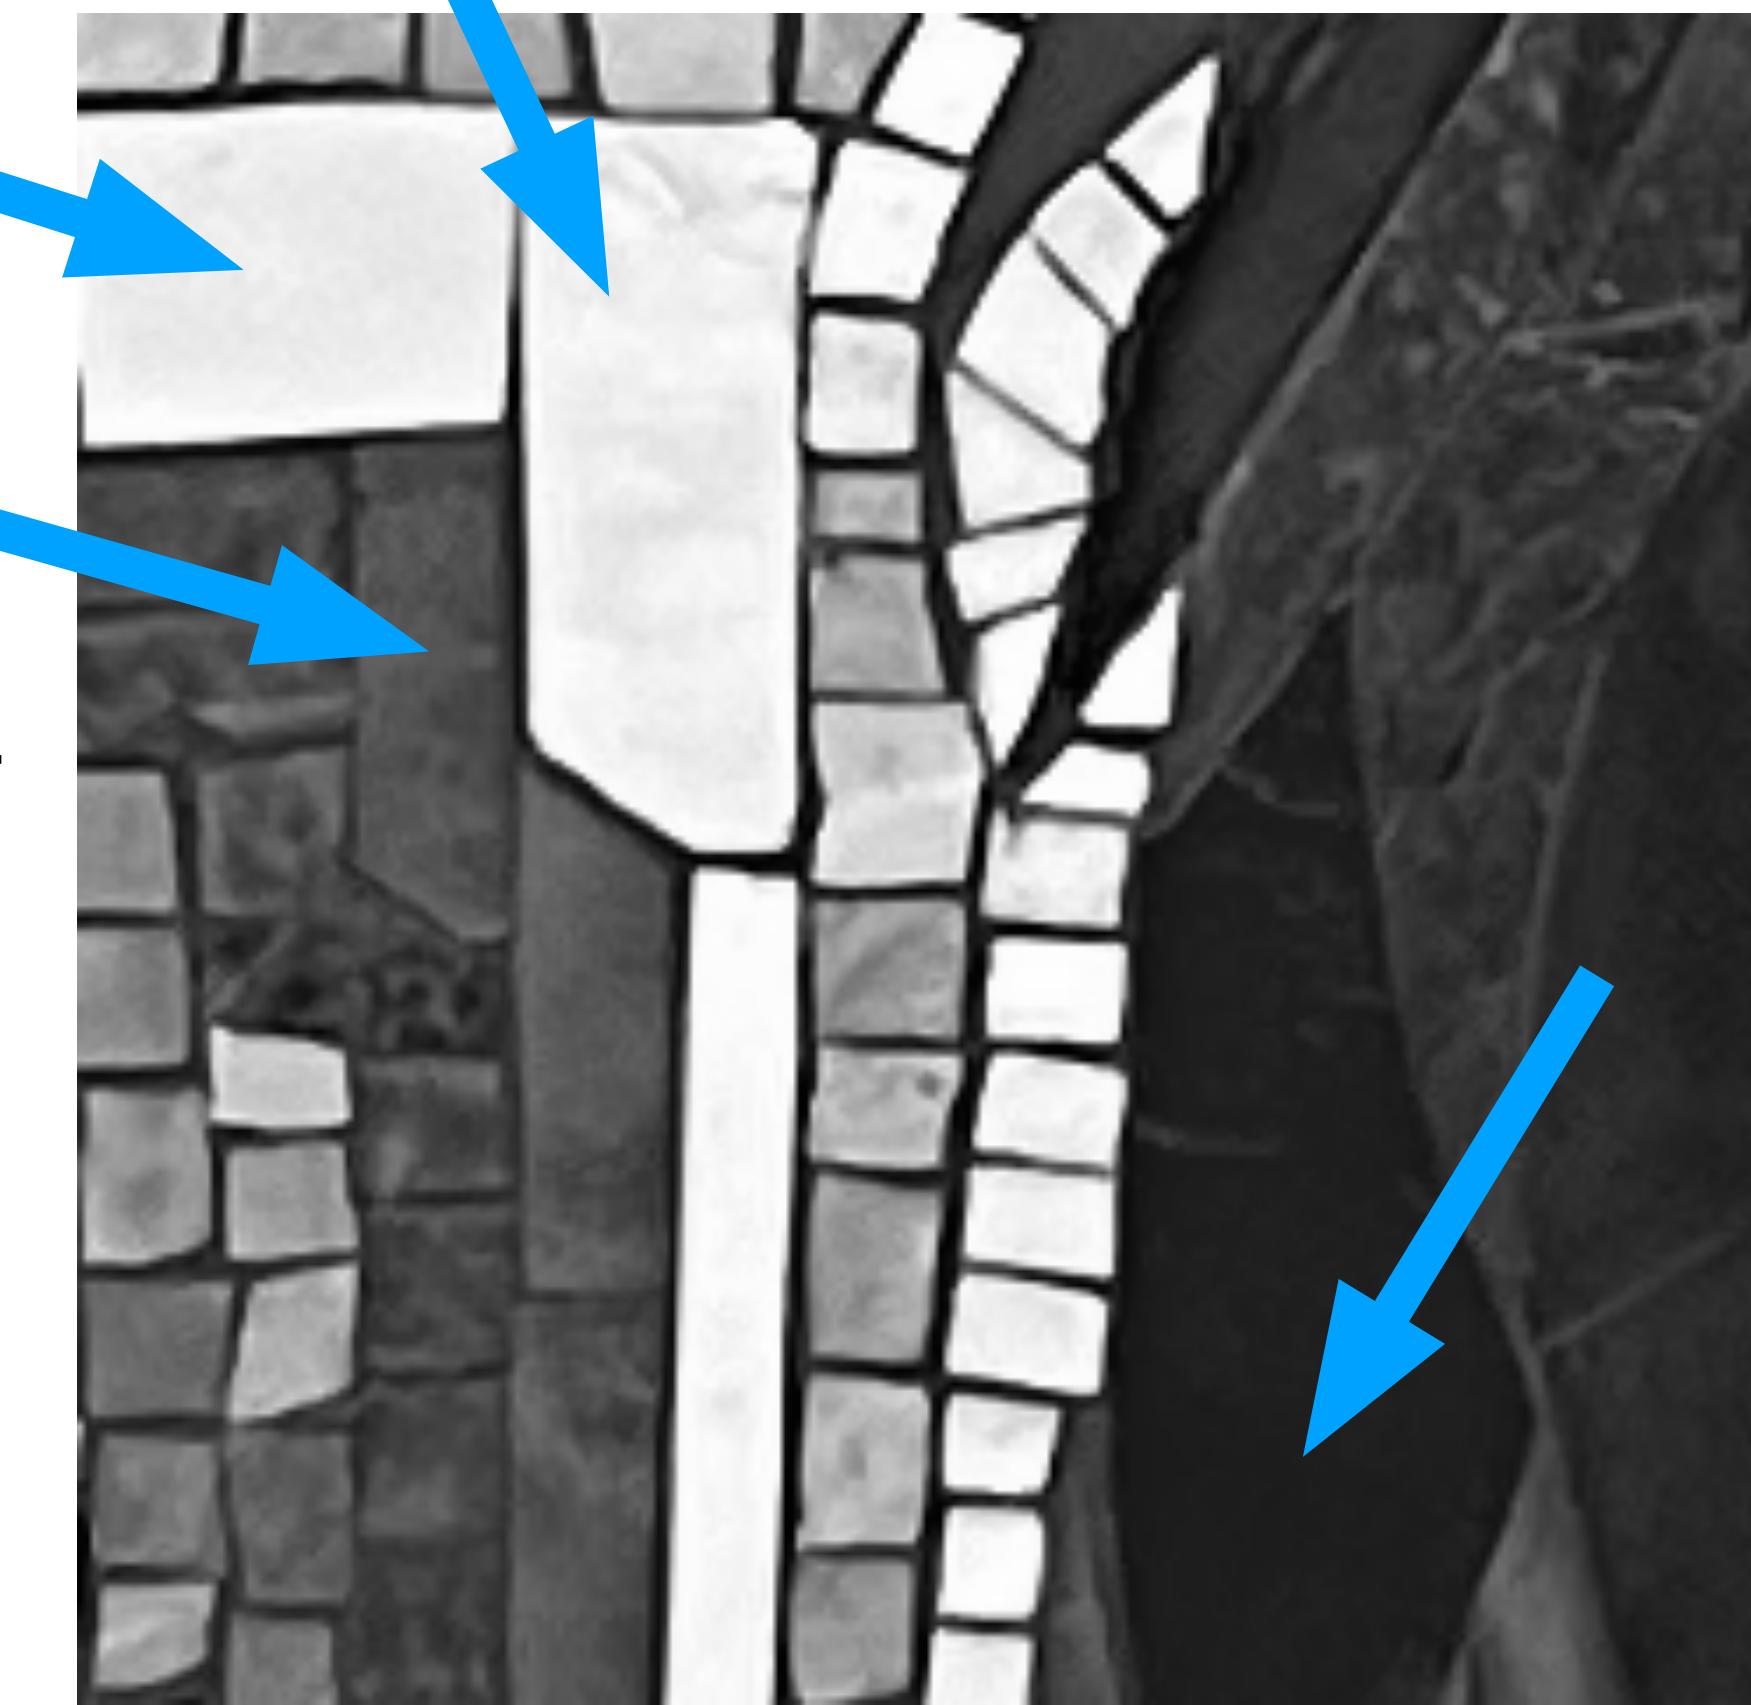
Focus on Y – Luminance

Large areas look
“uniform”



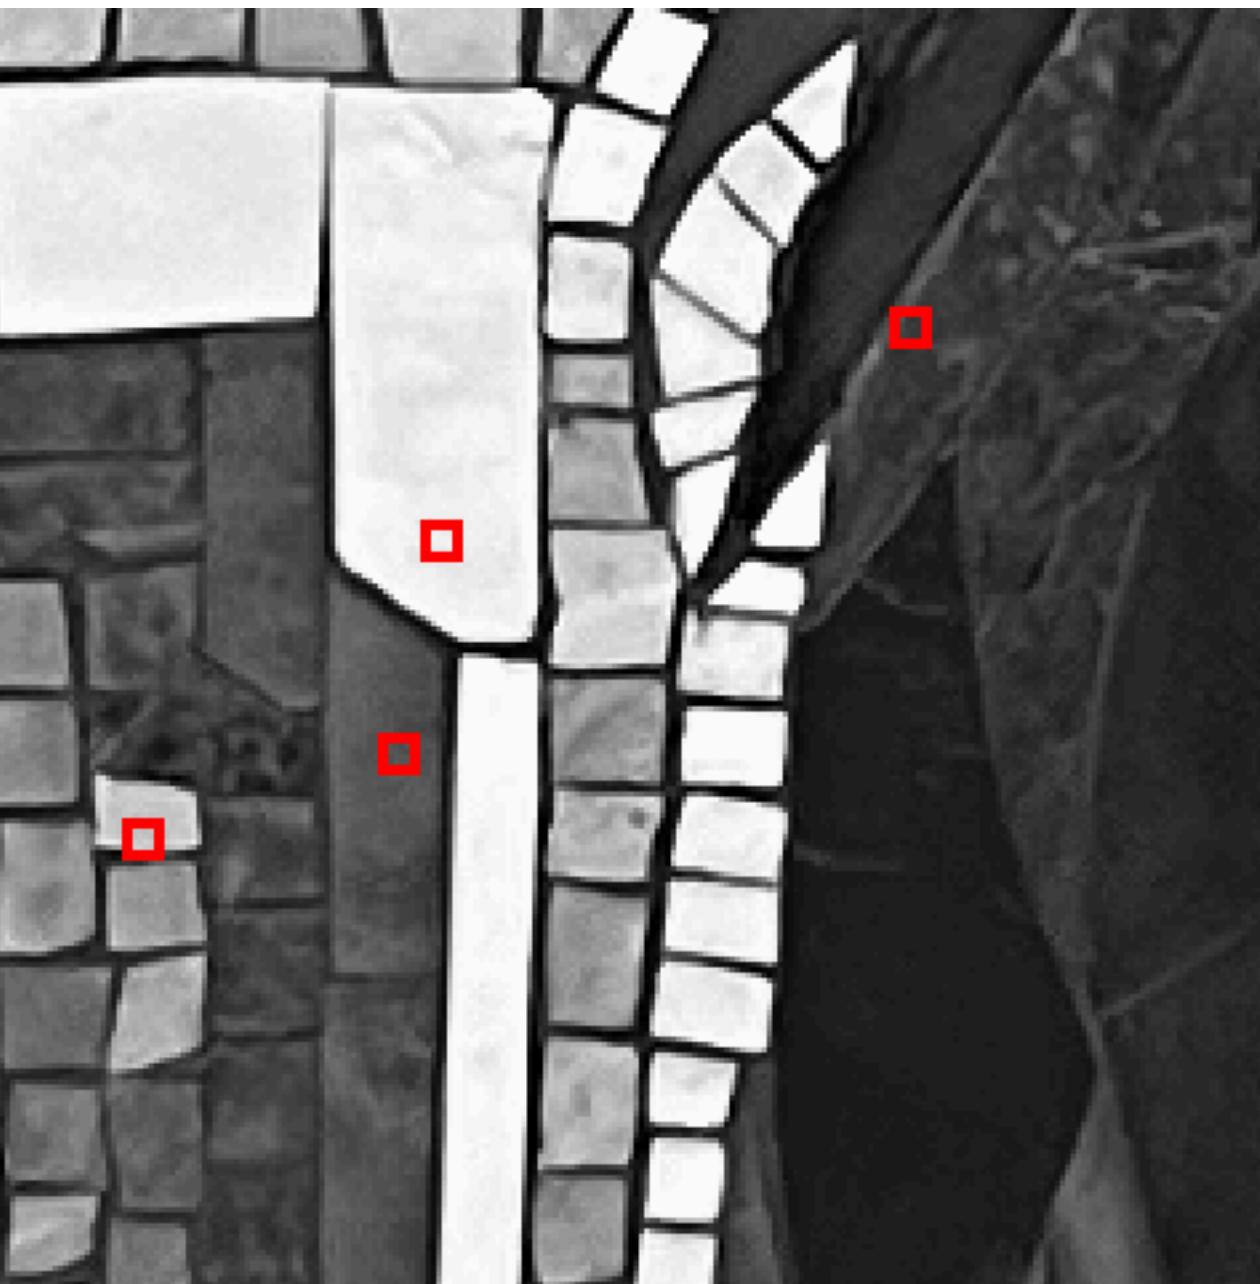
Even though this image
has lots of detail

When we zoom in, over
small blocks there is a
lot of uniformity



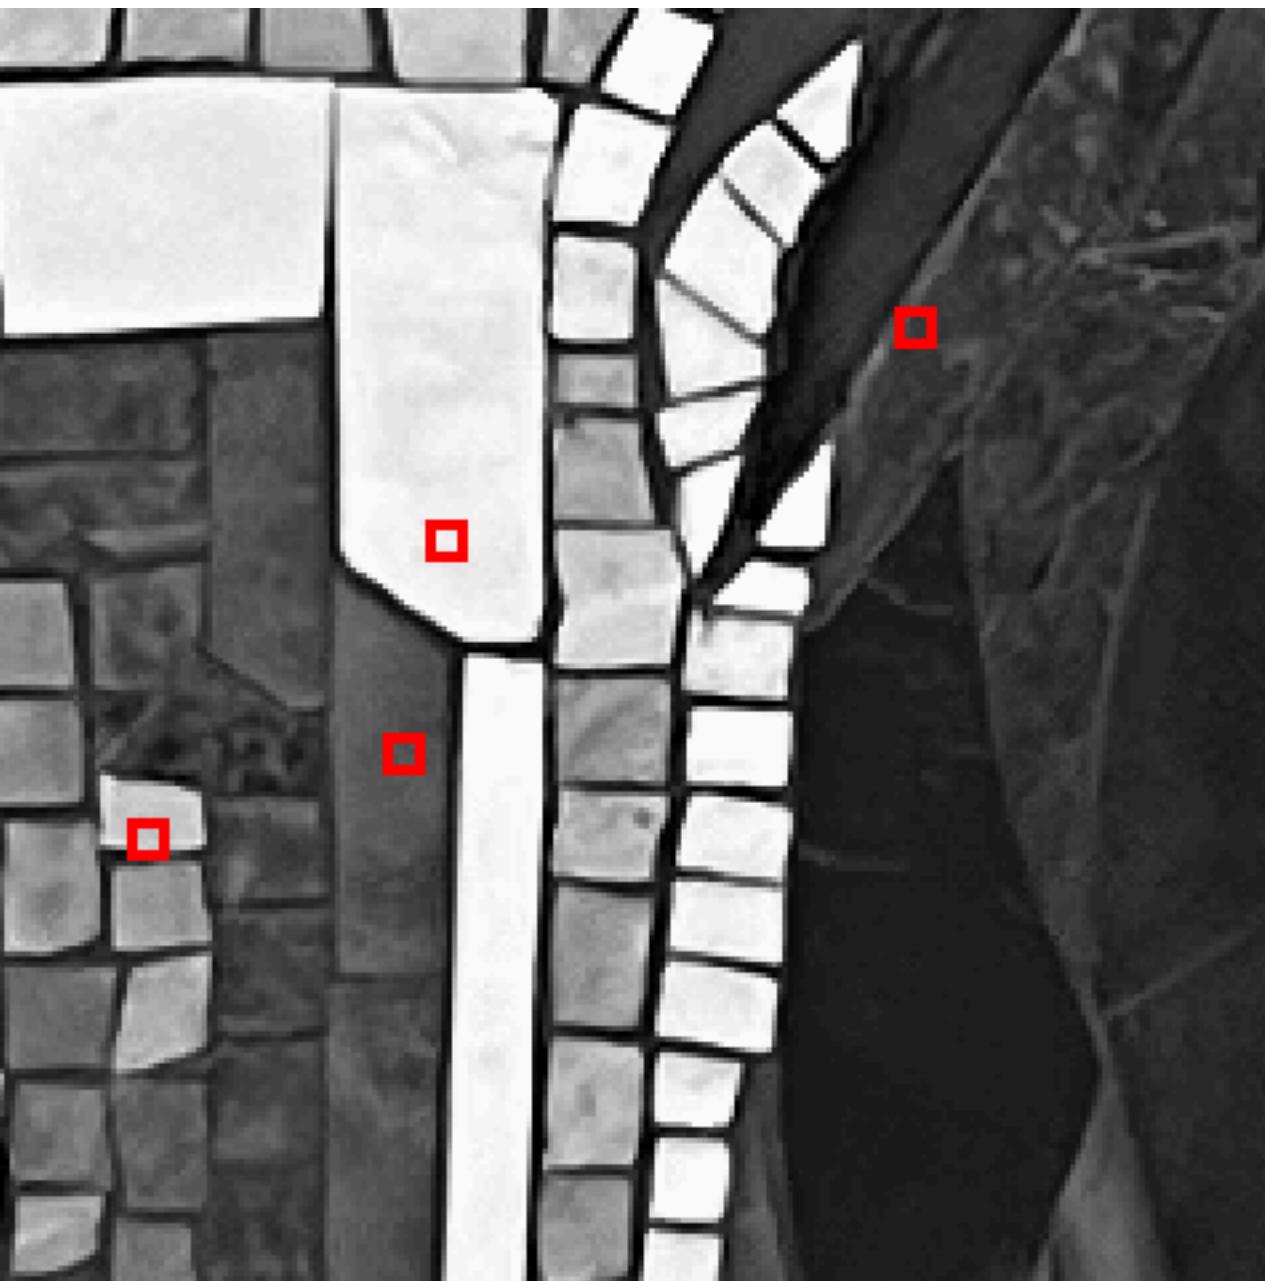
240 x 240

Exploit Uniformity over 8x8 blocks

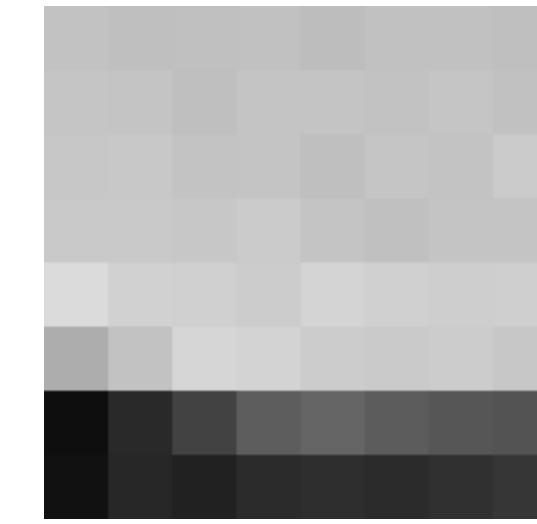


Want to keep detail, but store less

Exploit Uniformity over 8x8 blocks

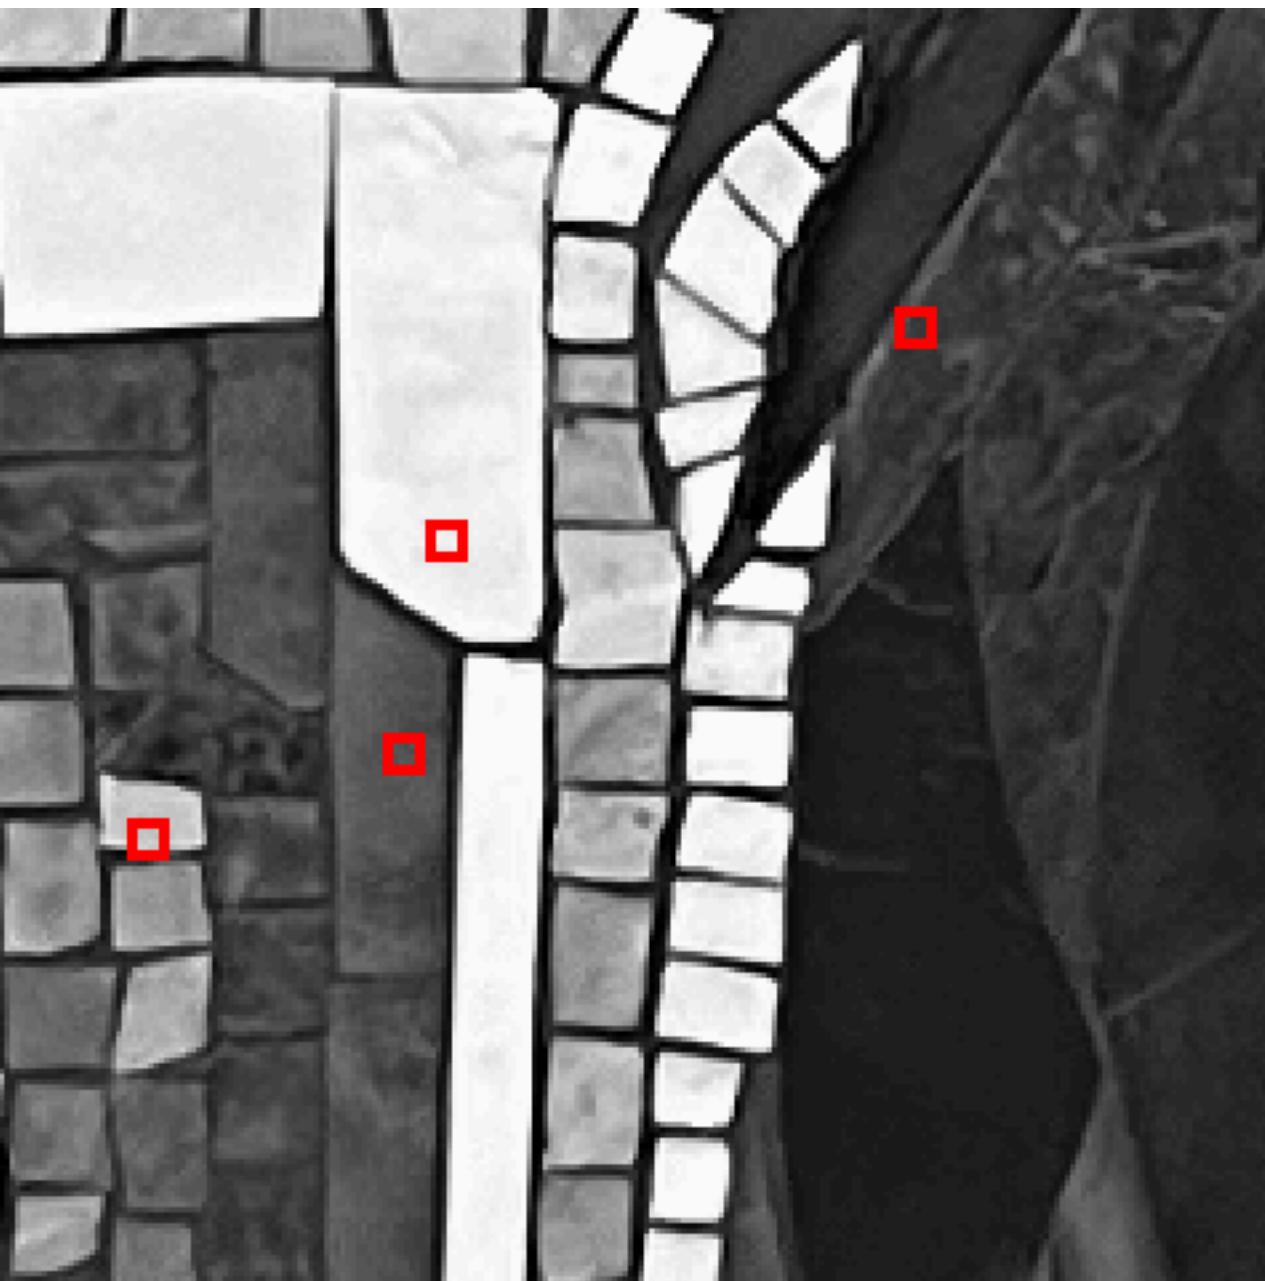


Many blocks look like
these, some texture,
but mostly samey



Want to keep detail, but store less

Exploit Uniformity over 8x8 blocks



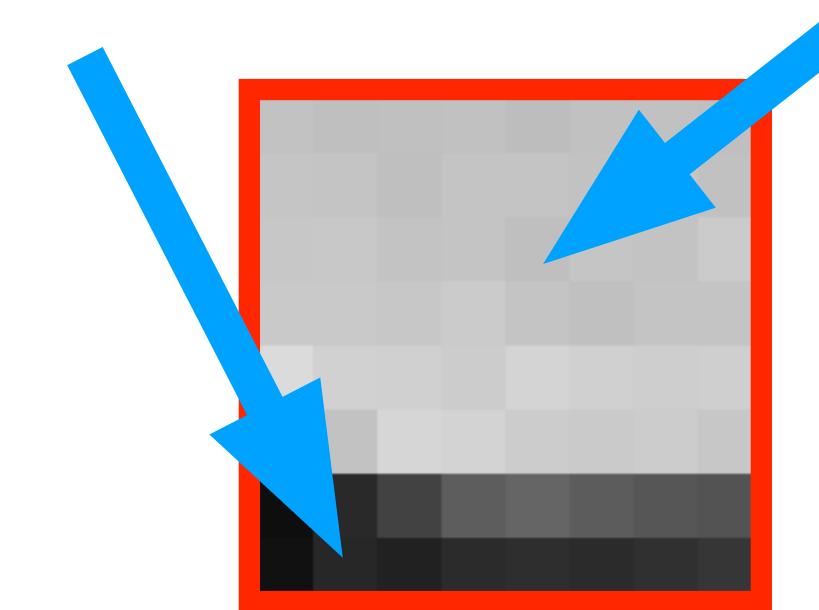
Many blocks look like
these, some texture,
but mostly samey



Dark down here



Light grey here

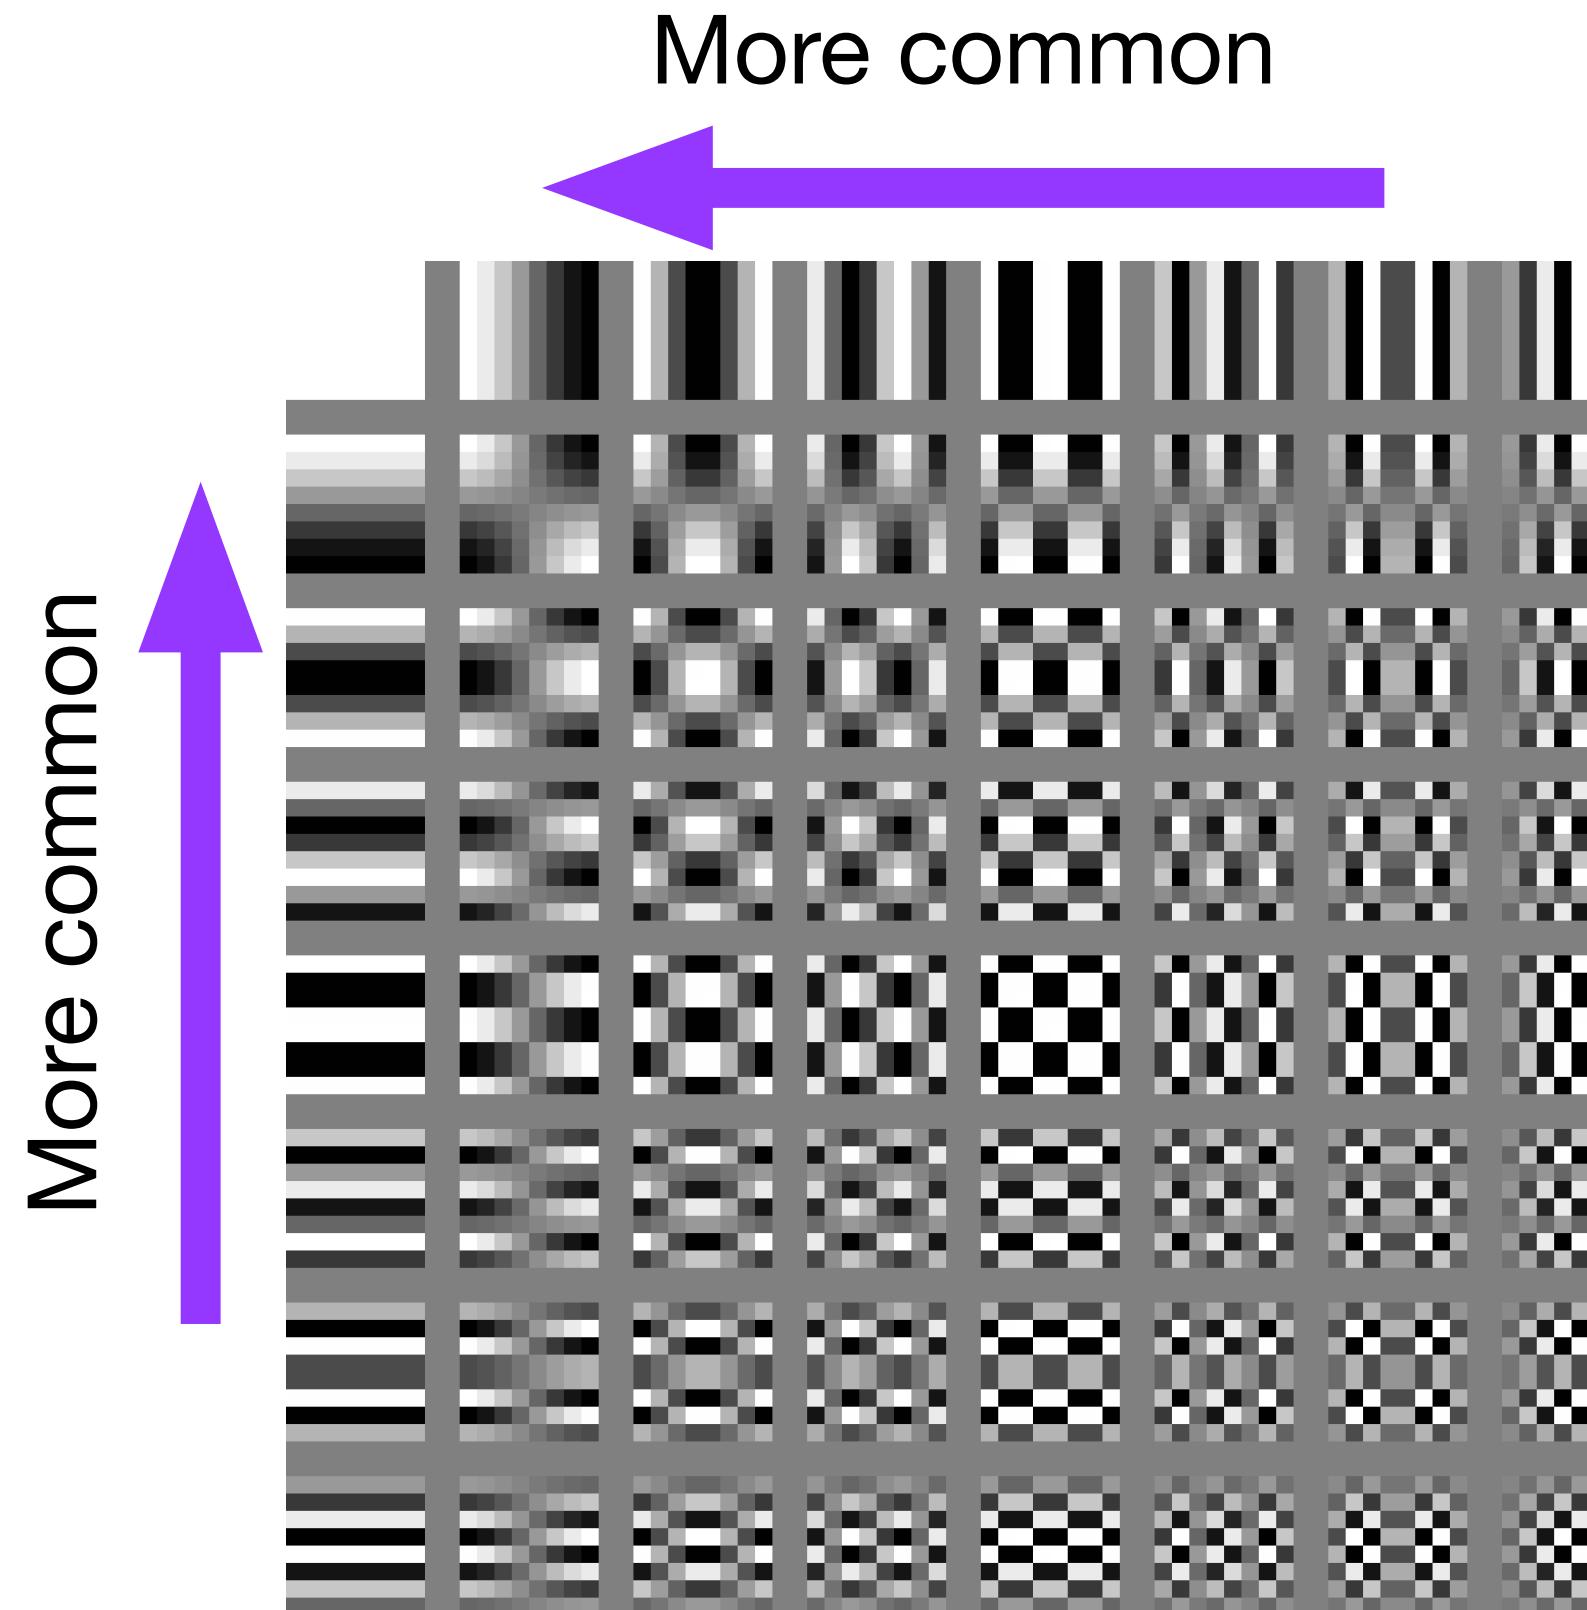


Even this block has
structure



Want to keep detail, but store less

Some Patterns are more common than others

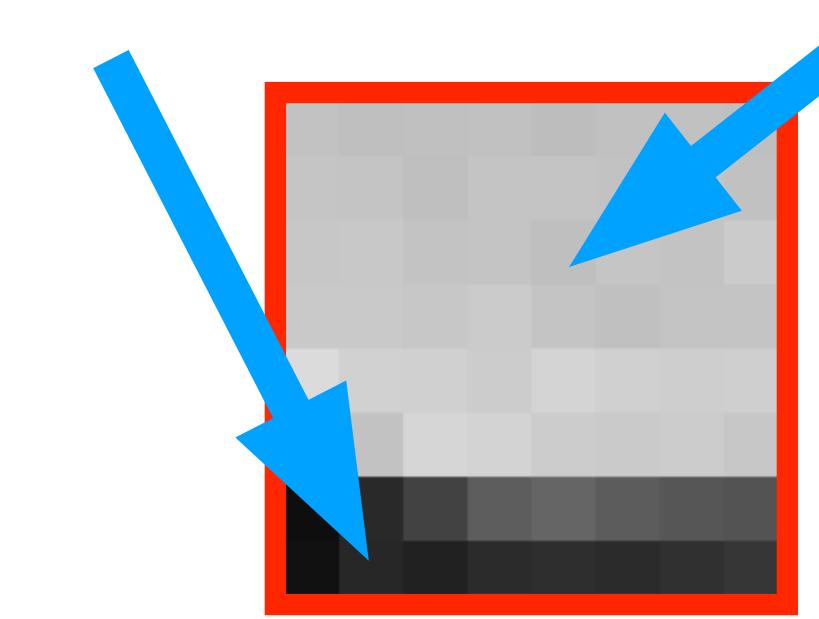


Many blocks look like these, some texture, but mostly samey



Dark down here

Light grey here

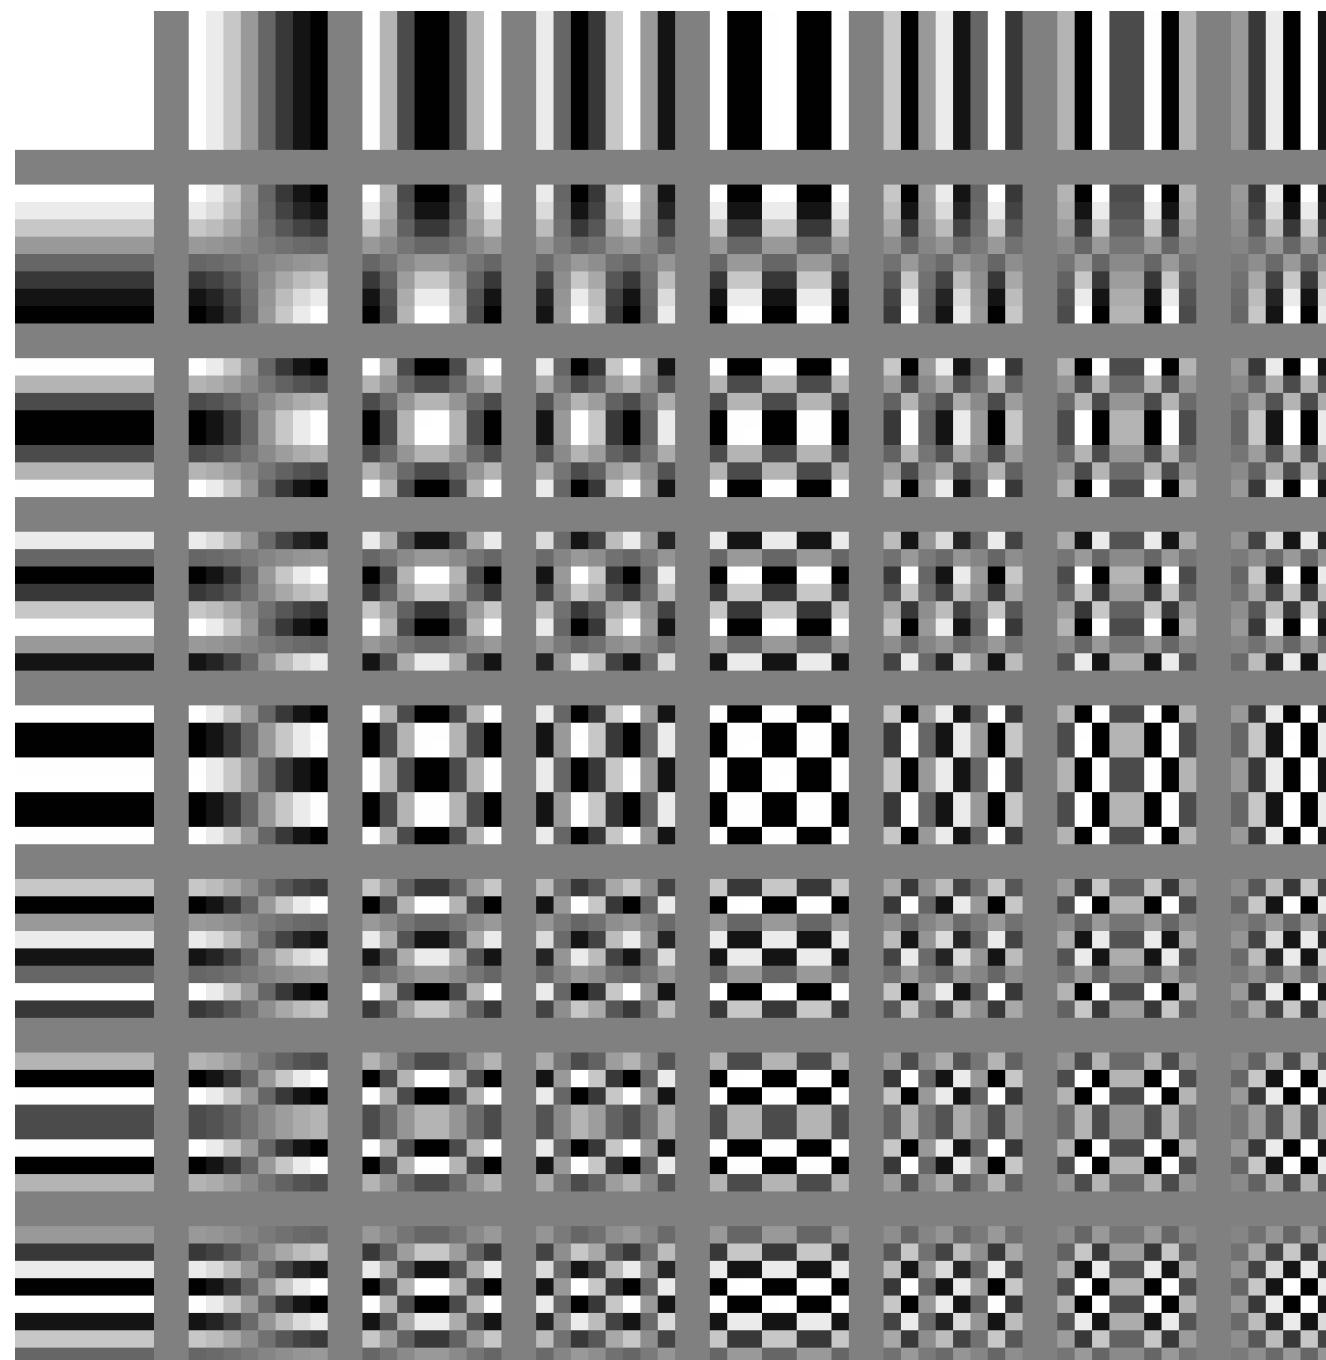


Even this block has structure



Want to keep detail, but store less

Discrete Cosine Transform (DCT)



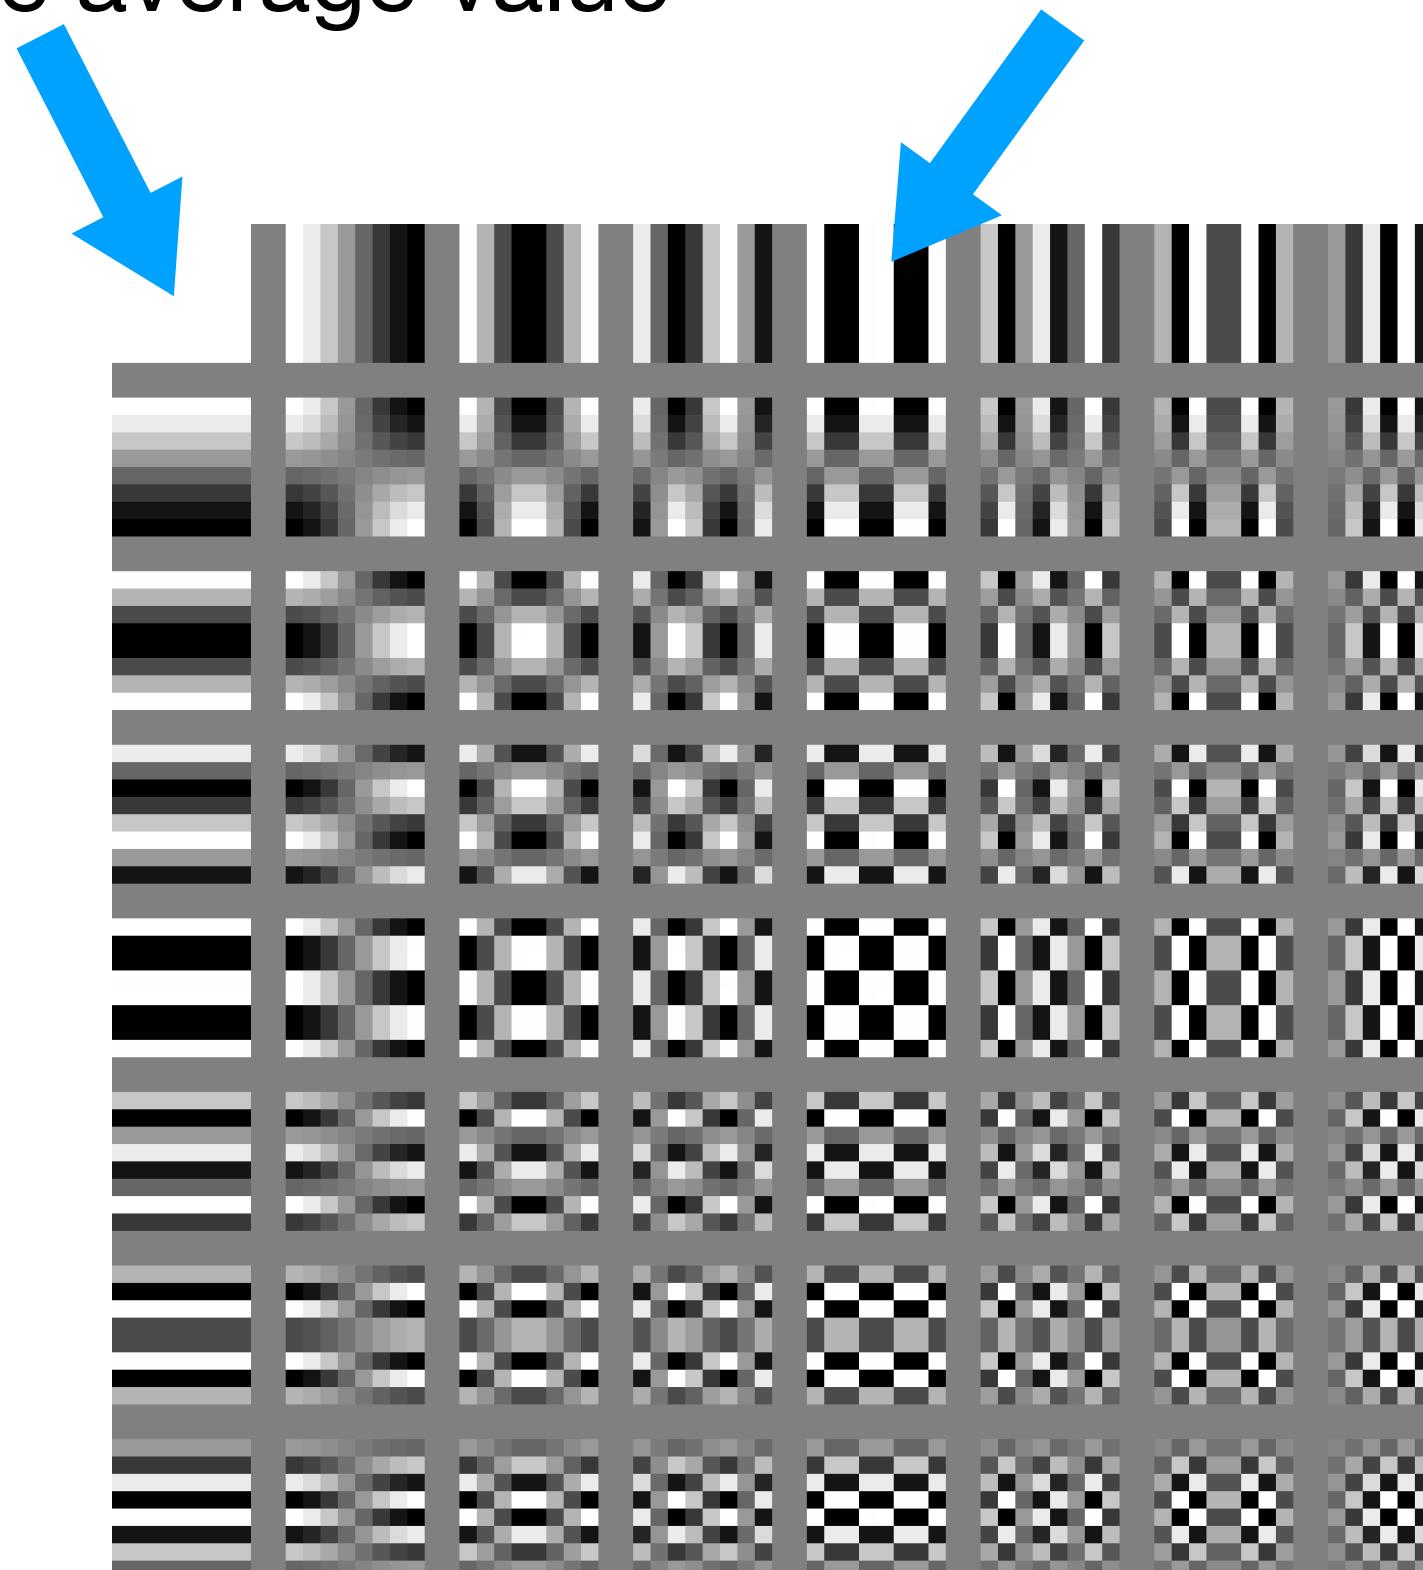
Can think of an 8x8 grid of pixels as a vector in $[255]^{64}$

The “usual” basis for this vector space is each individual pixel

The DCT gives us a different basis

Discrete Cosine Transform (DCT)

The average value



The average difference between pixels
in these regions

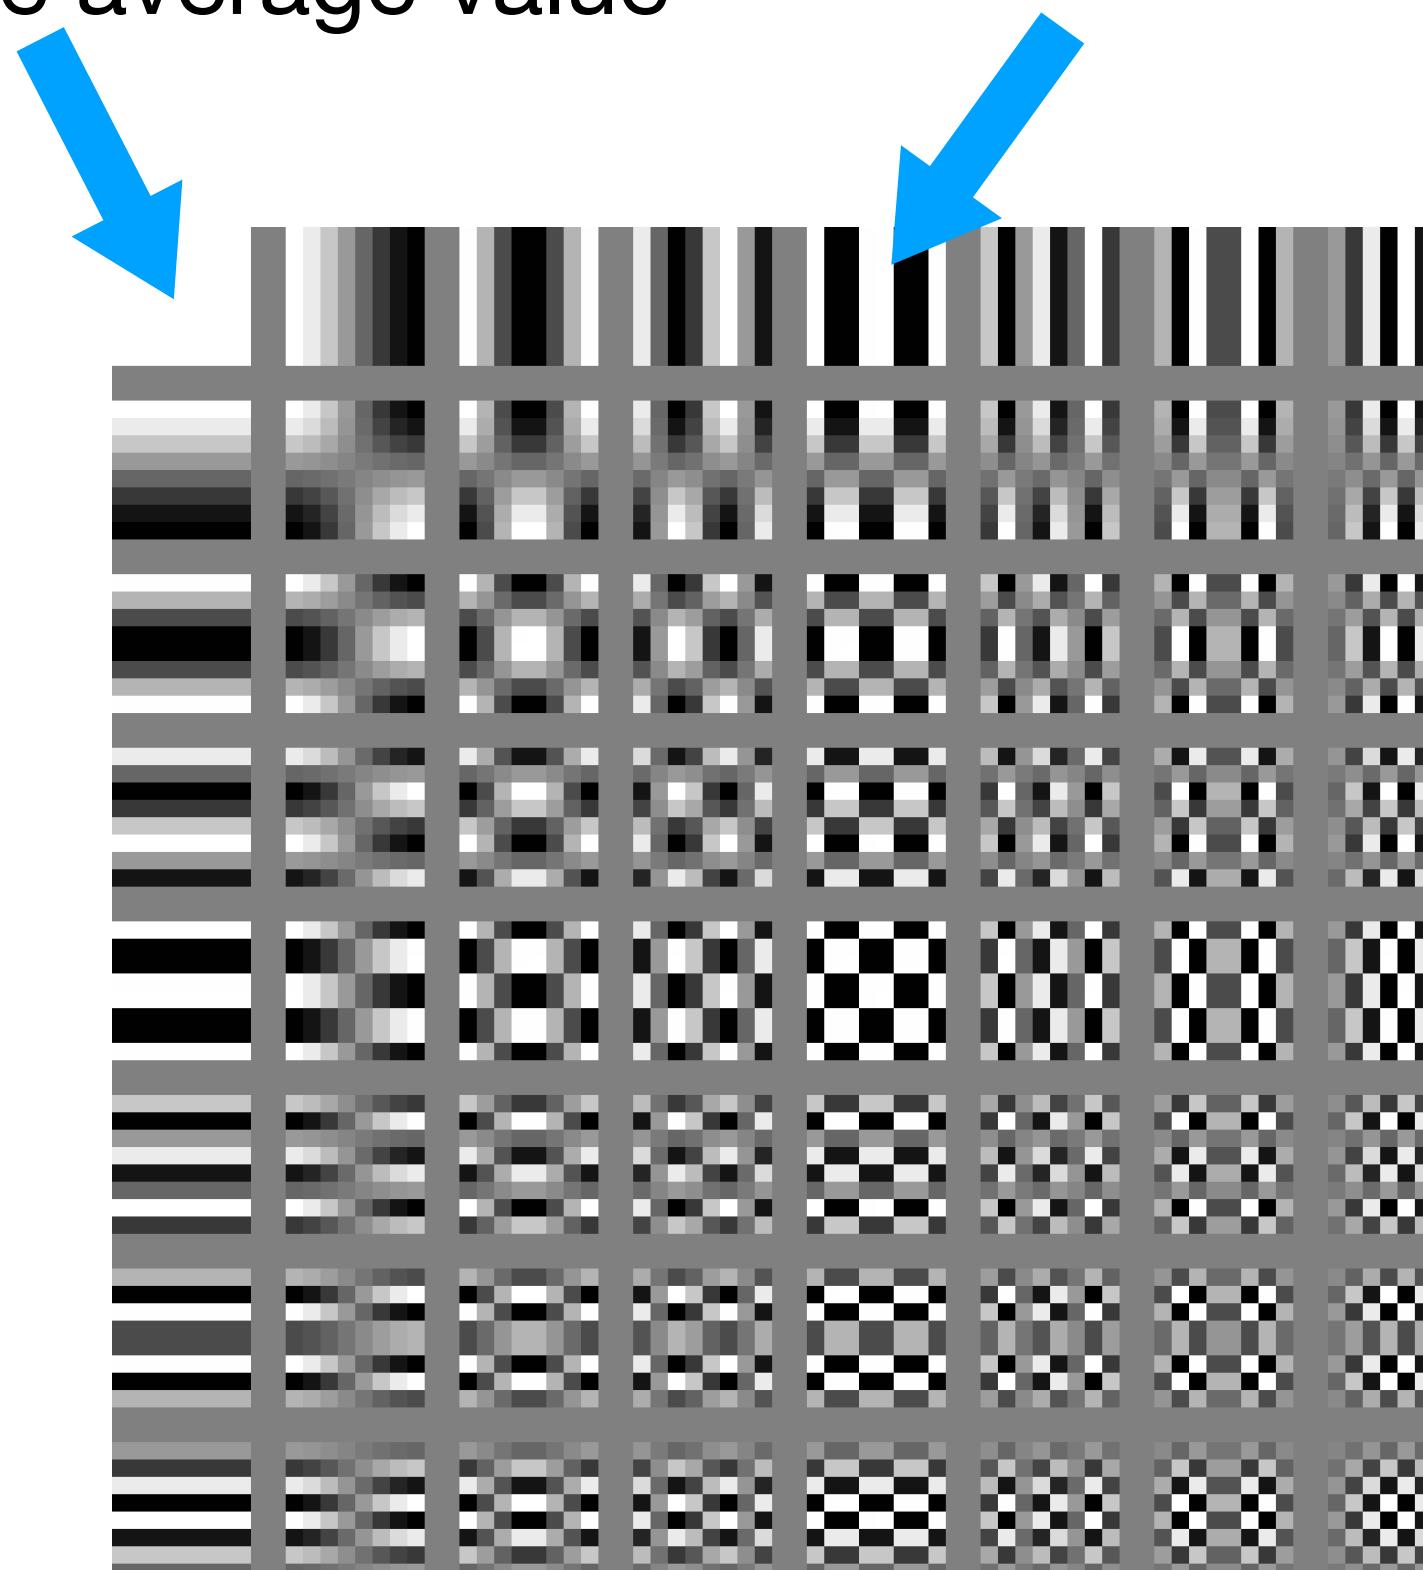
Can think of an 8x8 grid of pixels as a vector in $[255]^{64}$

The “usual” basis for this vector space is each individual pixel

The DCT gives us a different basis

Discrete Cosine Transform (DCT)

The average value



The average difference between pixels
in these regions

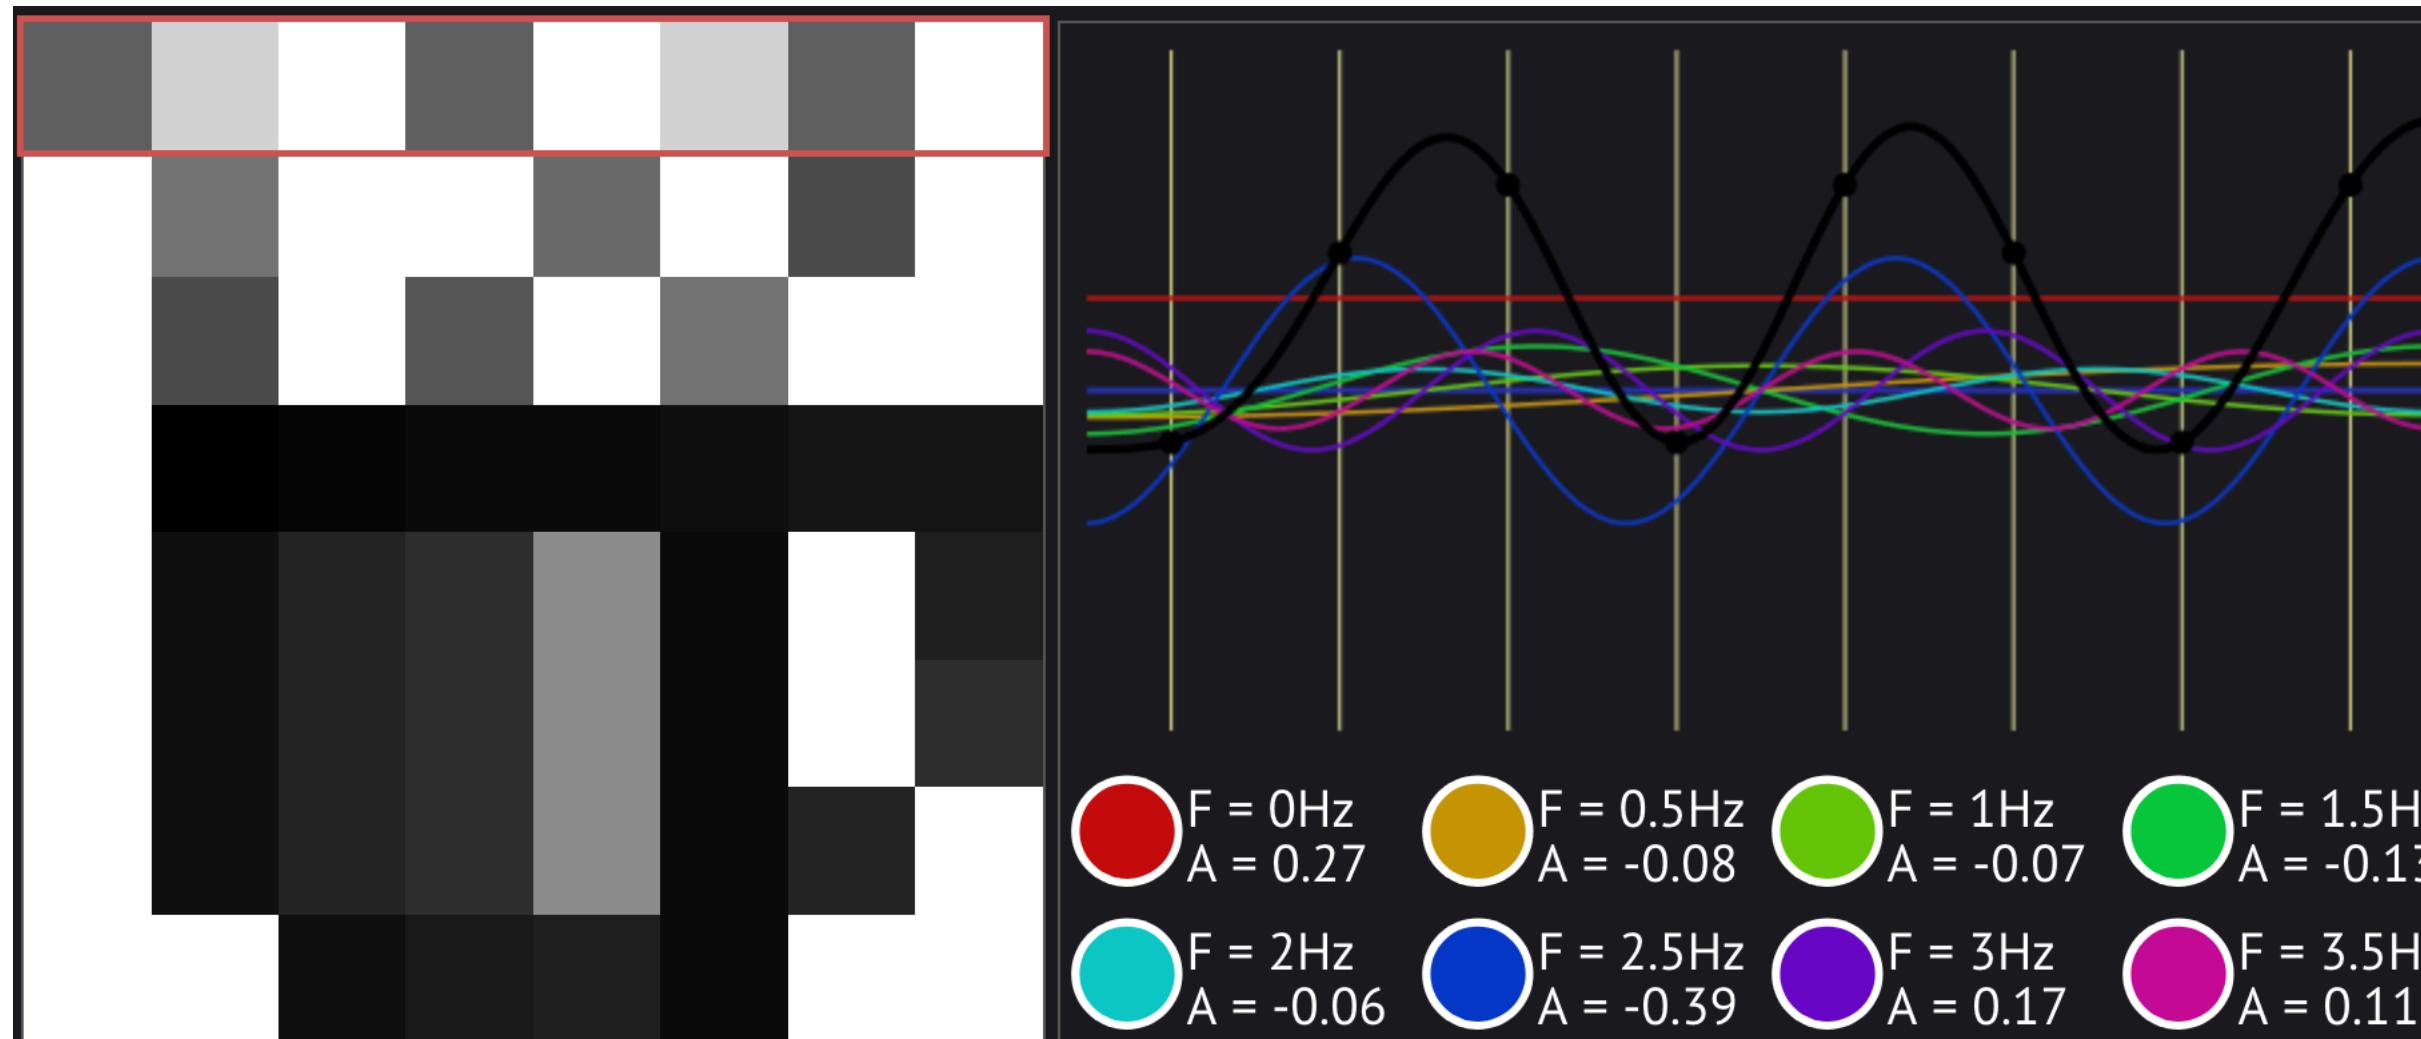
Can think of an 8x8 grid of pixels as a vector in $[255]^{64}$

The “usual” basis for this vector space is each individual pixel

The DCT gives us a different basis

We can represent the 8x8 block as a weighted
average of these patterns

Discrete Cosine Transform (DCT)



Express a given row as the integer intersections of a sum of cosine functions

$$X_k = \sum_{n=0}^{N-1} x_n \cos\left[\frac{\pi}{N} \left(n + \frac{1}{2}\right) k\right], \quad k = 0, 1, \dots, N-1$$

A type of Fourier analysis

[https://alexdowad.github.io/
visualizing-the-idct/](https://alexdowad.github.io/visualizing-the-idct/)



Discrete Cosine Transform (DCT)



Discrete Cosine Transform (DCT)

387	2	-2	1	-1	-1	2	1
9	6	0	-1	0	-2	0	-1
-7	0	4	-1	-1	-4	-3	1
2	1	-5	4	-1	-3	3	-1
5	-3	-4	-1	-2	-2	1	0
3	3	-5	-5	0	-1	-2	1
-1	-2	0	0	1	-1	2	-1
1	0	0	2	0	2	-1	1



Discrete Cosine Transform (DCT)

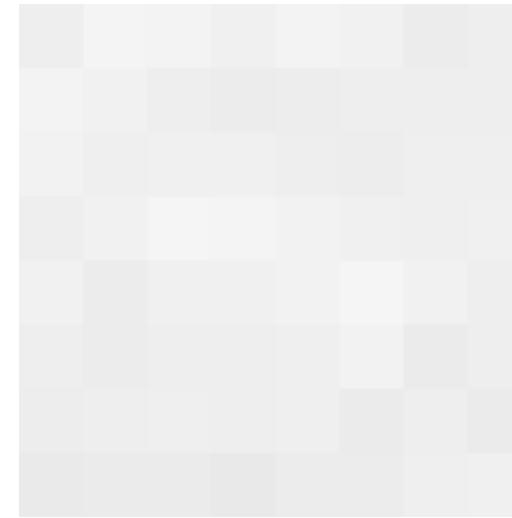
387	2	-2	1	-1	-1	2	1
9	6	0	-1	0	-2	0	-1
-7	0	4	-1	-1	-4	-3	1
2	1	-5	4	-1	-3	3	-1
5	-3	-4	-1	-2	-2	1	0
3	3	-5	-5	0	-1	-2	1
-1	-2	0	0	1	-1	2	-1
1	0	0	2	0	2	-1	1



Most of the coeffs are very small (except top left)

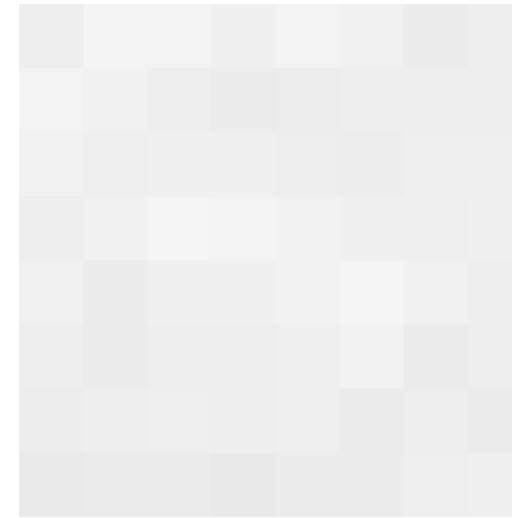
Discrete Cosine Transform (DCT)

-519	76	5	-6	-6	3	1	2
24	24	-5	-27	-7	-3	-4	-2
-10	-14	-22	-16	9	0	-2	-1
3	6	-18	0	8	7	-1	1
0	-10	-4	0	13	-5	-6	5
1	1	-3	5	0	-7	-4	4
-2	-4	0	0	-1	-4	-2	2
1	0	0	-1	-2	-2	0	1



Discrete Cosine Transform (DCT)

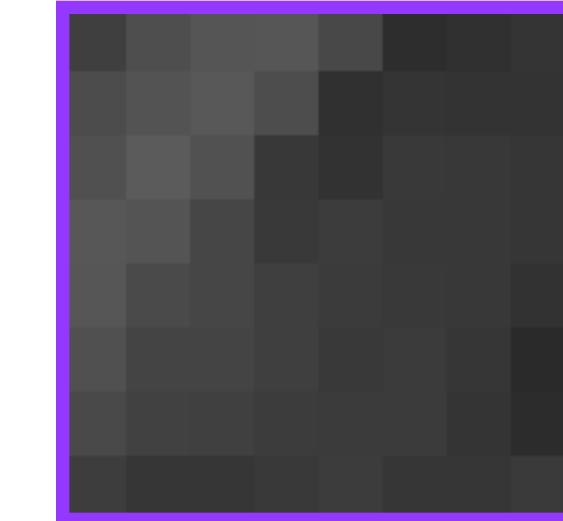
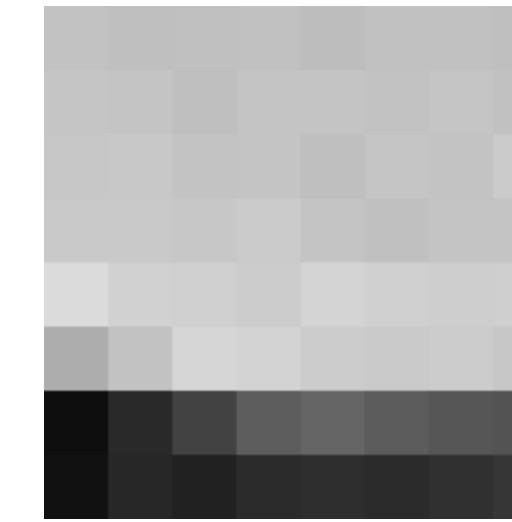
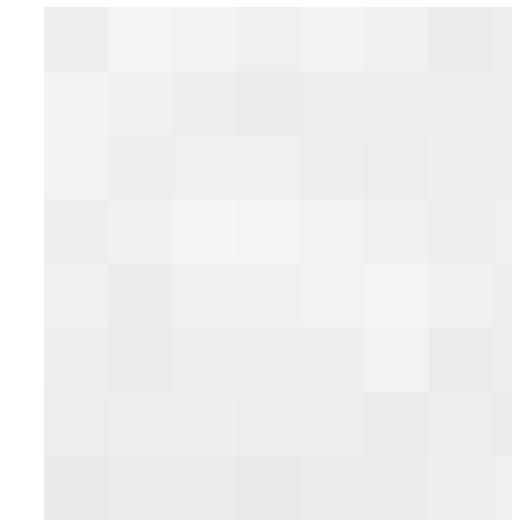
-519	76	5	-6	-6	3	1	2
24	24	-5	-27	-7	-3	-4	-2
-10	-14	-22	-16	9	0	-2	-1
3	6	-18	0	8	7	-1	1
0	-10	-4	0	13	-5	-6	5
1	1	-3	5	0	-7	-4	4
-2	-4	0	0	-1	-4	-2	2
1	0	0	-1	-2	-2	0	1



Larger coeffs in general, but still concentrated in the top left

Discrete Cosine Transform (DCT)

-519	76	5	-6	-6	3	1	2
24	24	-5	-27	-7	-3	-4	-2
-10	-14	-22	-16	9	0	-2	-1
3	6	-18	0	8	7	-1	1
0	-10	-4	0	13	-5	-6	5
1	1	-3	5	0	-7	-4	4
-2	-4	0	0	-1	-4	-2	2
1	0	0	-1	-2	-2	0	1

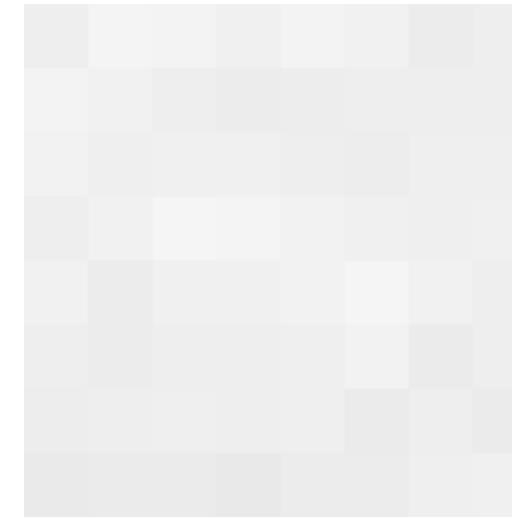


This block has more “texture”

Larger coeffs in general, but still concentrated in the top left

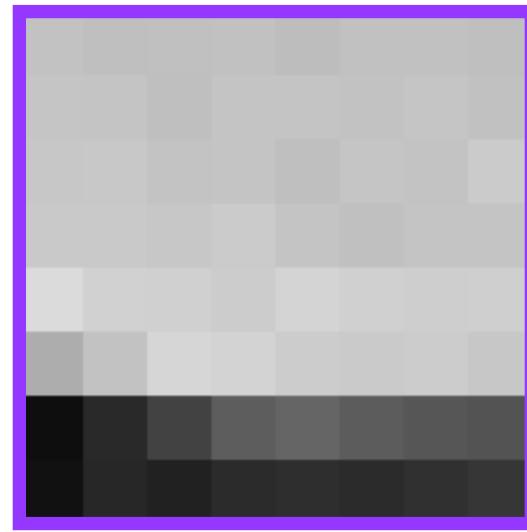
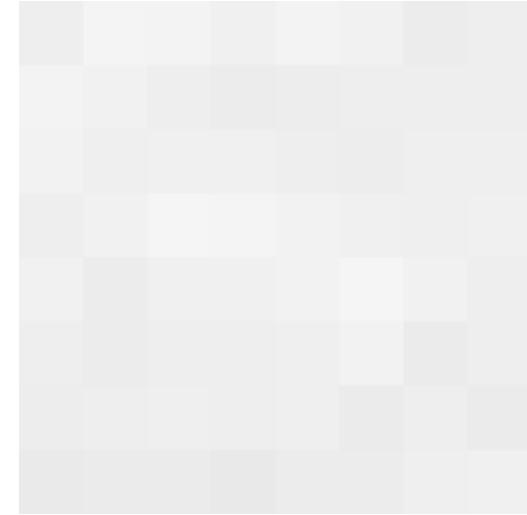
Discrete Cosine Transform (DCT)

280	-29	-23	-11	0	-2	-5	-3
349	42	33	13	3	7	5	-1
-299	-26	-11	-2	-2	-3	-5	-3
157	-2	-16	-2	-5	5	5	6
-25	22	23	8	4	-1	-1	1
-65	-21	-22	-6	1	2	0	-3
75	14	14	-9	-7	0	3	-2
-43	-10	1	9	5	-3	4	0



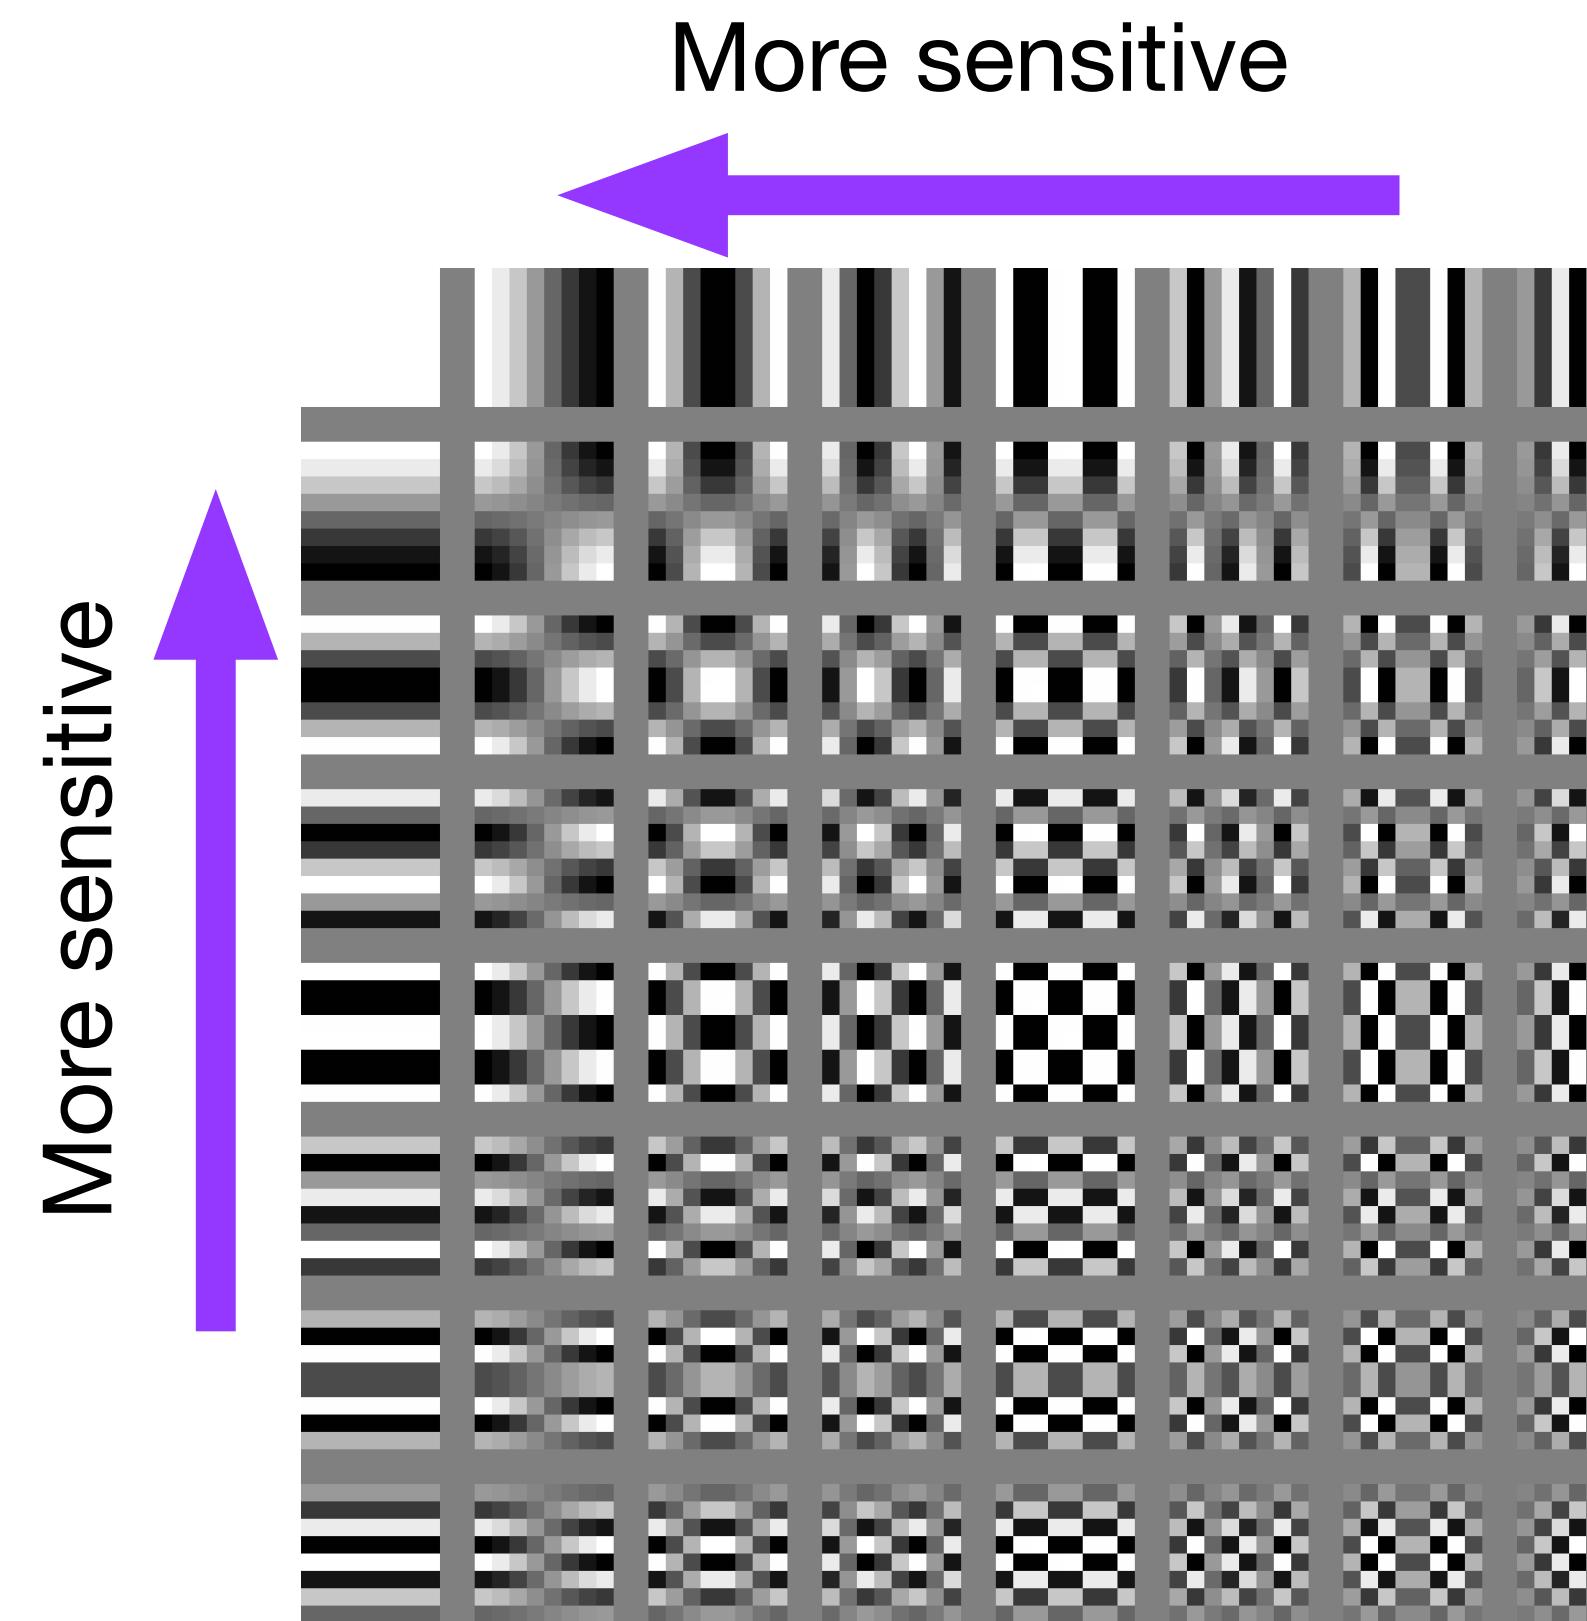
Discrete Cosine Transform (DCT)

280	-29	-23	-11	0	-2	-5	-3
349	42	33	13	3	7	5	-1
-299	-26	-11	-2	-2	-3	-5	-3
157	-2	-16	-2	-5	5	5	6
-25	22	23	8	4	-1	-1	1
-65	-21	-22	-6	1	2	0	-3
75	14	14	-9	-7	0	3	-2
-43	-10	1	9	5	-3	4	0



The pattern is still large-scale and so is mostly represented in the upper left

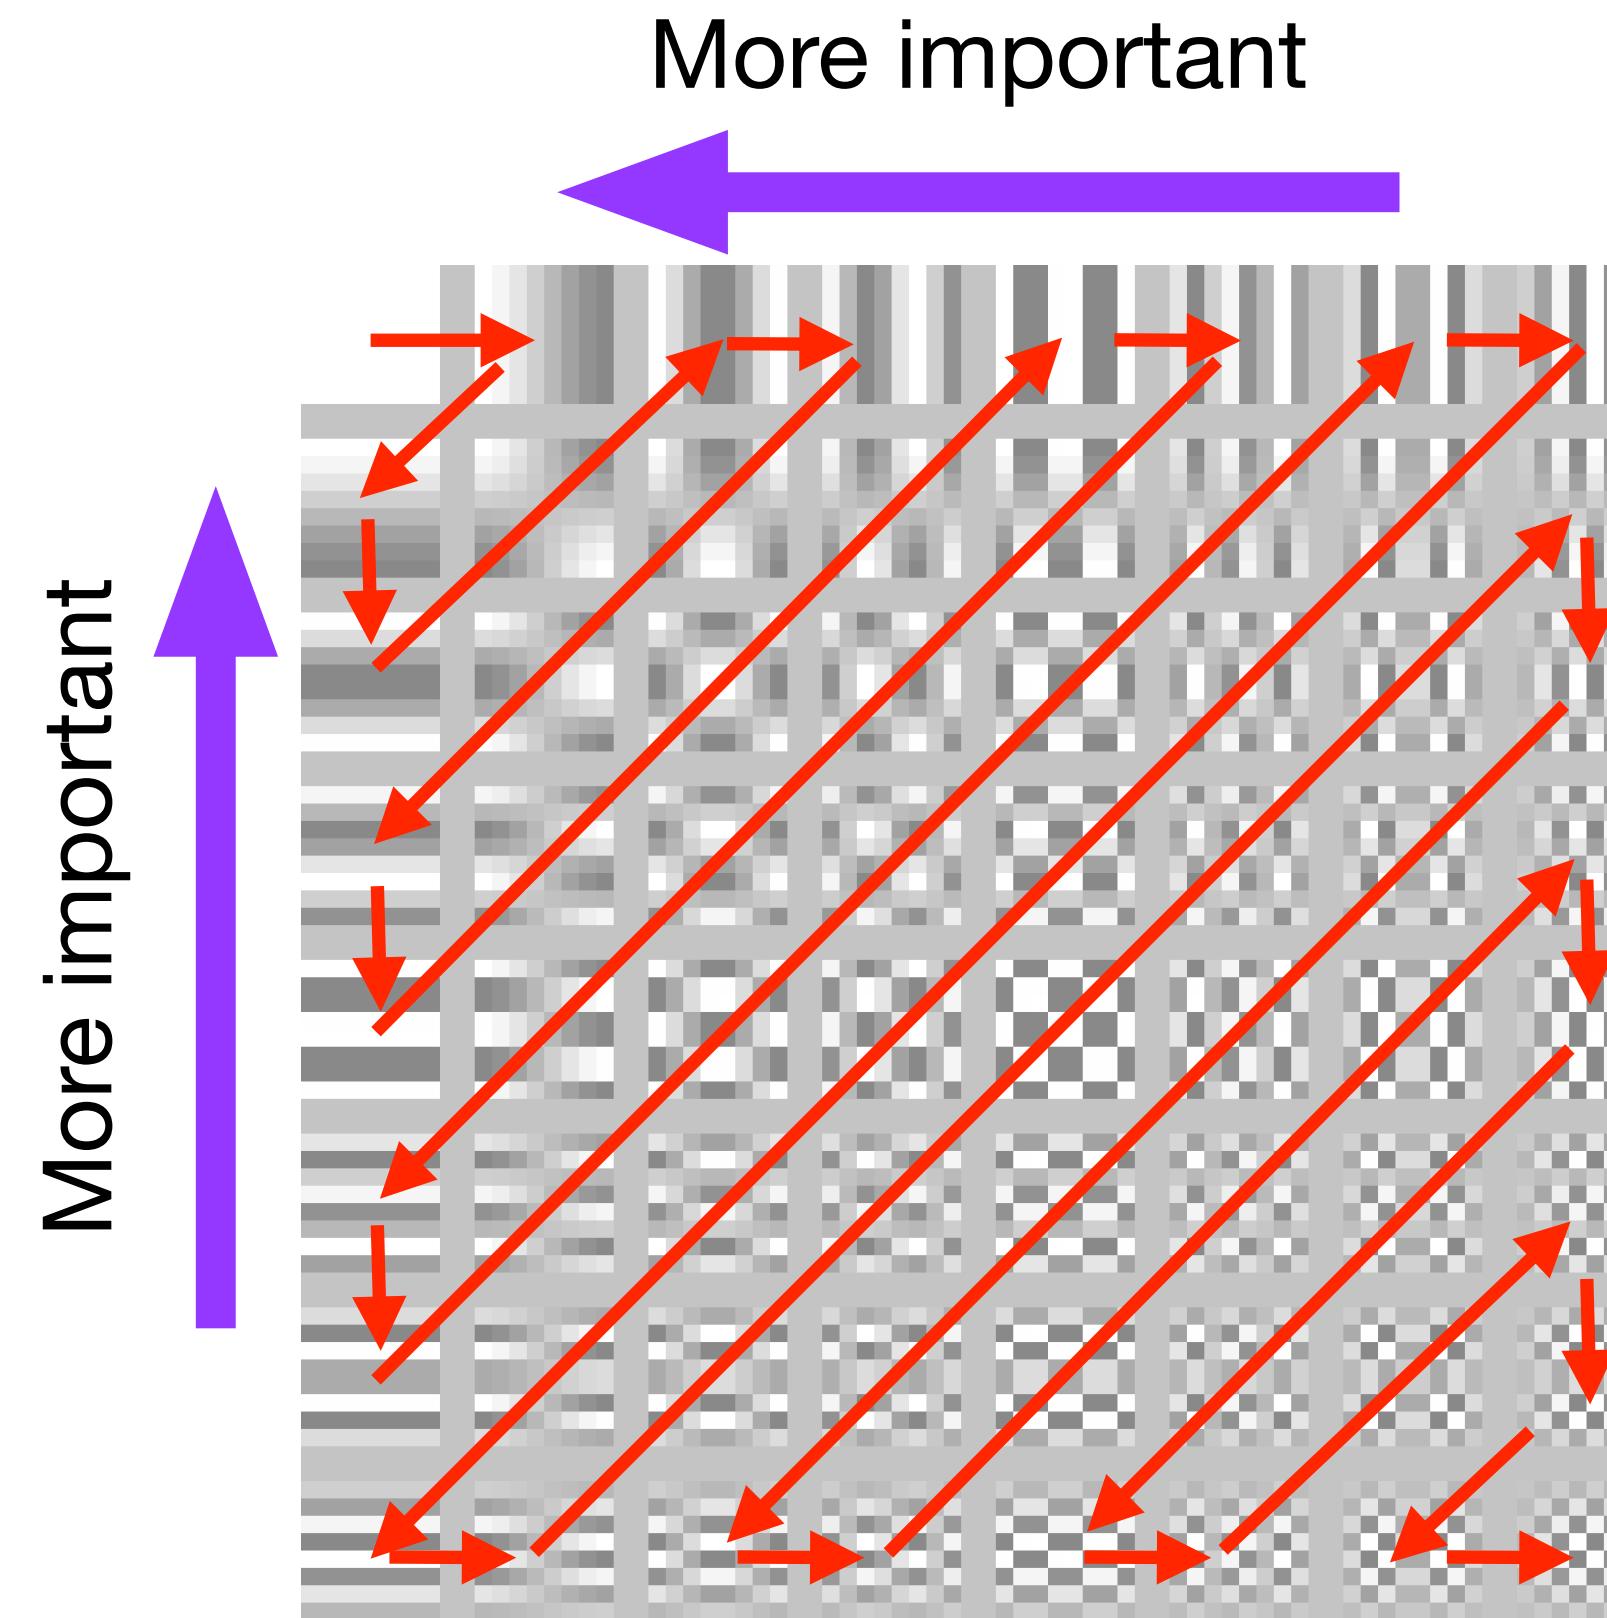
Human Perception and DCT



The human eye is more sensitive to
“blockier” signal

So we perceive the upper left coeffs much
more than the lower right

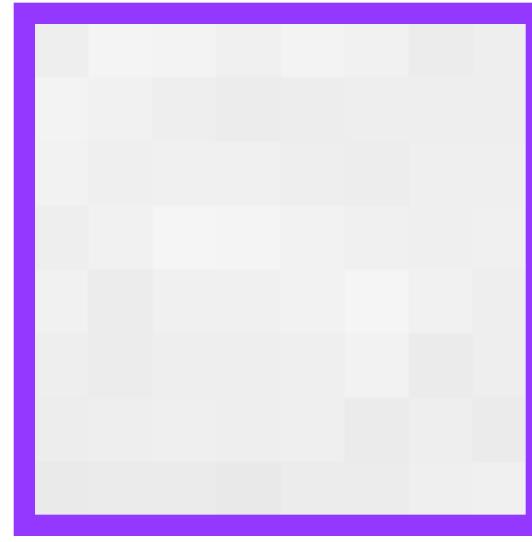
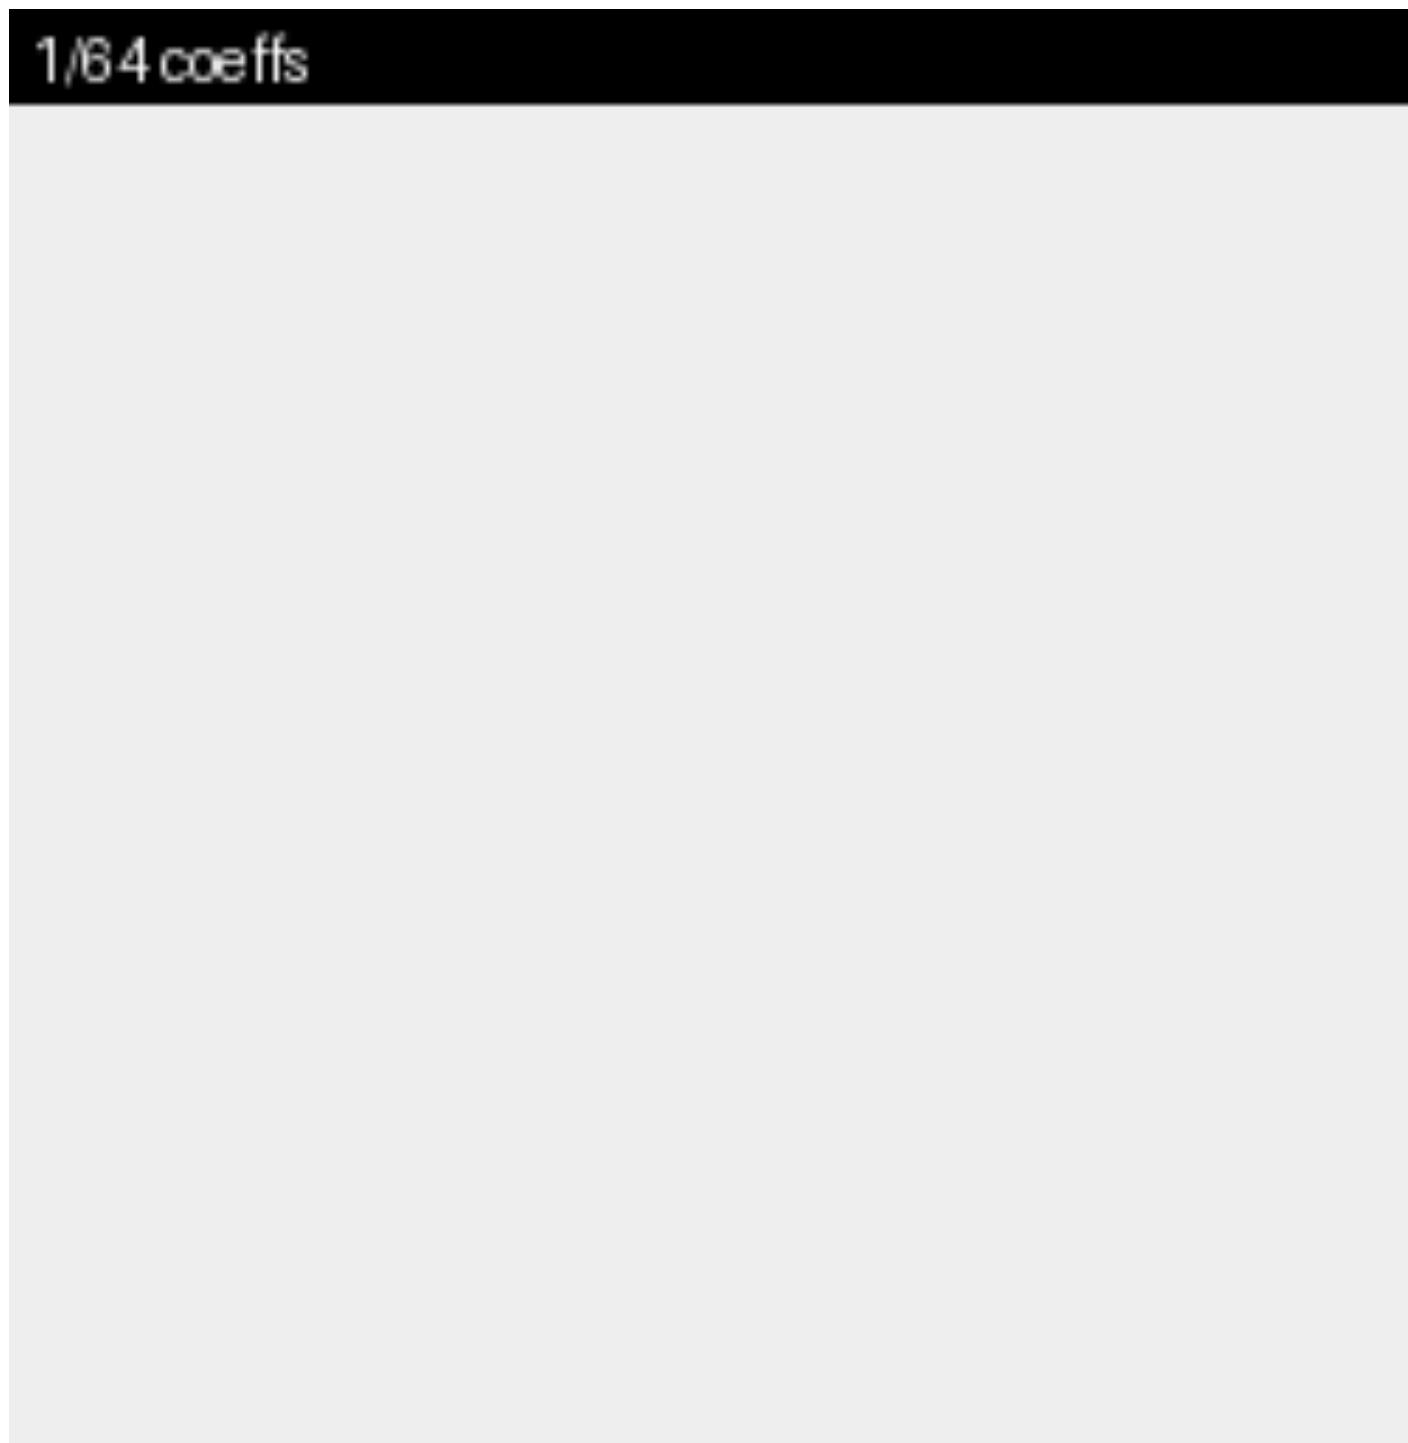
Zig-Zag Path



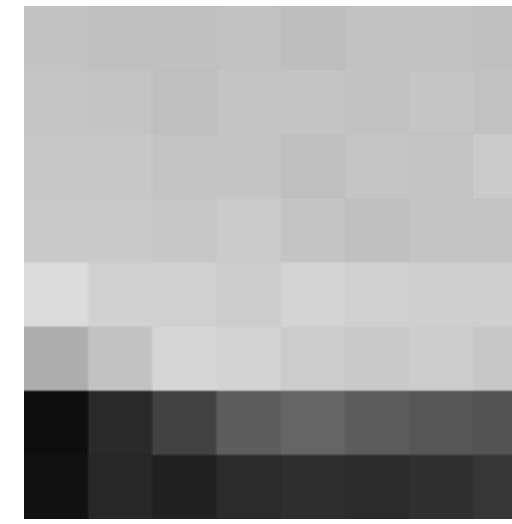
Write out the coefficients in zig-zag order:

More important coefficients first

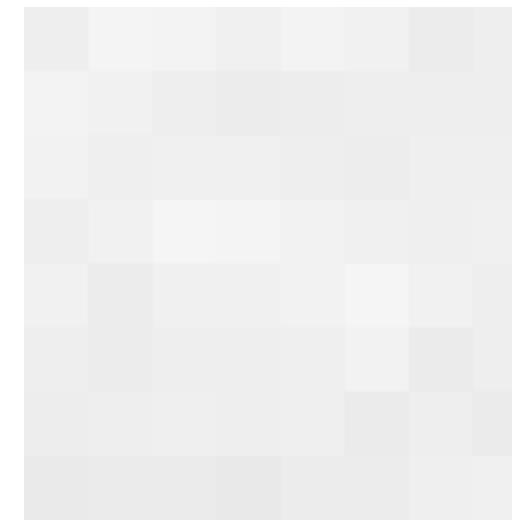
Zig-Zag Path



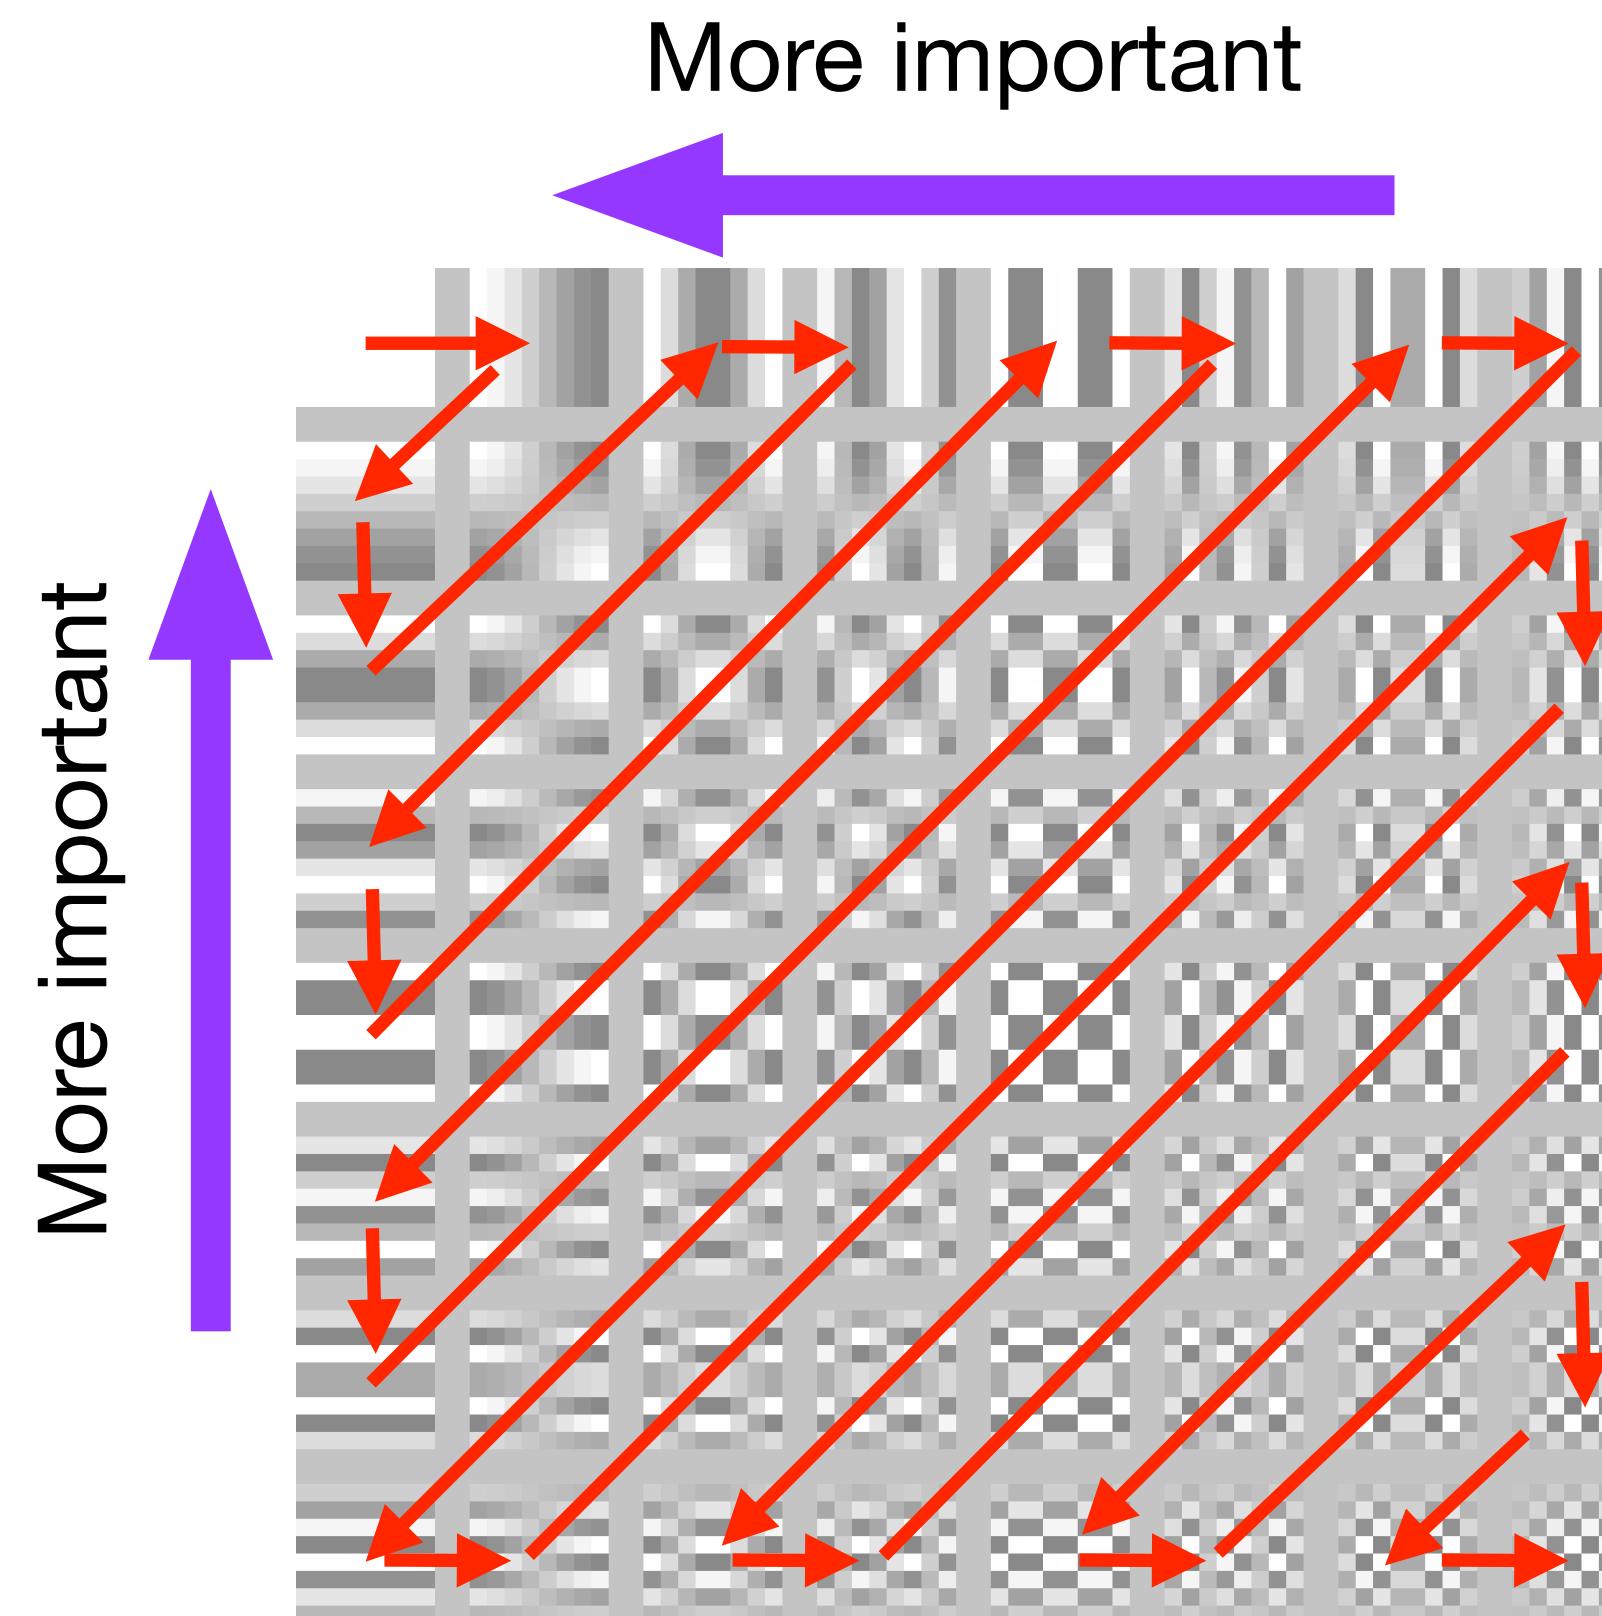
Zig-Zag Path



Zig-Zag Path



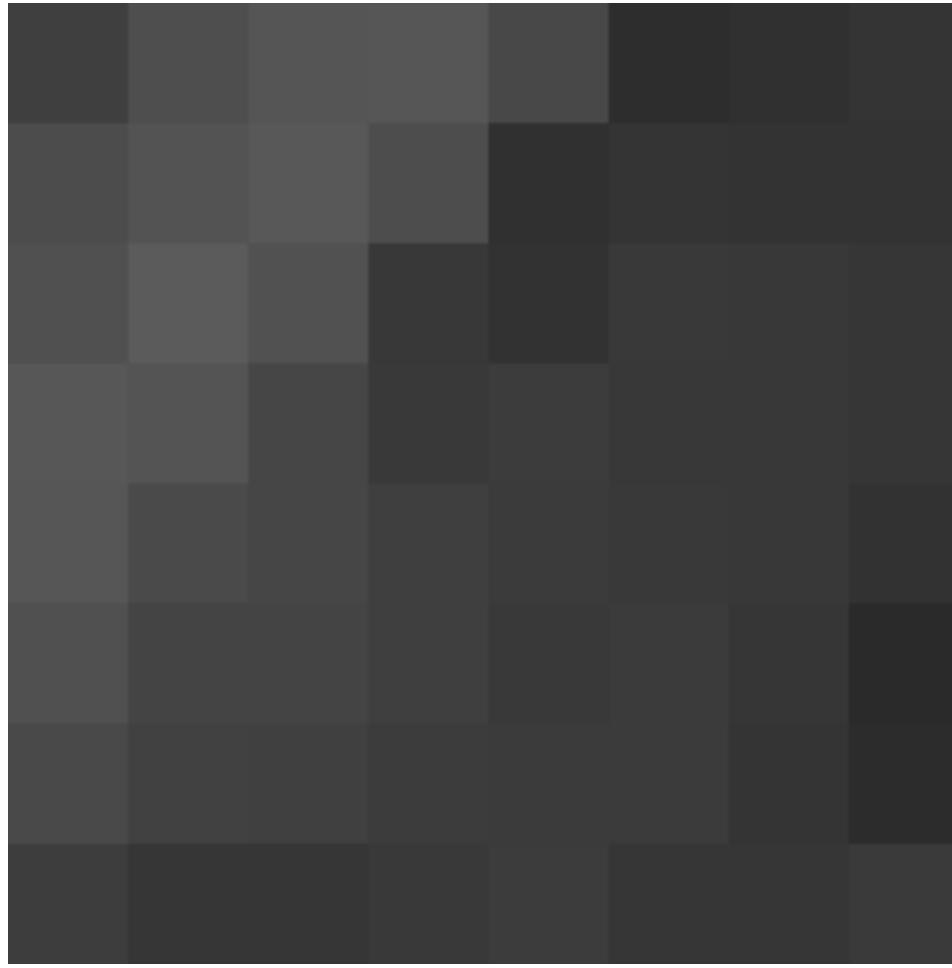
Zig-Zag and Quantization



Use more bits to store values in
upper left, less in lower right

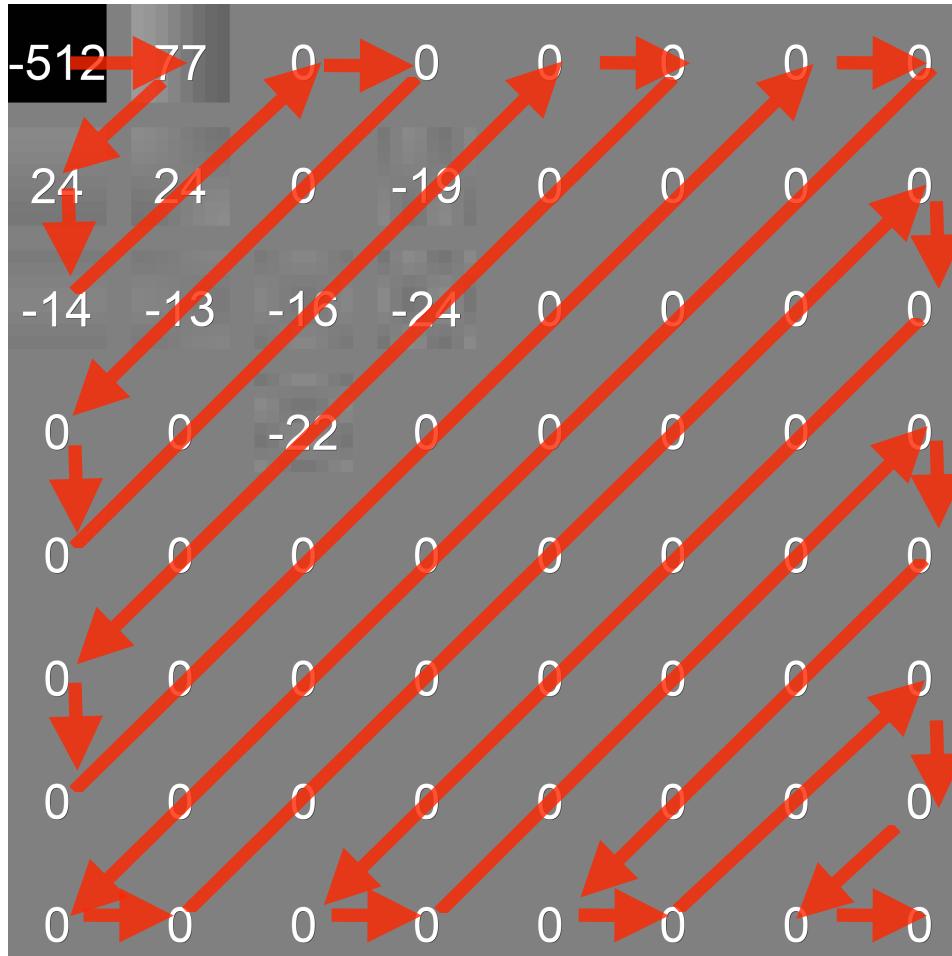
Important side effect:
Many coeffs round to 0

Run-Length Encoding



Efficiently store all these zeros

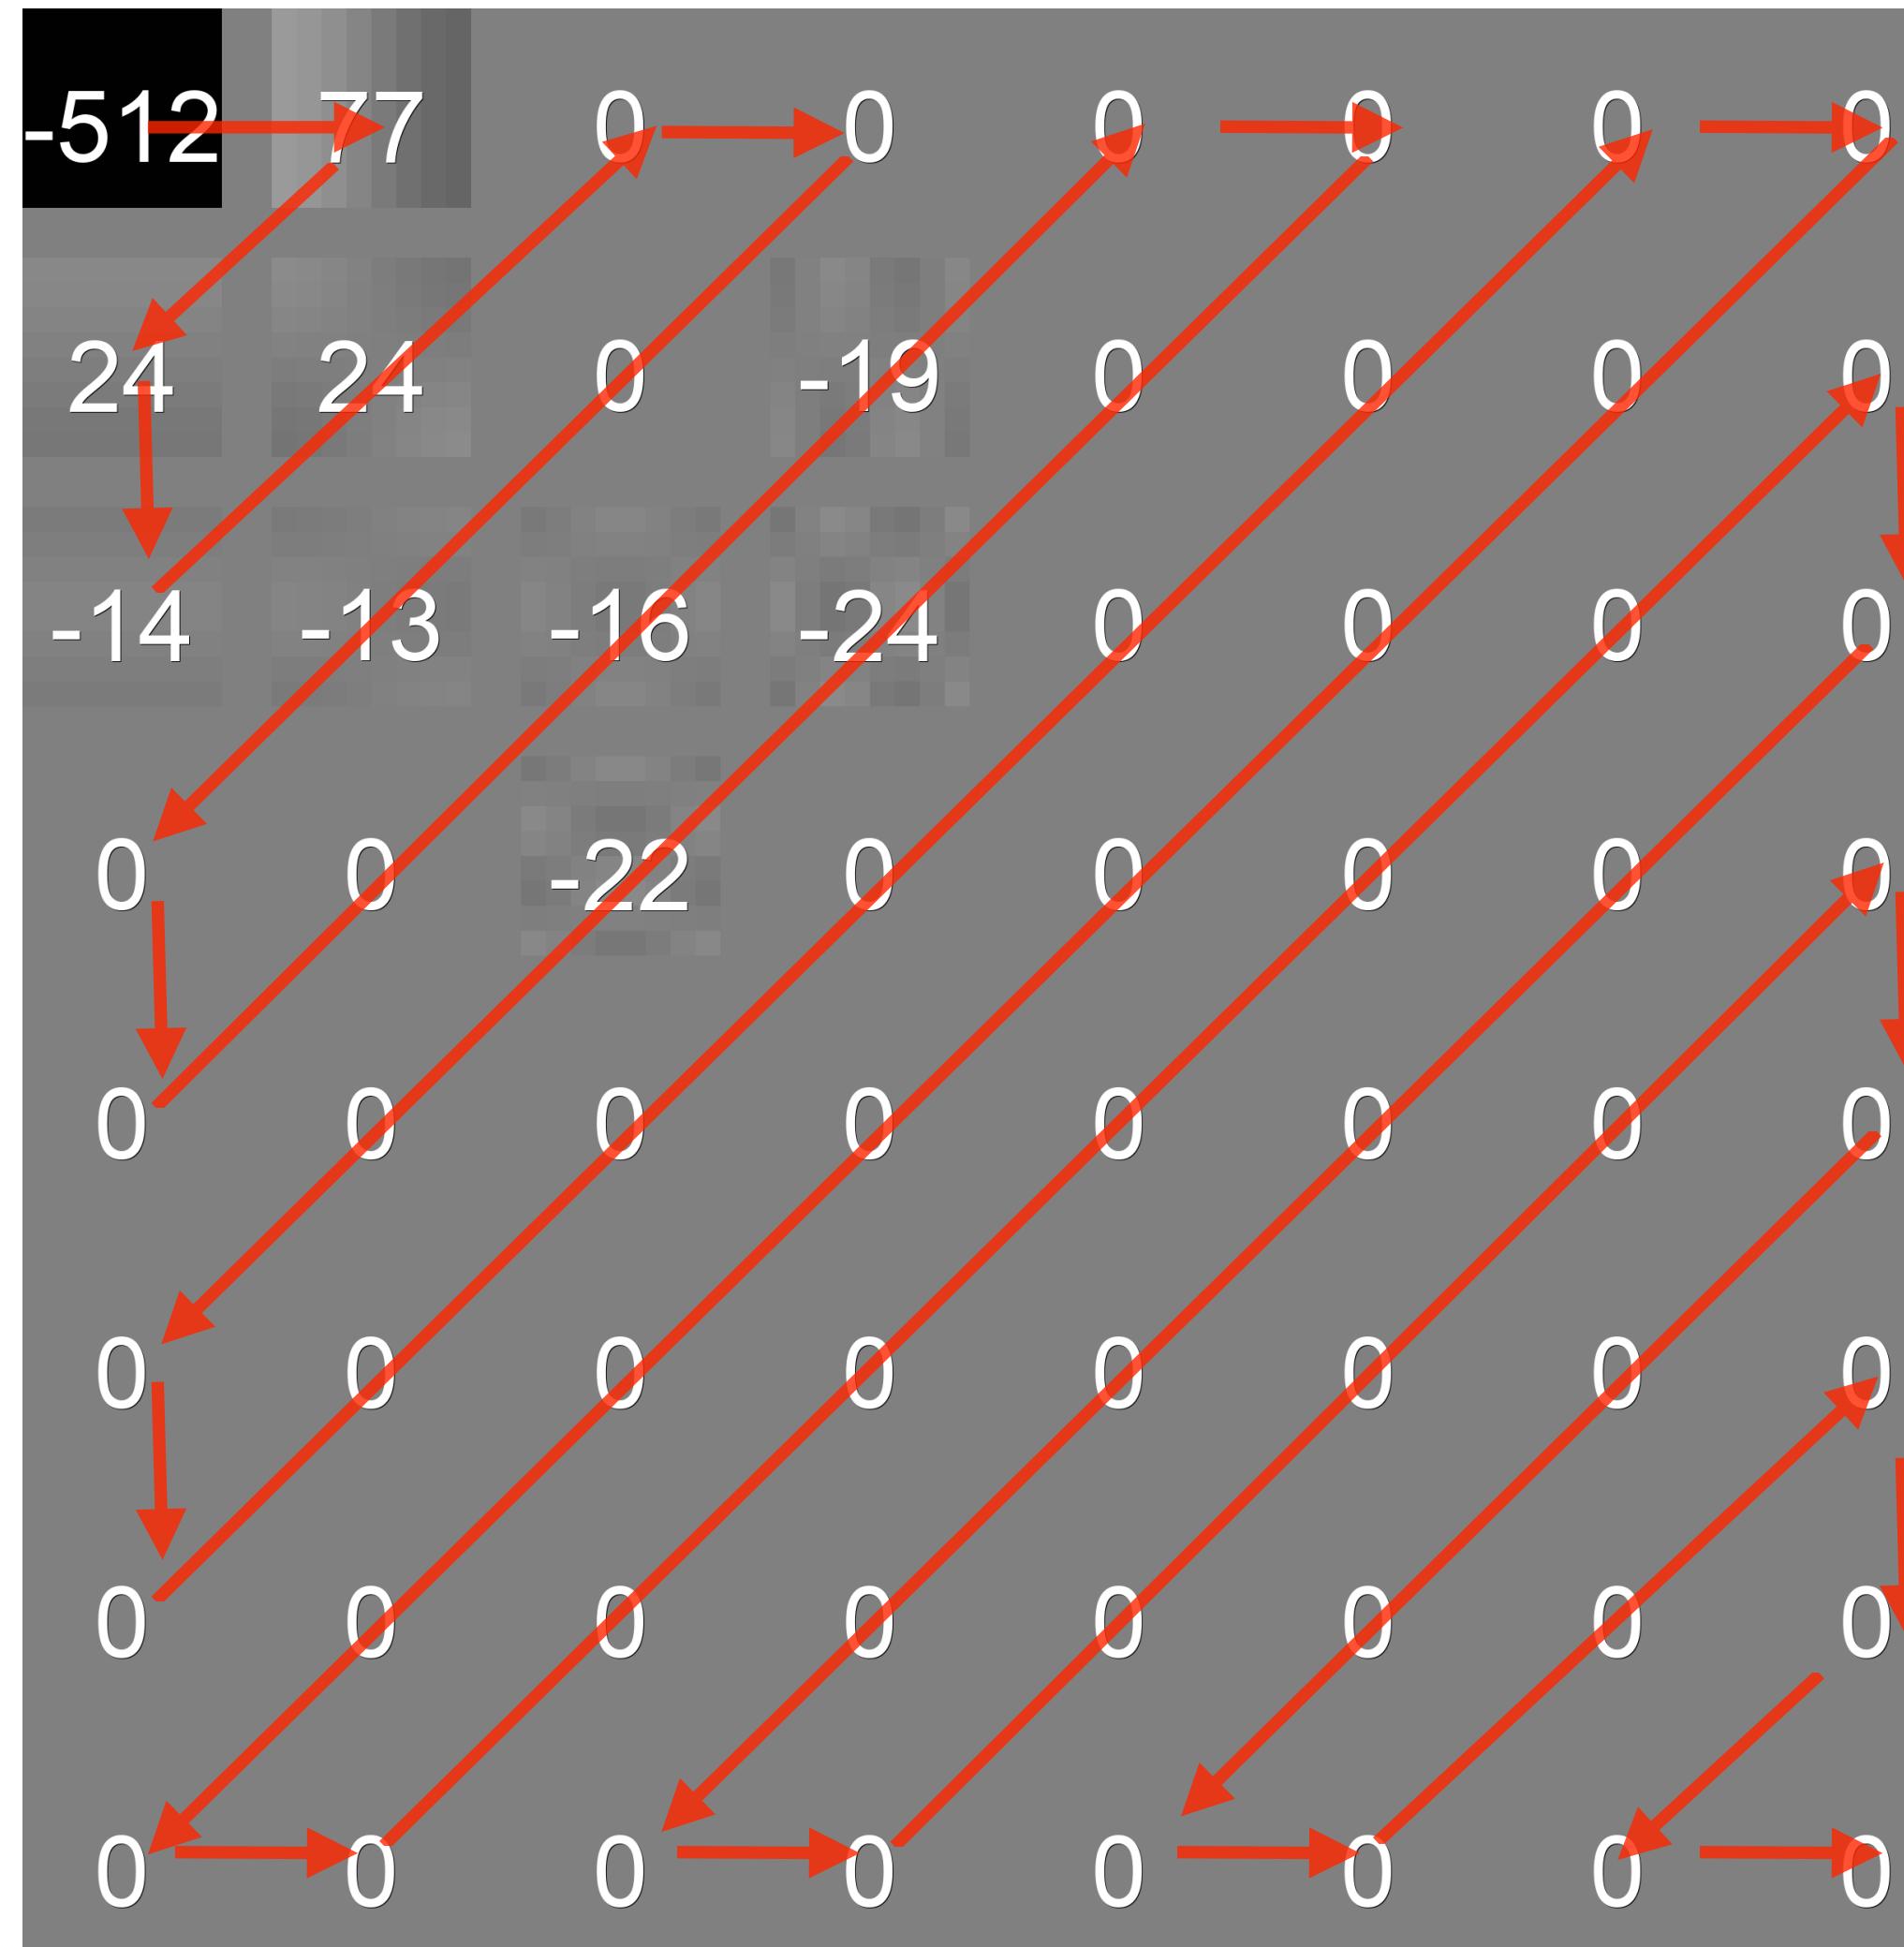
In each entry store
(# zeros, next non-zero value)



(0,-512) (0,77) (0,24) (0,14) (0,24) (3,-13)
(3,-16) (0,19) (3,-24) (0,-22) EOB

EOB (end of block) means no more non-zero values

Run-Length Encoding



Efficiently store all these zeros

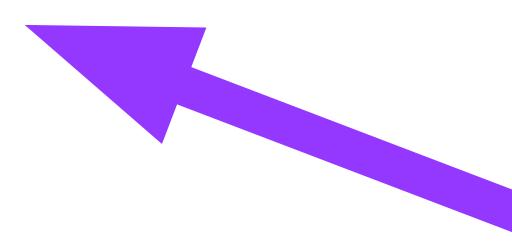
In each entry store
(# zeros, next non-zero value)

(0,-512) (0,77) (0,24) (0,14) (0,24) (3,-13)
(3,-16) (0,19) (3,-24) (0,-22) EOB

EOB (end of block) means no
more non-zero values

Huffman Code (Finally!)

(0,-512) (0,77) (0,24) (0,14) (0,24) (3,-13)
(3,-16) (0,19) (3,-24) (0,-22) EOB



These RLE values are symbols

Build a Huffman code
across the image

